

Assignment-3_Apply-KNN

February 27, 2019

1 Assignment-3: Apply K-NN on Amazon Fine Food Reviews DataSet

2 1.0 Introduction

(i).The k-nearest neighbors algorithm (k-NN) is a non-parametric method used for classification and regression predictive problem.

(ii).It is more widely used in classification problems in the industry.

3 2.0 Objective

To Predict the Polarity of Amazon Fine Food Review Using K-Nearst Neighbour Algorithm.

4 3.0 Importing All Required Library

```
In [1]: %matplotlib inline
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import cross_val_score
```

```

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

```

```

from tqdm import tqdm
import os
import warnings
warnings.filterwarnings("ignore")

```

```

C:\Users\User\Anaconda3\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows;
warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

```

5 4.0 Importing Amazon Fine Food Review Dataset

```

In [2]: if os.path.isfile("final.sqlite"):
        conn=sqlite3.connect("final.sqlite")
        Data=pd.read_sql_query("select * from Reviews where Score!=3",conn)
        conn.close()
    else :
        print("Error Importing the file")

```

```

In [3]: # Printing some data of DataFrame

```

```
Data.head(2)
```

```

Out[3]:
   index  Id  ProductId  UserId  ProfileName \
0  138706  150524  0006641040  ACITT7DI6IDDL  shari zychinski
1  138688  150506  0006641040  A2IW4PEEK02R0U          Tracy

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      0                      0       1  939340800
1                      1                      1       1  1194739200

   Summary \
0  EVERY book is educational
1  Love the book, miss the hard cover version

   Text \
0  this witty little book makes my son laugh at l...
1  I grew up reading these Sendak books, and watc...

   CleanedText
0  witti littl book make son laugh loud recit car...
1  grew read sendak book watch realli rosi movi i...

```

6 5.0 Information About DataSet

```
In [4]: print("\nNumber of Reviews: ",Data["Text"].count())
        print("\nNumber of Users: ",len(Data["UserId"].unique())) # Unique returns 1-D array o
        print("\nNumber of Products: ",len(Data["ProductId"].unique()))
        print("\nShape of Data: ", Data.shape)
        print("\nColumn Name of DataSet : ",Data.columns)
        print("\n\nNumber of Attributes/Columns in data: 12")
        print("\nNumber of Positive Reviews : ", Data['Score'].value_counts()[1])
        print("\nNumber of Negative Reviews : ", Data['Score'].value_counts()[0])
```

Number of Reviews: 364171

Number of Users: 243414

Number of Products: 65442

Shape of Data: (364171, 12)

Column Name of DataSet : Index(['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
'Summary', 'Text', 'CleanedText'],
dtype='object')

Number of Attributes/Columns in data: 12

Number of Positive Reviews : 307061

Number of Negative Reviews : 57110

```
In [5]: print("\nNumber of Reviews: ",Data["Text"].count())
```

Number of Reviews: 364171

6.0.1 5.1 Attribute Information About DataSet

- 1.Id - A unique value starts from 1
- 2.ProductId - A unique identifier for the product
- 3.UserId - A unique identifier for the user
- 4.ProfileName - Name of user profile
- 5.HelpfulnessNumerator - Number of users who found the review helpful
- 6.HelpfulnessDenominator - Number of users who indicated whether they found the review helpful or not
- 7.Score - Rating 0 or 1

- 8. Time - Timestamp for the review
- 9. Summary - Brief summary of the review
- 10. Text - Text of the review
- 11. Cleaned Text - Text that only alphabets

6.1 6.0 Due to Limited Hardware Resource we will limit our analysis on 20000 data points only.

In [6]: *# To randomly sample 10k points from both class*

```
data_pos = Data[Data["Score"] == 1].sample(n = 10000)
data_neg = Data[Data["Score"] == 0].sample(n = 10000)
final_20k = pd.concat([data_pos, data_neg])
final_20k.shape
```

Out[6]: (20000, 12)

```
In [7]: Y = final_20k['Score'].values
        X = final_20k['CleanedText'].values
```

6.1.1 7.0 Splitting DataSet into Train and Test Data

```
In [8]: from sklearn.model_selection import train_test_split
        # X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, shuffle=False)
        X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33) # this is random

        print("Shape of Train and Test Dataset")
        print(X_train.shape, y_train.shape)
        print(X_test.shape, y_test.shape)
```

```
Shape of Train and Test Dataset
(13400,) (13400,)
(6600,) (6600,)
```

7 8.0 Defining Some Function

7.0.1 8.1 Train Data Confusion Matrix Plot

```
In [9]: def trainconfusionmatrix(knn,X_train,y_train):
        print("Confusion Matrix for Train set")
        cm=confusion_matrix(y_train, knn.predict(X_train))
        class_label = ["negative", "positive"]
        df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
        sns.heatmap(df_cm, annot = True, fmt = "d")
        plt.title("Train Confusiion Matrix")
        plt.xlabel("Predicted Label")
        plt.ylabel("True Label")
        plt.show()
```

7.0.2 8.2 Test Data Confusion Matrix Plot

```
In [10]: def testconfusionmatrix(knn,X_test,y_test):
    print("Confusion Matrix for Test set")
    cm=confusion_matrix(y_test, knn.predict(X_test))
    class_label = ["negative", "positive"]
    df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
    sns.heatmap(df_cm, annot = True, fmt = "d")
    plt.title("Test Confusiion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()
```

7.0.3 8.3 ROC-AUC Curve Plot

```
In [11]: def plot_auc_roc(knn,X_train,X_test,y_train,y_test):
    train_fpr, train_tpr, thresholds = roc_curve(y_train, neigh.predict_proba(X_train)[:,1])
    test_fpr, test_tpr, thresholds = roc_curve(y_test, neigh.predict_proba(X_test)[:,1])

    plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
    plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
    plt.legend()
    plt.xlabel("K: hyperparameter")
    plt.ylabel("AUC")
    plt.title("ROC CURVE PLOTS")
    plt.show()
```

7.0.4 8.4 Error Plot

```
In [12]: def error_plot(neighbours,mse):
    plt.title('Error Plot')
    plt.xlabel('K')
    plt.ylabel('misscalssifiaction error')
    plt.plot(neighbours,mse)
```

7.0.5 8.5 Cross Validation Using Kd Tree Algorithm

```
In [13]: def knn_cv_kd(X_train,y_train,neighbours):

    cv_scores=[]

    for k in neighbours:
        knn = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree')
        scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy') #
        cv_scores.append(scores.mean())

    mse = [1-x for x in cv_scores] # calculating misscalssification_error = 1 - accuracy

    best_k = neighbours[mse.index(min(mse))] #returns k corresponding to minimum error
```

```
return mse,best_k
```

7.0.6 8.6 Cross Validation Using Brute Algorithm

```
In [14]: def knn_cv_brute(X_train,y_train,neighbours):
```

```
    cv_scores=[]
```

```
    for k in neighbours:
```

```
        knn = KNeighborsClassifier(n_neighbors=k,algorithm='brute')
```

```
        scores = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy') #
```

```
        cv_scores.append(scores.mean())
```

```
    mse = [1-x for x in cv_scores] # calculating missclassification_error = 1 - accuracy
```

```
    best_k = neighbours[mse.index(min(mse))] #returns k corresponding to minimum error
```

```
    return mse,best_k
```

7.0.7 8.7 Accuracy

```
In [15]: def accuracy(model,X_train,Y_train,X_test,Y_test):
```

```
    prediction=model.predict(X_test)
```

```
    training_accuracy=model.score(X_train,Y_train) # accuracy_score(y_train ,neigh.p
```

```
    training_error=1-training_accuracy
```

```
    test_accuracy=accuracy_score(Y_test,prediction)
```

```
    test_error=1-test_accuracy
```

```
    return training_accuracy,training_error,test_accuracy,test_error
```

8 9.0 Bags of Words

```
In [16]: vectorizer = CountVectorizer()
```

```
    vectorizer.fit(X_train) # fit has to happen only on train data
```

```
    # we use the fitted CountVectorizer to convert the text to vector
```

```
    X_train_bow = vectorizer.transform(X_train)
```

```
    X_test_bow = vectorizer.transform(X_test)
```

```
    print("Shape of Train , Test and Cross Validation Data After vectorizations")
```

```
    print(X_train_bow.shape, y_train.shape)
```

```
    print(X_test_bow.shape, y_test.shape)
```

```
Shape of Train , Test and Cross Validation Data After vectorizations
```

```
(13400, 15042) (13400,)
```

```
(6600, 15042) (6600,)
```

```
In [17]: type(X_train_bow)
```

```
Out[17]: scipy.sparse.csr.csr_matrix
```

8.0.1 9.1 Brute Force Algorithm


9.1.1 Finding Optimal Value of Hyperparameter(k)

```
In [18]: import numpy as np
```

```
neighbours=np.arange(1,100,2)  
mse,best_k = knn_cv_brute(X_train_bow,y_train,neighbours)
```

```
In [19]: error_plot(neighbours,mse)  
print("Best value of K found for Brute Force Algorithm Implementaion is : ",best_k)
```

Best value of K found for Brute Force Algorithm Implementaion is : 55



9.1.2 Training the model

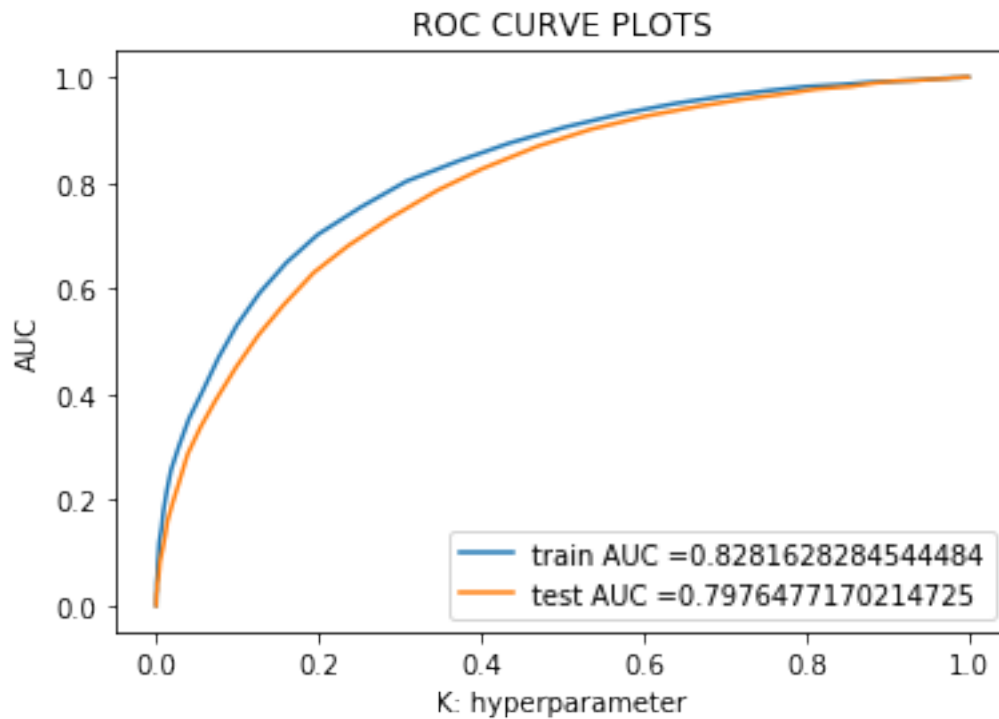
```
In [20]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='brute')
        neigh.fit(X_train_bow, y_train)
```

```
Out[20]: KNeighborsClassifier(algorithm='brute', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=55, p=2,
                             weights='uniform')
```

9.1.3 Evaluating the performance of model

(A). Roc-Auc Plot


```
In [21]: plot_auc_roc(neigh,X_train_bow,X_test_bow,y_train,y_test)
```



(B). Confusion Matrix Plot on Train Data

```
In [22]: trainconfusionmatrix(neigh,X_train_bow,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix on Test Data

```
In [23]: testconfusionmatrix(neigh,X_test_bow,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [24]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_bow)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 74.0

Train Error in %: 26.0

Test Accuracy in %: 71.0

Test Error in % : 28.999999999999996

(E). Classification Report

```
In [25]: print("Classification Report: \n")
prediction=neigh.predict(X_test_bow)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.60	0.68	3333
1	0.67	0.83	0.74	3267
micro avg	0.71	0.71	0.71	6600
macro avg	0.72	0.71	0.71	6600
weighted avg	0.72	0.71	0.71	6600

8.0.2 9.2 KD-Tree Algorithm


9.2.1 Finding Optimal Value of Hyperparameter(k)

```
In [26]: import numpy as np
```

```
neighbours=np.arange(1,100,2)  
mse,best_k = knn_cv_kd(X_train_bow,y_train,neighbours)
```

```
In [27]: error_plot(neighbours,mse)  
print("Best value of K found for KD Tree Algorithm Implementaion is : ",best_k)
```

Best value of K found for KD Tree Algorithm Implementaion is : 55



9.2.2 Training the model


```
In [28]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='kd_tree')  
        neigh.fit(X_train_bow, y_train)
```

```
Out[28]: KNeighborsClassifier(algorithm='kd_tree', leaf_size=30, metric='minkowski',  
                             metric_params=None, n_jobs=None, n_neighbors=55, p=2,  
                             weights='uniform')
```

9.2.3 Evaluating the performance of model

(A). Roc-Auc Plot


```
In [29]: plot_auc_roc(neigh,X_train_bow,X_test_bow,y_train,y_test)
```



(B). Confusion Matrix Plot on Train Data

```
In [30]: trainconfusionmatrix(neigh,X_train_bow,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix Plot on Test Data

```
In [31]: testconfusionmatrix(neigh,X_test_bow,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [32]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_bow)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 74.0

Train Error in %: 26.0

Test Accuracy in %: 71.0

Test Error in % : 28.999999999999996

(E). Classification Report

```
In [33]: print("Classification Report: \n")
prediction=neigh.predict(X_test_bow)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.60	0.68	3333
1	0.67	0.83	0.74	3267
micro avg	0.71	0.71	0.71	6600
macro avg	0.72	0.71	0.71	6600
weighted avg	0.72	0.71	0.71	6600

9 10.0 TF-IDF

```
In [34]: vectorizer = TfidfVectorizer(ngram_range=(1,2))
         vectorizer.fit(X_train) # fit has to happen only on train data

         # we use the fitted CountVectorizer to convert the text to vector
         X_train_TF = vectorizer.transform(X_train)
         X_test_TF = vectorizer.transform(X_test)
```

```
In [35]: print("After vectorizations")
         print(X_train_TF.shape, y_train.shape)
         print(X_test_TF.shape, y_test.shape)
```

```
After vectorizations
(13400, 304884) (13400,)
(6600, 304884) (6600,)
```

9.0.1 10.1 Brute Force Algorithm


10.1.1 Finding Optimal Value of Hyperparameter(k)

```
In [36]: import numpy as np

         neighbours=np.arange(1,100,2)
         mse,best_k = knn_cv_brute(X_train_TF,y_train,neighbours)

In [37]: error_plot(neighbours,mse)
         print("Best value of K found for Brute Force Algorithm Implementaion is : ",best_k)
```

Best value of K found for Brute Force Algorithm Implementaion is : 97



10.1.2 Training the model


```
In [38]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='brute')
         neigh.fit(X_train_TF, y_train)
```

```
Out[38]: KNeighborsClassifier(algorithm='brute', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=97, p=2,
                             weights='uniform')
```

10.1.3 Evaluating the performance of model

(A). Roc-Auc Plot


```
In [39]: plot_auc_roc(neigh,X_train_TF,X_test_TF,y_train,y_test)
```

(B). Confusion Matrix Plot on Train Data

```
In [40]: trainconfusionmatrix(neigh,X_train_TF,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix Plot on Test Data

```
In [41]: testconfusionmatrix(neigh,X_test_TF,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [42]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_TF)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 80.0

Train Error in %: 20.0

Test Accuracy in %: 80.0

Test Error in % : 20.0

(E). Classification Report

```
In [43]: print("Classification Report: \n")
prediction=neigh.predict(X_test_TF)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.76	0.79	3333
1	0.78	0.83	0.80	3267
micro avg	0.80	0.80	0.80	6600
macro avg	0.80	0.80	0.80	6600
weighted avg	0.80	0.80	0.80	6600

9.0.2 10.2 KD Tree Algorithm


10.2.1 Finding Optimal Value of Hyperparameter(k)

```
In [44]: import numpy as np
```

```
neighbours=np.arange(1,100,2)  
mse,best_k = knn_cv_kd(X_train_TF,y_train,neighbours)
```

```
In [45]: error_plot(neighbours,mse)  
print("Best value of K found for KD Tree Algorithm Implementaion is : ",best_k)
```

Best value of K found for KD Tree Algorithm Implementaion is : 97



10.2.2 Training the model


```
In [46]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='kd_tree')
        neigh.fit(X_train_TF, y_train)
```

```
Out[46]: KNeighborsClassifier(algorithm='kd_tree', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=97, p=2,
                             weights='uniform')
```

10.2.3 Evaluating the performance of model

(A). Roc-Auc Plot


```
In [47]: plot_auc_roc(neigh,X_train_TF,X_test_TF,y_train,y_test)
```



(B). Confusion Matrix Plot on Train Data

```
In [48]: trainconfusionmatrix(neigh,X_train_TF,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix Plot on Test Data

```
In [49]: testconfusionmatrix(neigh,X_test_TF,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [50]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_TF)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 80.0

Train Error in %: 20.0

Test Accuracy in %: 80.0

Test Error in % : 20.0

(E). Classification Report

```
In [51]: print("Classification Report: \n")
prediction=neigh.predict(X_test_TF)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.76	0.79	3333
1	0.78	0.83	0.80	3267
micro avg	0.80	0.80	0.80	6600
macro avg	0.80	0.80	0.80	6600
weighted avg	0.80	0.80	0.80	6600

10 11.0 Word To Vector

```
In [16]: list_of_Train_sent=[]
         list_of_Test_sent=[]

         for sent in X_train:
             list_of_Train_sent.append(sent.split())

         for sent in X_test:
             list_of_Test_sent.append(sent.split())
```

```
In [17]: len(list_of_Train_sent)
```

```
Out[17]: 13400
```

```
In [18]: Train_model=Word2Vec(list_of_Train_sent,min_count=5,size=50, workers=4)
         Test_model=Word2Vec(list_of_Test_sent,min_count=5,size=50, workers=4)
```

10.0.1 11.1 Avg Word2Vec

```
In [55]: import numpy as np

         Train_vectors = []
         for sent in list_of_Train_sent:
             sent_vec = np.zeros(50)
             cnt_words = 0
             for word in sent:
                 try:
                     vec = Train_model.wv[word]
                     sent_vec += vec
                     cnt_words += 1
                 except:
                     pass
             if cnt_words!=0:
                 sent_vec /= cnt_words
```



```

Train_vectors.append(sent_vec)
Train_vectors = np.nan_to_num(Train_vectors)

```

In [56]: `import numpy as np`

```

Test_vectors=[]
for sent in list_of_Test_sent:
    sent_vec=np.zeros(50)
    cnt_words=0
    for word in sent:
        try:
            vec=Test_model.wv[word]
            sent_vec+=vec
            cnt_words+=1
        except:
            pass
    if cnt_words!=0:
        sent_vec/=cnt_words
    Test_vectors.append(sent_vec)
Test_vectors=np.nan_to_num(Test_vectors)

```

In [57]: `print("Shape of Test Vectors : ",Test_vectors.shape)`

Shape of Test Vectors : (6600, 50)

In [58]: `X_train_AWV = Train_vectors`
`X_test_AWV = Test_vectors`

In [59]: `print(X_train_AWV.shape, y_train.shape)`
`print(X_test_AWV.shape, y_test.shape)`

(13400, 50) (13400,)
(6600, 50) (6600,)

10.0.2 10.2 Brute Force Algorithm

10.2.1 Finding Optimal Value of Hyperparameter(k)

In [60]: `import numpy as np`


```

neighbours=np.arange(1,100,2)
mse,best_k = knn_cv_brute(X_train_AWV,y_train,neighbours)

```

In [61]: `error_plot(neighbours,mse)`
`print("Best value of K found for Brute Force Algorithm Implementaion is : ",best_k)`

Best value of K found for Brute Force Algorithm Implementaion is : 37



10.2.2 Training the model


```
In [62]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='brute')
         neigh.fit(X_train_AWV, y_train)
```

```
Out[62]: KNeighborsClassifier(algorithm='brute', leaf_size=30, metric='minkowski',
                             metric_params=None, n_jobs=None, n_neighbors=37, p=2,
                             weights='uniform')
```

10.2.3 Evaluating the performance of model

(A). Roc-Auc Plot


```
In [63]: plot_auc_roc(neigh,X_train_AWV,X_test_AWV,y_train,y_test)
```



(B). Confusion Matrix Plot on Train Data

```
In [64]: trainconfusionmatrix(neigh,X_train_AWV,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix Plot on Test Data

```
In [65]: testconfusionmatrix(neigh,X_test_AWV,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [66]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_AWV)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 78.0

Train Error in %: 22.0

Test Accuracy in %: 57.99999999999999

Test Error in % : 42.0

(E). Classification Report

```
In [67]: print("Classification Report: \n")
prediction=neigh.predict(X_test_AWV)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.33	0.44	3333
1	0.55	0.84	0.66	3267
micro avg	0.58	0.58	0.58	6600
macro avg	0.61	0.58	0.55	6600
weighted avg	0.61	0.58	0.55	6600

10.0.3 10.3 KD Tree Algorithm


10.3.1 Finding Optimal Value of Hyperparameter(k)

```
In [68]: import numpy as np
```

```
neighbours=np.arange(1,100,2)  
mse,best_k = knn_cv_kd(X_train_AWV,y_train,neighbours)
```

```
In [69]: error_plot(neighbours,mse)  
print("Best value of K found for KD Tree Algorithm Implementaion is : ",best_k)
```

Best value of K found for KD Tree Algorithm Implementaion is : 37



10.3.2 Training the model


```
In [70]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='kd_tree')  
        neigh.fit(X_train_AWV, y_train)
```

```
Out[70]: KNeighborsClassifier(algorithm='kd_tree', leaf_size=30, metric='minkowski',  
                             metric_params=None, n_jobs=None, n_neighbors=37, p=2,  
                             weights='uniform')
```

10.3.3 Evaluating the performance of model

(A). Roc-Auc Plot


```
In [71]: plot_auc_roc(neigh,X_train_AWV,X_test_AWV,y_train,y_test)
```



(B). Confusion Matrix Plot on Train Data

```
In [72]: trainconfusionmatrix(neigh,X_train_AWV,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix Plot on Test Data

```
In [73]: testconfusionmatrix(neigh,X_test_AWV,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [74]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_AWV)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 78.0

Train Error in %: 22.0

Test Accuracy in %: 57.999999999999999

Test Error in % : 42.0

(E). Classification Report

```
In [75]: print("Classification Report: \n")
prediction=neigh.predict(X_test_AWV)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.33	0.44	3333
1	0.55	0.84	0.66	3267
micro avg	0.58	0.58	0.58	6600
macro avg	0.61	0.58	0.55	6600
weighted avg	0.61	0.58	0.55	6600

11 11.0 TF-IDF Word To Vector

```
In [19]: model = TfidfVectorizer()
         model.fit(X_train)
         X_Train_TF = model.transform(X_train)
         X_Test_TF = model.transform(X_test)
```

```
In [20]: print("Shape of Train Data After TFIDF : ",X_Train_TF.shape)
         print("Shape of Test Data After TFIDF : ",X_Test_TF.shape)
```

```
Shape of Train Data After TFIDF : (13400, 15367)
Shape of Test Data After TFIDF : (6600, 15367)
```

```
In [21]: TFIDF_Feature=model.get_feature_names()
         print(len(TFIDF_Feature))
         print(TFIDF_Feature[0:20])
```

```
15367
```

```
['aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaa', 'aaaaaaaaagghh', 'aaaaahhhhhhhhhhhhhhhhh', 'aa
```

```
In [22]: from tqdm import tqdm
         Train_TFIDF_W2V_Vectors=[]
         row=0
         for sent in tqdm(list_of_Train_sent):
             sent_vec=np.zeros(50)
             weight=0
             for word in sent:
                 try :
                     w2v_vec=Train_model.wv[word]
                     tfidf_vec=X_Train_TF[row,TFIDF_Feature.index(word)]
                     sent_vec+=(w2v_vec*tfidf_vec)
                     weight+=tfidf_vec
```

```

        except :
            pass
    if weight!=0:
        sent_vec/=weight
    Train_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1

```

100%|| 13400/13400 [02:56<00:00, 75.98it/s]

```

In [23]: Test_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Test_sent):
    sent_vec=np.zeros(50)
    weight=0

    for word in sent:
        try:
            w2v_vec=Test_model.wv[word]
            tfidf_vec=X_Test_TF(row,TFIDF_Feature.index(word))
            sent_vec+=(w2v_vec*tfidf_vec)
            weight+=tfidf

        except :
            pass

    if weight!=0:
        sent_vec/=weight
    Test_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1

```

100%|| 6600/6600 [01:03<00:00, 103.35it/s]

```

In [25]: Test_tfidfw2v_vectors=np.nan_to_num(Test_TFIDF_W2V_Vectors)
Train_tfidfw2v_vectors=np.nan_to_num(Train_TFIDF_W2V_Vectors)

```

```

In [28]: X_train_TfIdfW2v=Train_tfidfw2v_vectors
X_test_TfIdfW2v=Test_tfidfw2v_vectors

```

11.0.1 11.1 Brute Force Algorithm

11.1.1 Finding Optimal Value of Hyperparameter(k)

```


In [29]: import numpy as np

neighbours=np.arange(1,100,2)
mse,best_k = knn_cv_brute(X_train_TfIdfW2v,y_train,neighbours)

```

```
In [30]: error_plot(neighbours,mse)
         print("Best value of K found for Brute Force Algorithm Implementaion is : ",best_k)
```

Best value of K found for Brute Force Algorithm Implementaion is : 95



11.1.2 Training the model


```
In [33]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='brute')
         neigh.fit(X_train_TfidfW2v, y_train)
```

```
Out [33]: KNeighborsClassifier(algorithm='brute', leaf_size=30, metric='minkowski',
                               metric_params=None, n_jobs=None, n_neighbors=95, p=2,
                               weights='uniform')
```

11.1.3 Evaluating the performance of model

(A). Roc-Auc Plot

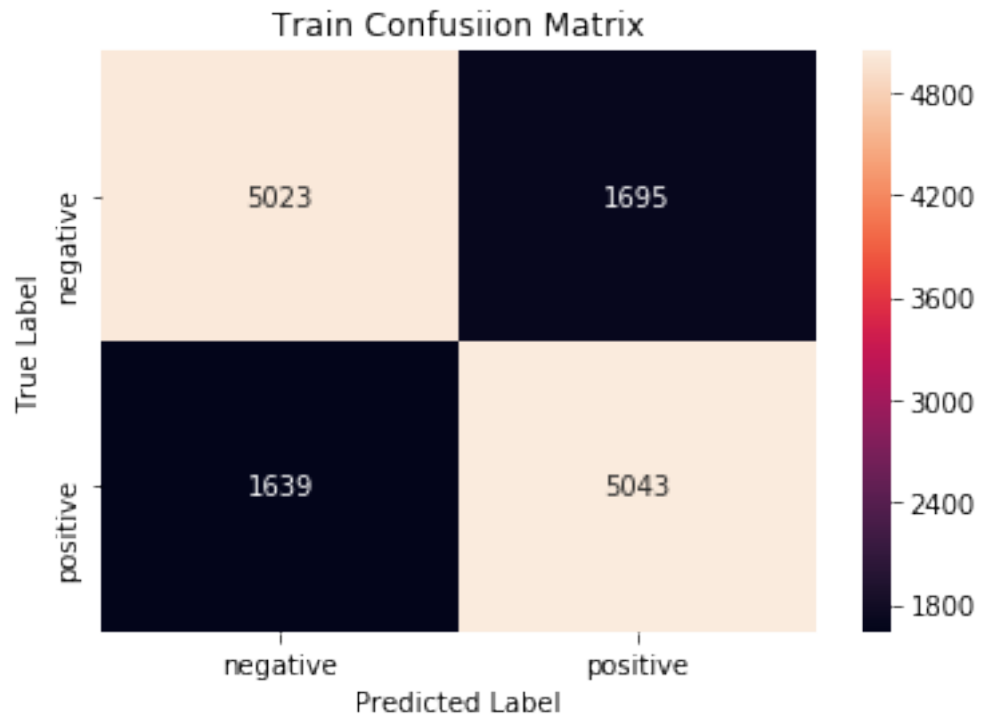
```
In [34]: plot_auc_roc(neigh,X_train_TfidfW2v,X_test_TfidfW2v,y_train,y_test)
```



(B). Confusion Matrix Plot on Train Data

```
In [35]: trainconfusionmatrix(neigh,X_train_TfIdfW2v,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix Plot on Test Data

```
In [36]: testconfusionmatrix(neigh,X_test_TfIdfW2v,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [37]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_TfIdfW2v)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 75.0

Train Error in %: 25.0

Test Accuracy in %: 50.0

Test Error in % : 50.0

(E). Classification Report

```
In [38]: print("Classification Report: \n")
prediction=neigh.predict(X_test_TfIdfW2v)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.50	1.00	0.66	3282
1	0.00	0.00	0.00	3318
micro avg	0.50	0.50	0.50	6600
macro avg	0.25	0.50	0.33	6600
weighted avg	0.25	0.50	0.33	6600

11.0.2 11.2 KD Tree Algorithm


11.2.1 Finding Optimal Value of Hyperparameter(k)

```
In [39]: import numpy as np
```

```
neighbours=np.arange(1,100,2)  
mse,best_k = knn_cv_kd(X_train_TfIdfW2v,y_train,neighbours)
```

```
In [40]: error_plot(neighbours,mse)  
print("Best value of K found for KD Tree Algorithm Implementaion is : ",best_k)
```

Best value of K found for KD Tree Algorithm Implementaion is : 95



11.2.2 Training the model


```
In [41]: neigh = KNeighborsClassifier(n_neighbors = best_k,algorithm='kd_tree')  
        neigh.fit(X_train_TfidfW2v, y_train)
```

```
Out[41]: KNeighborsClassifier(algorithm='kd_tree', leaf_size=30, metric='minkowski',  
                             metric_params=None, n_jobs=None, n_neighbors=95, p=2,  
                             weights='uniform')
```

11.2.3 Evaluating the performance of model

(A). Roc-Auc Plot


```
In [42]: plot_auc_roc(neigh,X_train_TfidfW2v,X_test_TfidfW2v,y_train,y_test)
```



(B). Confusion Matrix Plot on Train Data

```
In [43]: trainconfusionmatrix(neigh,X_train_TfidfW2v,y_train)
```


Confusion Matrix for Train set



(C). Confusion Matrix Plot on Test Data

```
In [44]: testconfusionmatrix(neigh,X_test_TfIdfW2v,y_test)
```

Confusion Matrix for Test set



(D). Accuracy and Error

```
In [45]: training_accuracy,training_error,test_accuracy,test_error = accuracy(neigh,X_train_TfIdfW2v)
print("Training Accuracy in %: ", round(training_accuracy,2)*100)
print("\nTrain Error in %: ", round(training_error,2)*100)
print("\nTest Accuracy in %: ", round(test_accuracy,2)*100)
print("\nTest Error in % : ", round(test_error,2)*100)
```

Training Accuracy in %: 75.0

Train Error in %: 25.0

Test Accuracy in %: 50.0

Test Error in % : 50.0

(E). Classification Report

```
In [46]: print("Classification Report: \n")
prediction=neigh.predict(X_test_TfIdfW2v)
print(classification_report(y_test, prediction))
```

Classification Report:

	precision	recall	f1-score	support
0	0.50	1.00	0.66	3282
1	0.00	0.00	0.00	3318
micro avg	0.50	0.50	0.50	6600
macro avg	0.25	0.50	0.33	6600
weighted avg	0.25	0.50	0.33	6600

12 12.0 Conclusion :

1. Report On Brute Force Algorithm

```
In [51]: from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["Vectorizer", "Hyperparameter", "AUC", "Train Accuracy (%)", "Test Accuracy (%)"]
```

```
x.add_row(["BOW", 55, 0.79, 74, 71])
```

```
x.add_row(["TF-IDF", 97, 0.88, 80, 80])
```

```
x.add_row(["W2V", 37, 0.64, 78, 57])
```

```
x.add_row(["TF-IDF W2V", 95, 0.50, 75, 50])
```

```
print(x)
```

Vectorizer	Hyperparameter	AUC	Train Accuracy (%)	Test Accuracy (%)
BOW	55	0.79	74	71
TF-IDF	97	0.88	80	80
W2V	37	0.64	78	57
TF-IDF W2V	95	0.5	75	50

2. Report on KD Tree Algorithm

```
In [52]: x = PrettyTable()
```

```
x.field_names = ["Vectorizer", "Hyperparameter", "AUC", "Train Accuracy (%)", "Test Accuracy (%)"]
```

```
x.add_row(["BOW", 55, 0.79, 74, 71])
```

```
x.add_row(["TF-IDF", 97, 0.88, 80, 80])
```

```
x.add_row(["W2V", 37, 0.64, 78, 57])
```

```
x.add_row(["TF-IDF W2V",95,0.50,75,50])
print(x)
```

Vectorizer	Hyperparameter	AUC	Train Accuracy (%)	Test Accuracy (%)
BOW	55	0.79	74	71
TF-IDF	97	0.88	80	80
W2V	37	0.64	78	57
TF-IDF W2V	95	0.5	75	50

3. 71 % Of Accuracy is achieved by the model in case of Bag of Words.

4. 80 % Of Accuracy is achieved by the model in case of TF-IDF Vectorizer.

5. Model is overfit in case of Average Word to vector.

6. Very small subset of Data is taken but still it took more time due to large dimension and time complexity of KNN.

7. Model behaviour in TF-IDF W2V is lenient towards one class .