

Assignment-5: Apply Support Vector Machine On Amazon Fine Food Reviews DataSet

Introduction

(i). "Support Vector Machine" (SVM) is a supervised machine learning algorithm which can be used for both classification or regression challenges. However, it is mostly used in classification problems.

(ii). The main goal of this algorithm is to find the optimal hyperplane.

Objective

To Predict the Polarity of Amazon Fine Food Review Using Support Vector Machine Algorithm.

Importing All Required Library

```

In [1]: %matplotlib inline
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
import math

from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import SGDClassifier
from sklearn.svm import SVC
from sklearn.calibration import CalibratedClassifierCV

from sklearn.metrics import classification_report
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import TimeSeriesSplit

from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import cross_val_score
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score

from gensim.models import Word2Vec
from gensim.models import KeyedVectors

import pickle

from tqdm import tqdm
import os
import warnings
warnings.filterwarnings("ignore")

```

C:\Users\User\Anaconda3\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows; aliasing chunkize to chunkize_serial
 warnings.warn("detected Windows; aliasing chunkize to chunkize_serial")

Importing Amazon Fine Food Review Dataset

```

In [2]: if os.path.isfile("final.sqlite"):
        conn=sqlite3.connect("final.sqlite")
        Data=pd.read_sql_query("select * from Reviews where Score!=3",conn)
        conn.close()
    else :
        print("Error Importing the file")

```

In [3]: `# Printing some data of DataFrame`

```
Data['Score'].value_counts()
```

Out[3]: 1 307061
0 57110
Name: Score, dtype: int64

Information About DataSet

```
In [4]: print("\nNumber of Reviews: ",Data["Text"].count())
print("\nNumber of Users: ",len(Data["UserId"].unique())) # Unique returns 1-D
array of unique values of DataFrame.
print("\nNumber of Products: ",len(Data["ProductId"].unique()))
print("\nShape of Data: ", Data.shape)
print("\nColumn Name of DataSet : ",Data.columns)
print("\n\nNumber of Attributes/Columns in data: 12")
print("\nNumber of Positive Reviews : ", Data['Score'].value_counts()[1])
print("\nNumber of Negative Reviews : ", Data['Score'].value_counts()[0])
```

Number of Reviews: 364171

Number of Users: 243414

Number of Products: 65442

Shape of Data: (364171, 12)

Column Name of DataSet : Index(['index', 'Id', 'ProductId', 'UserId', 'ProfileName',
'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
'Summary', 'Text', 'CleanedText'],
dtype='object')

Number of Attributes/Columns in data: 12

Number of Positive Reviews : 307061

Number of Negative Reviews : 57110

```
In [5]: print("\nNumber of Reviews: ",Data["Text"].count())
```

Number of Reviews: 364171

Attribute Information About DataSet

- 1.Id - A unique value starts from 1
- 2.ProductId - A unique identifier for the product
- 3.UserId - A unique identifier for the user
- 4.ProfileName - Name of user profile
- 5.HelpfulnessNumerator - Number of users who found the review helpful
- 6.HelpfulnessDenominator - Number of users who indicated whether they found the review helpful or not
- 7.Score - Rating 0 or 1
- 8.Time - Timestamp for the review
- 9.Summary - Brief summary of the review
- 10.Text - Text of the review
- 11.Cleaned Text - Text that only alphabets

```
In [6]: # Sorting on the basis of Time Parameter
Data.sort_values('Time',inplace=True)
```

```
In [7]: Data=Data.head(100000)
```

```
In [8]: Y = Data['Score']
X = Data['CleanedText']
```

Splitting DataSet into Train and Test Data

```
In [9]: from sklearn.model_selection import train_test_split
# X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, sh
uffle=False): this is for time series split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33,shuff
le=False) # this is random splitting

print("Shape of Train and Test Dataset for 100k points")
print(X_train.shape, Y_train.shape)
print(X_test.shape, Y_test.shape)
```

```
Shape of Train and Test Dataset for 100k points
(67000,) (67000,)
(33000,) (33000,)
```

Defining Some Function

Train Data Confusion Matrix Plot

```
In [10]: def trainconfusionmatrix(model,X_train,y_train):
    print("Confusion Matrix for Train set")
    cm=confusion_matrix(y_train, model.predict(X_train))
    class_label = ["negative", "positive"]
    df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
    sns.heatmap(df_cm, annot = True, fmt = "d")
    plt.title("Train Confusiion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()
```

Test Data Confusion Matrix Plot

```
In [11]: def testconfusionmatrix(model,X_test,y_test):
    print("Confusion Matrix for Test set")
    cm=confusion_matrix(y_test, model.predict(X_test))
    class_label = ["negative", "positive"]
    df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
    sns.heatmap(df_cm, annot = True, fmt = "d")
    plt.title("Test Confusiion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()
```

ROC-AUC Curve (RBF SVM)

```
In [12]: def plot_auc_roc(model,X_train,X_test,y_train,y_test):
    train_fpr, train_tpr, thresholds = roc_curve(y_train,model.predict_proba(X_train)[:,:1])
    test_fpr, test_tpr, thresholds = roc_curve(y_test,model.predict_proba(X_test)[:,:1])

    plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
    plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
    plt.legend()
    plt.xlabel("Hyperameter (C)")
    plt.ylabel("AUC")
    plt.title("ROC CURVE PLOTS")
    plt.show()
```

Plot Linear

```
In [13]: def plot_l(gsv):
    res=gsv.cv_results_
    cv_data=res['mean_test_score']
    train_data=res['mean_train_score']
    alpha=[0.0001,0.001,0.01,0.1,1,10,100,1000,10000]
    log_c= [math.log10(x) for x in alpha]

    plt.xlabel("Hyperparameter (Log(Alpha))",fontsize=15)
    plt.ylabel("Mean_Score",fontsize=15)
    plt.title('Mean Score v/s C',fontsize=15)
    plt.plot(log_c,cv_data,label="Cross-Validation")
    plt.plot(log_c,train_data,label="Train")
    plt.legend()
    plt.show()
```

Plot RBF

```
In [14]: def plot_r(gsv):
    C=[0.0001,0.001,0.01,0.1,1,10,100,1000,10000]
    res=gsv.cv_results_
    cv_data=res['mean_test_score']
    train_data=res['mean_train_score']
    log_c= [math.log10(x) for x in C]

    plt.xlabel("Hyperparameter (Log(C))",fontsize=15)
    plt.ylabel("Mean_Score",fontsize=15)
    plt.title('Mean Score v/s C',fontsize=15)
    plt.plot(log_c,cv_data,label="Cross-Validation")
    plt.plot(log_c,train_data,label="Train")
    plt.legend()
    plt.show()
```

ROC-AUC Curve (Linear SVM)

```
In [15]: def plot_auc_roc_L(model,X_train,X_test,y_train,y_test):

    Clf = CalibratedClassifierCV(model,cv='prefit')
    Clf.fit(X_train,y_train)
    train_fpr, train_tpr, thresholds = roc_curve(y_train,Clf.predict_proba(X_train)[:,-1])
    test_fpr, test_tpr, thresholds = roc_curve(y_test,Clf.predict_proba(X_test)[:,-1])

    plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
    plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
    plt.legend()
    plt.xlabel("Hyperparameter (C)")
    plt.ylabel("AUC")
    plt.title("ROC CURVE PLOTS")
    plt.show()
```

GridSearchCV Linear

```
In [16]: def Grid_SearchCV(X_train,Y_train,regularisation_parameter):

    param_grid= {'alpha':[0.0001,0.001,0.01,0.1,1,10,100,1000,10000],
                  'penalty':[regularisation_parameter]}
    tscv = TimeSeriesSplit(n_splits=10)

    gsv=GridSearchCV(SGDClassifier(shuffle=False,loss='hinge',class_weight="balanced"), param_grid, scoring = 'roc_auc', cv=tscv,n_jobs = -1,pre_dispatch=2)
    gsv.fit(X_train,Y_train)

    return gsv
```

GridSearchCV RBF

```
In [17]: def Grid_SearchCV_RBF(X_train,Y_train):

    param_grid = {'C':[0.0001,0.001,0.01,0.1,1,10,100,1000,10000]}
    tscv = TimeSeriesSplit(n_splits=5)

    gsv=GridSearchCV(SVC(class_weight="balanced"),param_grid,scoring="roc_auc",cv=tscv,verbose=1,n_jobs=-1)
    gsv.fit(X_train,Y_train)

    return gsv
```

30 Informative Feature

```

In [18]: def show_30_informative_feature(vectorizer,model,n=30):

    Weights_Index = model.coef_[0].argsort()
    length = len(Weights_Index)
    # For Negative Class
    neg_class=Weights_Index[:30]

    neg_feat=[vectorizer.get_feature_names()[x] for x in neg_class]
    neg_prob=[model.coef_[0][x] for x in neg_class]

    neg_zip=list(zip(neg_feat,neg_prob))
    neg_zip.sort()

    # For Positive Class
    pos_class = Weights_Index[-30:]

    pos_feat=[vectorizer.get_feature_names()[x] for x in pos_class]
    pos_prob=[model.coef_[0][x] for x in pos_class]
    pos_zip=list(zip(pos_feat,pos_prob))
    pos_zip.sort()

    top=zip(pos_zip,neg_zip)

    print("{0:20}{1:55}{2:20}".format("S.N","Positive","Negative"))
    print("_"*90)
    i=1
    for (fn_1,coef_1), (fn_2,coef_2) in top:
        print("%d.\t\t%.3f\t\t%-30s\t\t%.3f\t\t%s" % (i,coef_1, fn_1, coef_2, fn_2
        ))
        i+=1

```

Bags of Words Vectorizer

```

In [19]: vectorizer = CountVectorizer()
vectorizer.fit(X_train) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_bow = vectorizer.transform(X_train)
X_train_bow=preprocessing.normalize(X_train_bow)

X_test_bow = vectorizer.transform(X_test)
X_test_bow=preprocessing.normalize(X_test_bow)

print("Shape of Train and Test Data After vectorizations")
print(X_train_bow.shape, Y_train.shape)
print(X_test_bow.shape, Y_test.shape)

```

```

Shape of Train and Test Data After vectorizations
(67000, 30737) (67000,)
(33000, 30737) (33000,)

```


Part 1 : Taking L1 as a Regularisation Parameter


Finding the best value Of hyperparameter (Alpha)

```
In [20]: gsv=Grid_SearchCV(X_train_bow,Y_train,'l1')  
  
print("Best HyperParameter: ",gsv.best_params_)  
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l1'}
Best Accuracy: 91.39%

Plot

```
In [21]: plot_l(gsv)
```



Training the model

```
In [23]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_Bow=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,cla
ss_weight="balanced")
Model_Bow.fit(X_train_bow,Y_train)
```

```
Out[23]: SGDClassifier(alpha=0.0001, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l1',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [24]: trainconfusionmatrix(Model_Bow,X_train_bow,Y_train)
```

Confusion Matrix for Train set




```
In [25]: testconfusionmatrix(Model_Bow,X_test_bow,Y_test)
```

Confusion Matrix for Test set



```
In [26]: plot_auc_roc_L(Model_Bow,X_train_bow,X_test_bow,Y_train,Y_test)
```



```
In [27]: print("Classification Report: \n")
y_pred=Model_Bow.predict(X_test_bow)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.43	0.89	0.58	4478
1	0.98	0.81	0.89	28522
micro avg	0.82	0.82	0.82	33000
macro avg	0.70	0.85	0.73	33000
weighted avg	0.91	0.82	0.85	33000

Displaying 30 most informative features

In [28]: `show_30_informative_feature(vectorizer,Model_Bow)`

S.N		Positive		Ne
gative				
1.	6.557	addict	-5.486	ancho
vi				
2.	3.099	alway	-5.576	aw
3.	5.918	amaz	-6.686	bland
4.	4.781	awesom	-4.886	brief
5.	3.891	beat	-7.289	clams
hel				
6.	3.725	beauti	-5.133	conce
pt				
7.	5.042	best	-7.386	credi
t				
8.	3.831	burton	-5.385	dirt
9.	6.243	delici	-6.302	disap
point				
10.	4.048	delight	-8.115	edit
11.	3.165	easi	-8.200	gophe
r				
12.	6.141	excel	-6.973	horch
ata				
13.	4.659	fantast	-10.263	horri
bl				
14.	3.686	fast	-4.066	lack
15.	3.750	favorit	-4.140	leg
16.	4.033	glad	-8.273	poor
17.	4.977	great	-7.422	retur
n				
18.	3.624	happi	-5.484	sept
19.	3.176	hook	-4.682	sorri
20.	3.119	keep	-5.333	stenc
h				
21.	3.467	love	-6.720	sucra
los				
22.	3.931	nice	-7.303	terri
bl				
23.	7.188	perfect	-4.645	threw
24.	3.628	refresh	-4.404	throw
n				
25.	3.335	rich	-5.133	trap
26.	4.295	satisfi	-6.764	twink
i				
27.	4.254	smooth	-5.329	unfor
tun				
28.	4.167	uniqu	-4.145	weak
29.	3.766	wonder	-7.557	wors
30.	5.052	yummi	-7.702	worst

Part 2 : Taking L2 as a Regularisation Parameter

Finding the best value Of hyperparameter (Alpha)


```
In [22]: gsv=Grid_SearchCV(X_train_bow,Y_train,'l2')

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l2'}
Best Accuracy: 92.20%

Plot

```
In [23]: plot_l(gsv)
```



Training the model

```
In [31]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_Bow=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,cla
ss_weight='balanced')
Model_Bow.fit(X_train_bow,Y_train)
```

```
Out[31]: SGDClassifier(alpha=0.0001, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l2',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [32]: trainconfusionmatrix(Model_Bow,X_train_bow,Y_train)
```

Confusion Matrix for Train set




```
In [33]: testconfusionmatrix(Model_Bow,X_test_bow,Y_test)
```

Confusion Matrix for Test set



```
In [34]: plot_auc_roc_L(Model_Bow,X_train_bow,X_test_bow,Y_train,Y_test)
```



```
In [35]: print("Classification Report: \n")
y_pred=Model_Bow.predict(X_test_bow)
print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.45	0.90	0.60	4478
1	0.98	0.83	0.90	28522
micro avg	0.84	0.84	0.84	33000
macro avg	0.72	0.86	0.75	33000
weighted avg	0.91	0.84	0.86	33000

Displaying 30 most informative features

In [36]: `show_30_informative_feature(vectorizer,Model_Bow)`

S.N		Positive		Negative
1.	2.893	addict	-2.919	aw
2.	2.360	alway	-2.134	away
3.	3.197	amaz	-2.315	bad
4.	2.507	awesom	-3.416	bland
5.	1.977	beat	-2.024	didnt
6.	4.094	best	-4.536	disap
point				
7.	4.449	delici	-1.933	disgu
st				
8.	2.361	easi	-2.119	guess
9.	4.225	excel	-2.056	hope
10.	2.416	fantast	-2.413	horri
bl				
11.	2.080	fast	-1.976	howev
12.	2.728	favorit	-2.306	lack
13.	2.383	glad	-2.317	mayb
14.	2.028	good	-2.377	money
15.	4.143	great	-2.103	perha
p				
16.	2.732	happi	-2.069	refun
d				
17.	2.326	keep	-2.037	retur
n				
18.	2.859	love	-1.915	sad
19.	2.835	nice	-2.064	sorri
20.	4.463	perfect	-2.667	stale
21.	2.222	pleasant	-2.100	stick
22.	1.984	quick	-1.929	stuck
23.	2.154	refresh	-2.467	taste
less				
24.	2.323	rich	-3.430	terri
bl				
25.	2.314	satisfi	-2.392	thoug
ht				
26.	2.632	smooth	-2.086	threw
27.	2.330	tasti	-3.094	unfor
tun				
28.	2.244	thank	-2.724	weak
29.	2.746	wonder	-2.064	wors
30.	2.673	yummi	-3.865	worst

TF-IDF Vectorizer

```
In [24]: vectorizer_tfidf=TfidfVectorizer()
vectorizer_tfidf.fit(X_train)
```

```
Out[24]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.float64'>, encoding='utf-8', input='content',
lowercase=True, max_df=1.0, max_features=None, min_df=1,
ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_idf=True,
stop_words=None, strip_accents=None, sublinear_tf=False,
token_pattern='(?u)\\b\\w+\\b', tokenizer=None, use_idf=True,
vocabulary=None)
```

```
In [25]: X_Train_Tfidf=vectorizer_tfidf.transform(X_train)
X_Train_Tfidf=preprocessing.normalize(X_Train_Tfidf)

X_Test_Tfidf=vectorizer_tfidf.transform(X_test)
X_Test_Tfidf=preprocessing.normalize(X_Test_Tfidf)
```

```
In [26]: print("Shape of Train and Test Data After vectorizations")
print(X_Train_Tfidf.shape, Y_train.shape)
print(X_Test_Tfidf.shape, Y_test.shape)
```

```
Shape of Train and Test Data After vectorizations
(67000, 30737) (67000,)
(33000, 30737) (33000,)
```

Taking L1 as a Regularisation Parameter

Finding the best value Of hyperparameter (Alpha)


```
In [27]: gsv=Grid_SearchCV(X_Train_Tfidf,Y_train,'l1')

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

```
Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l1'}
Best Accuracy: 91.45%
```

Plot

In [28]: `plot_l(gsv)`



Training the model

```
In [21]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_Tfidf=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,
class_weight='balanced')
Model_Tfidf.fit(X_Train_Tfidf,Y_train)
```

```
Out[21]: SGDClassifier(alpha=0.0001, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l1',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [22]: trainconfusionmatrix(Model_Tfidf,X_Train_Tfidf,Y_train)
```

Confusion Matrix for Train set




```
In [23]: testconfusionmatrix(Model_Tfidf,X_Test_Tfidf,Y_test)
```

Confusion Matrix for Test set



```
In [24]: plot_auc_roc_L(Model_Tfidf,X_Train_Tfidf,X_Test_Tfidf,Y_train,Y_test)
```



```
In [25]: print("Classification Report: \n")
y_pred=Model_Tfidf.predict(X_Test_Tfidf)
print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.44	0.89	0.59	4478
1	0.98	0.83	0.90	28522
micro avg	0.83	0.83	0.83	33000
macro avg	0.71	0.86	0.74	33000
weighted avg	0.91	0.83	0.85	33000

Displaying 30 most informative features

In [28]: `show_30_informative_feature(vectorizer_tfidf,Model_Tfidf)`

S.N	Positive		Negative	
1.	6.218	addict	-5.533	alvin
2.	3.493	alway	-5.740	billi
3.	6.243	amaz	-5.652	bland
4.	4.914	awesom	-8.018	brief
5.	3.631	beauti	-12.165	categ
6.	8.510	best	-7.380	cinem
7.	5.288	burton	-15.401	clams
8.	8.874	delici	-8.837	credi
9.	4.007	delight	-6.060	debbi
10.	7.494	excel	-6.335	defic
11.	4.363	fantast	-7.808	dirt
12.	5.369	favorit	-6.838	disap
13.	3.787	find	-7.074	edit
14.	4.131	glad	-8.417	gazil
15.	4.704	good	-6.761	ghoul
16.	9.627	great	-13.176	gophe
17.	4.196	happi	-8.771	horch
18.	3.498	hotter	-8.035	horri
19.	3.977	keep	-6.861	leg
20.	6.823	love	-6.551	misle
21.	5.009	nice	-5.475	scrip
22.	3.390	often	-11.039	sept
23.	9.882	perfect	-8.171	spamm
24.	4.544	satisfi	-7.132	stenc
25.	4.150	smooth	-5.628	sucra
26.	3.599	tasti	-5.817	terri
27.	3.788	tim	-6.206	titan
28.	4.407	uniqu	-5.556	unfor
29.	4.719	wonder	-7.890	wors
30.	5.013	yummi	-6.649	worst

Taking L2 as a Regularisation Parameter


Finding the best value Of hyperparameter (Alpha)

```
In [29]: gsv=Grid_SearchCV(X_Train_Tfidf,Y_train,'l2')  
  
print("Best HyperParameter: ",gsv.best_params_)  
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'alpha': 0.0001, 'penalty': 'l2'}
Best Accuracy: 92.57%

Plot

```
In [30]: plot_l(gsv)
```



Training the mode

```
In [48]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_Tfidf=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,class_weight='balanced')
Model_Tfidf.fit(X_Train_Tfidf,Y_train)
```

```
Out[48]: SGDClassifier(alpha=0.0001, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l2',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [49]: trainconfusionmatrix(Model_Tfidf,X_Train_Tfidf,Y_train)
```

Confusion Matrix for Train set





```
In [50]: testconfusionmatrix(Model_Tfidf,X_Test_Tfidf,Y_test)
```

Confusion Matrix for Test set



```
In [51]: plot_auc_roc_L(Model_Tfidf,X_Train_Tfidf,X_Test_Tfidf,Y_train,Y_test)
```



```
In [52]: print("Classification Report: \n")
y_pred=Model_Tfidf.predict(X_Test_Tfidf)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.47	0.90	0.62	4478
1	0.98	0.84	0.91	28522
micro avg	0.85	0.85	0.85	33000
macro avg	0.73	0.87	0.76	33000
weighted avg	0.91	0.85	0.87	33000

Displaying 30 most informative features

```
In [53]: show_30_informative_feature(vectorizer_tfidf,Model_Tfidf)
```

S.N		Positive		Negative
1.	2.749	addict	-2.662	aw
2.	2.614	alway	-1.930	away
3.	2.992	amaz	-2.193	bad
4.	2.450	awesom	-3.264	bland
5.	2.087	beat	-2.160	didnt
6.	5.369	best	-4.585	disap
point				
7.	4.792	delici	-1.957	disgu
st				
8.	2.346	easi	-2.081	gross
9.	4.339	excel	-2.012	guess
10.	2.252	fantast	-2.115	hope
11.	2.046	fast	-2.416	horri
bl				
12.	3.172	favorit	-2.484	lack
13.	2.483	find	-2.234	mayb
14.	2.164	glad	-2.366	money
15.	3.430	good	-1.959	perha
p				
16.	6.535	great	-1.970	retur
n				
17.	2.700	happi	-2.098	sorri
18.	2.560	keep	-2.445	stale
19.	4.743	love	-1.997	stick
20.	3.208	nice	-2.004	stuck
21.	4.640	perfect	-2.033	tast
22.	2.132	quick	-2.474	taste
less				
23.	2.132	refresh	-3.163	terri
bl				
24.	2.070	rich	-2.535	thoug
ht				
25.	2.201	satisfi	-2.050	threw
26.	2.554	smooth	-2.916	unfor
tun				
27.	2.429	tasti	-2.585	weak
28.	2.123	thank	-2.208	wors
29.	3.127	wonder	-3.548	worst
30.	2.548	yummi	-1.936	yuck

Word To Vector

```
In [31]: list_of_Train_sent=[]
list_of_Test_sent=[]

for sent in X_train:
    list_of_Train_sent.append(sent.split())

for sent in X_test:
    list_of_Test_sent.append(sent.split())
```

```
In [32]: model=Word2Vec(list_of_Train_sent,min_count=5,size=50, workers=4)
```

Average Word To Vector

```
In [33]: import numpy as np

Train_vectors = []
for sent in list_of_Train_sent:
    sent_vec = np.zeros(50)
    cnt_words = 0
    for word in sent:
        try:
            vec = model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    if cnt_words!=0:
        sent_vec /= cnt_words
    Train_vectors.append(sent_vec)
Train_vectors = np.nan_to_num(Train_vectors)
```

```
In [34]: import numpy as np

Test_vectors=[]
for sent in list_of_Test_sent:
    sent_vec=np.zeros(50)
    cnt_words=0
    for word in sent:
        try:
            vec=model.wv[word]
            sent_vec+=vec
            cnt_words+=1
        except:
            pass
    if cnt_words!=0:
        sent_vec/=cnt_words
    Test_vectors.append(sent_vec)
Test_vectors=np.nan_to_num(Test_vectors)
```

```
In [35]: print("Shape of Test Vectors : ",Test_vectors.shape)
```

Shape of Test Vectors : (33000, 50)

```
In [36]: X_Train_Awv=Train_vectors  
X_Test_Awv=Test_vectors
```

```
In [37]: print(X_Train_Awv.shape, Y_train.shape)  
print(X_Test_Awv.shape, Y_test.shape)
```

(67000, 50) (67000,)
(33000, 50) (33000,)

Taking L1 as a Regularisation Parameter


Finding the best value Of hyperparameter (Alpha)

```
In [38]: gsv=Grid_SearchCV(X_Train_Awv,Y_train,"l1")  
  
print("Best HyperParameter: ",gsv.best_params_)  
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'alpha': 0.001, 'penalty': 'l1'}
Best Accuracy: 88.50%

Plot

```
In [39]: plot_l(gsv)
```



Training Model

```
In [62]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_Awv=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,cla
ss_weight='balanced')
Model_Awv.fit(X_Train_Awv,Y_train)
```

```
Out[62]: SGDClassifier(alpha=0.001, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l1',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [63]: trainconfusionmatrix(Model_Awv,X_Train_Awv,Y_train)
```

Confusion Matrix for Train set




```
In [64]: trainconfusionmatrix(Model_Awv,X_Test_Awv,Y_test)
```

Confusion Matrix for Train set



```
In [65]: plot_auc_roc_L(Model_Awv,X_Train_Awv,X_Test_Awv,Y_train,Y_test)
```



```
In [66]: print("Classification Report: \n")
y_pred=Model_Awv.predict(X_Test_Awv)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.34	0.89	0.49	4478
1	0.98	0.73	0.83	28522
micro avg	0.75	0.75	0.75	33000
macro avg	0.66	0.81	0.66	33000
weighted avg	0.89	0.75	0.79	33000

Taking L2 as a Regularisation Parameter

Finding the best value Of hyperparameter (Alpha)

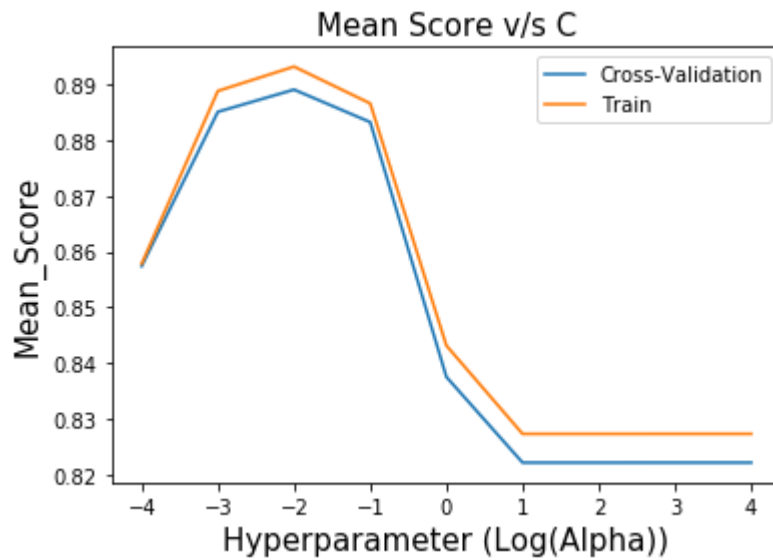
```
In [40]: gsv=Grid_SearchCV(X_Train_Awv,Y_train,"l2")

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'alpha': 0.01, 'penalty': 'l2'}
 Best Accuracy: 88.90%

Plot

In [41]: `plot_l(gsv)`



Training the model

```
In [68]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_Awv=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,cla
ss_weight='balanced')
Model_Awv.fit(X_Train_Awv,Y_train)
```

```
Out[68]: SGDClassifier(alpha=0.01, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l2',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [69]: trainconfusionmatrix(Model_Awv,X_Train_Awv,Y_train)
```

Confusion Matrix for Train set




```
In [70]: testconfusionmatrix(Model_Awv,X_Test_Awv,Y_test)
```

Confusion Matrix for Test set



```
In [71]: plot_auc_roc_L(Model_Awv,X_Train_Awv,X_Test_Awv,Y_train,Y_test)
```



```
In [72]: print("Classification Report: \n")
y_pred=Model_Awv.predict(X_Test_Awv)
print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.36	0.88	0.51	4478
1	0.97	0.75	0.85	28522
micro avg	0.77	0.77	0.77	33000
macro avg	0.67	0.81	0.68	33000
weighted avg	0.89	0.77	0.80	33000

TF-IDF Word To Vector

```
In [42]: TFIDF_Feature=vectorizer_tfidf.get_feature_names()
print(len(TFIDF_Feature))
print(TFIDF_Feature[0:20])
```

30737

```
['aaa', 'aaaaaaaaagghh', 'aaaaah', 'aaaaahhhhhhhhhhhhhhhhh', 'aaaah', 'aaah',
'aachen', 'aad', 'aadp', 'aafco', 'aagh', 'aah', 'aahh', 'aand', 'aardvark',
'ab', 'aback', 'abandon', 'abaolut', 'abattoir']
```

```
In [43]: from tqdm import tqdm
Train_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Train_sent):
    sent_vec=np.zeros(50)
    weight=0
    for word in sent:
        try :
            w2v_vec=model.wv[word]
            tfidf_vec=X_Train_Tfidf[row,TFIDF_Feature.index(word)]
            sent_vec+=(w2v_vec*tfidf_vec)
            weight+=tfidf_vec

        except :
            pass
    if weight!=0:
        sent_vec/=weight
    Train_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1
```

[illegible]

```
In [44]: Test_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Test_sent):
    sent_vec=np.zeros(50)
    weight=0

    for word in sent:
        try:
            w2v_vec=model.wv[word]
            tfidf_vec=X_Test_Tfidf(row,TFIDF_Feature.index(word))
            sent_vec+=(w2v_vec*tfidf_vec)
            weight+=tfidf

        except :
            pass

    if weight!=0:
        sent_vec/=weight
    Test_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1
```

```
100%|███████████████████████████████████| 33000/33000 [11:02<00:00, 49.84it/  
s]
```

```
In [45]: Train_TFIDF_W2V_Vectors = np.nan_to_num(Train_TFIDF_W2V_Vectors)
Test_TFIDF_W2V_Vectors = np.nan_to_num(Test_TFIDF_W2V_Vectors)
```

```
In [46]: X_Train_TfidfW2v=Train_TFIDF_W2V_Vectors
X_Test_TfidfW2v=Test_TFIDF_W2V_Vectors
```

```
In [47]: print(X_Train_TfidfW2v.shape, Y_train.shape)
print(X_Test_TfidfW2v.shape, Y_test.shape)
```

```
(67000, 50) (67000,)
(33000, 50) (33000,)
```

Taking L1 as a Regularisation Parameter

Finding the best value Of hyperparameter (Alpha)


```
In [48]: gsv=Grid_SearchCV(X_Train_TfidfW2v,Y_train,"l1")

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

```
Best HyperParameter: {'alpha': 0.01, 'penalty': 'l1'}
Best Accuracy: 85.66%
```

Plot

```
In [49]: plot_l(gsv)
```



Training Model

```
In [47]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_TfidfW2v=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,
class_weight='balanced')
Model_TfidfW2v.fit(X_Train_TfidfW2v,Y_train)
```

```
Out[47]: SGDClassifier(alpha=0.001, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l1',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [48]: trainconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Train set




```
In [49]: testconfusionmatrix(Model_TfidfW2v,X_Test_TfidfW2v,Y_test)
```

Confusion Matrix for Test set



```
In [86]: plot_auc_roc_L(Model_TfidfW2v,X_Train_TfidfW2v,X_Test_TfidfW2v,Y_train,Y_test)
```



```
In [87]: print("Classification Report: \n")
y_pred=Model_TfidfW2v.predict(X_Test_TfidfW2v)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.14	1.00	0.24	4478
1	0.00	0.00	0.00	28522
micro avg	0.14	0.14	0.14	33000
macro avg	0.07	0.50	0.12	33000
weighted avg	0.02	0.14	0.03	33000

Taking L2 as a Regularisation Parameter

Finding the best value Of hyperparameter (Alpha)


```
In [50]: gsv=Grid_SearchCV(X_Train_TfidfW2v,Y_train,"l1")

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'alpha': 0.01, 'penalty': 'l1'}
 Best Accuracy: 85.66%

Plot

In [51]: `plot_1(gsv)`



Training the model

```
In [89]: Best_Param=gsv.best_params_
C=Best_Param['alpha']
Penalty = Best_Param['penalty']


Model_TfidfW2v=SGDClassifier(alpha=C,penalty=Penalty,loss='hinge',shuffle=False,
class_weight='balanced')
Model_TfidfW2v.fit(X_Train_TfidfW2v,Y_train)
```

```
Out[89]: SGDClassifier(alpha=0.001, average=False, class_weight='balanced',
early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=None,
n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l1',
power_t=0.5, random_state=None, shuffle=False, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Evaluating the performance of model


```
In [90]: trainconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Train set




```
In [91]: testconfusionmatrix(Model_TfidfW2v,X_Test_TfidfW2v,Y_train)
```

Confusion Matrix for Test set



```
In [92]: plot_auc_roc_L(Model_TfidfW2v,X_Train_TfidfW2v,X_Test_TfidfW2v,Y_train,Y_test)
```



```
In [93]: print("Classification Report: \n")
y_pred=Model_TfidfW2v.predict(X_Test_TfidfW2v)
print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.14	1.00	0.24	4478
1	0.00	0.00	0.00	28522
micro avg	0.14	0.14	0.14	33000
macro avg	0.07	0.50	0.12	33000
weighted avg	0.02	0.14	0.03	33000

Using RBF Kernel

```
In [52]: Data_RBF = Data.head(20000)
```

```
In [53]: Y = Data_RBF['Score']
X = Data_RBF['CleanedText']
```

```
In [54]: from sklearn.model_selection import train_test_split
# X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, shuffle=False): this is for time series split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33, shuffle=False) # this is random splitting

print("Shape of Train and Test Dataset for 100k points")
print(X_train.shape, Y_train.shape)
print(X_test.shape, Y_test.shape)
```

```
Shape of Train and Test Dataset for 100k points
(13400,) (13400,)
(6600,) (6600,)
```

Bag Of Words

```
In [55]: vectorizer = CountVectorizer(max_features=500)
vectorizer.fit(X_train) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_bow = vectorizer.transform(X_train)
X_train_bow=preprocessing.normalize(X_train_bow)

X_test_bow = vectorizer.transform(X_test)
X_test_bow=preprocessing.normalize(X_test_bow)

print("Shape of Train and Test Data After vectorizations")
print(X_train_bow.shape, Y_train.shape)
print(X_test_bow.shape, Y_test.shape)
```

```
Shape of Train and Test Data After vectorizations
(13400, 500) (13400,)
(6600, 500) (6600,)
```

Finding the best value of hyperparameter Alpha

```
In [56]: gsv=Grid_SearchCV_RBF(X_train_bow,Y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 5 folds for each of 9 candidates, totalling 45 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 45 out of 45 | elapsed: 9.2min finished
```

```
Best HyperParameter: {'C': 100}
Best Accuracy: 87.13%
```

Plot

```
In [57]: plot_r(gsv)
```



Training the model


```
In [100]: Best_Param=gsv.best_params_  
C = Best_Param['C']  
  
Model_Bow=SVC(C=C,probability=True,class_weight='balanced')  
Model_Bow.fit(X_train_bow,Y_train)
```

```
Out[100]: SVC(C=100, cache_size=200, class_weight='balanced', coef0=0.0,  
decision_function_shape='ovr', degree=3, gamma='auto_deprecated',  
kernel='rbf', max_iter=-1, probability=True, random_state=None,  
shrinking=True, tol=0.001, verbose=False)
```

Evaluating the performance of model


```
In [101]: trainconfusionmatrix(Model_Bow,X_train_bow,Y_train)
```

Confusion Matrix for Train set




```
In [102]: testconfusionmatrix(Model_Bow,X_test_bow,Y_test)
```

Confusion Matrix for Test set



```
In [103]: plot_auc_roc(Model_Bow,X_train_bow,X_test_bow,Y_train,Y_test)
```



```
In [104]: print("Classification Report: \n")
y_pred=Model_Bow.predict(X_test_bow)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.35	0.80	0.48	705
1	0.97	0.82	0.89	5895
micro avg	0.82	0.82	0.82	6600
macro avg	0.66	0.81	0.69	6600
weighted avg	0.90	0.82	0.85	6600

TF-IDF Vectorizer

```
In [58]: vectorizer_tfidf=TfidfVectorizer()
vectorizer_tfidf.fit(X_train)
```

```
Out[58]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
dtype=<class 'numpy.float64'>, encoding='utf-8', input='content',
lowercase=True, max_df=1.0, max_features=None, min_df=1,
ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_idf=True,
stop_words=None, strip_accents=None, sublinear_tf=False,
token_pattern='(?u)\\b\\w\\w+\\b', tokenizer=None, use_idf=True,
vocabulary=None)
```

```
In [59]: X_Train_Tfidf=vectorizer_tfidf.transform(X_train)
X_Train_Tfidf=preprocessing.normalize(X_Train_Tfidf)

X_Test_Tfidf=vectorizer_tfidf.transform(X_test)
X_Test_Tfidf=preprocessing.normalize(X_Test_Tfidf)
```

```
In [60]: print("Shape of Train and Test Data After vectorizations")
print(X_Train_Tfidf.shape, Y_train.shape)
print(X_Test_Tfidf.shape, Y_test.shape)
```

```
Shape of Train and Test Data After vectorizations
(13400, 15698) (13400,)
(6600, 15698) (6600,)
```

Finding the best value of hyperparameter Alpha

```
In [61]: gsv=Grid_SearchCV_RBF(X_Train_Tfidf,Y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 5 folds for each of 9 candidates, totalling 45 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 45 out of 45 | elapsed: 14.2min finished
```

```
Best HyperParameter: {'C': 1000}
Best Accuracy: 90.33%
```

Plot

```
In [62]: plot_r(gsv)
```



Training the model


```
In [109]: Best_Param=gsv.best_params_  
C = Best_Param['C']  
  
Model_Tfidf=SVC(C=C,probability=True,class_weight='balanced')  
Model_Tfidf.fit(X_Train_Tfidf,Y_train)
```

```
Out[109]: SVC(C=1000, cache_size=200, class_weight='balanced', coef0=0.0,  
decision_function_shape='ovr', degree=3, gamma='auto_deprecated',  
kernel='rbf', max_iter=-1, probability=True, random_state=None,  
shrinking=True, tol=0.001, verbose=False)
```

Evaluating the performance of model


```
In [112]: trainconfusionmatrix(Model_Tfidf,X_Train_Tfidf,Y_train)
```

Confusion Matrix for Train set




```
In [113]: testconfusionmatrix(Model_Tfidf,X_Test_Tfidf,Y_test)
```

Confusion Matrix for Test set



```
In [114]: plot_auc_roc(Model_Tfidf,X_Train_Tfidf,X_Test_Tfidf,Y_train,Y_test)
```



```
In [115]: print("Classification Report: \n")
y_pred=Model_Tfidf.predict(X_Test_Tfidf)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.43	0.80	0.56	705
1	0.97	0.87	0.92	5895
micro avg	0.86	0.86	0.86	6600
macro avg	0.70	0.84	0.74	6600
weighted avg	0.92	0.86	0.88	6600

Word To Vector

```
In [63]: list_of_Train_sent=[]
list_of_Test_sent=[]

for sent in X_train:
    list_of_Train_sent.append(sent.split())

for sent in X_test:
    list_of_Test_sent.append(sent.split())
```

```
In [64]: model=Word2Vec(list_of_Train_sent,min_count=5,size=50, workers=4)
```

```
In [65]: import numpy as np

Train_vectors = []
for sent in list_of_Train_sent:
    sent_vec = np.zeros(50)
    cnt_words = 0
    for word in sent:
        try:
            vec = model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    if cnt_words!=0:
        sent_vec /= cnt_words
    Train_vectors.append(sent_vec)
Train_vectors = np.nan_to_num(Train_vectors)
```

```
In [66]: import numpy as np

Test_vectors=[]
for sent in list_of_Test_sent:
    sent_vec=np.zeros(50)
    cnt_words=0
    for word in sent:
        try:
            vec=model.wv[word]
            sent_vec+=vec
            cnt_words+=1
        except:
            pass
    if cnt_words!=0:
        sent_vec/=cnt_words
    Test_vectors.append(sent_vec)
Test_vectors=np.nan_to_num(Test_vectors)
```

```
In [67]: print("Shape of Test Vectors : ",Test_vectors.shape)
```

Shape of Test Vectors : (6600, 50)

```
In [68]: X_Train_Awv=Train_vectors
X_Test_Awv=Test_vectors
```

```
In [69]: print(X_Train_Awv.shape, Y_train.shape)
print(X_Test_Awv.shape, Y_test.shape)
```

(13400, 50) (13400,)
(6600, 50) (6600,)

Finding the best hyperparameter (C)

```
In [70]: gsv=Grid_SearchCV_RBF(X_Train_Awv,Y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 5 folds for each of 9 candidates, totalling 45 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 45 out of 45 | elapsed: 5.5min finished

Best HyperParameter: {'C': 10}
Best Accuracy: 86.11%

Plot

```
In [71]: plot_r(gsv)
```



Training the model

```
In [124]: Best_Param=gsv.best_params_
C = Best_Param['C']


Model_Awv=SVC(C=C,probability=True,class_weight='balanced')
Model_Awv.fit(X_Train_Awv,Y_train)
```

```
Out[124]: SVC(C=10, cache_size=200, class_weight='balanced', coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
kernel='rbf', max_iter=-1, probability=True, random_state=None,
shrinking=True, tol=0.001, verbose=False)
```

Evaluating the performance of model


```
In [125]: trainconfusionmatrix(Model_Awv,X_Train_Awv,Y_train)
```

Confusion Matrix for Train set




```
In [126]: testconfusionmatrix(Model_Awv,X_Test_Awv,Y_test)
```

Confusion Matrix for Test set



```
In [127]: plot_auc_roc(Model_Awv,X_Train_Awv,X_Test_Awv,Y_train,Y_test)
```



```
In [128]: print("Classification Report: \n")
y_pred=Model_Awv.predict(X_Test_Awv)
print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.29	0.78	0.43	705
1	0.97	0.77	0.86	5895
micro avg	0.78	0.78	0.78	6600
macro avg	0.63	0.78	0.64	6600
weighted avg	0.90	0.78	0.81	6600

Tf-IDF Word To Vector

```
In [72]: TFIDF_Feature=vectorizer_tfidf.get_feature_names()
print(len(TFIDF_Feature))
print(TFIDF_Feature[0:20])
```

15698

['aaaaah', 'aafco', 'aagh', 'aah', 'ab', 'aback', 'abandon', 'abba', 'abc', 'abdomen', 'abdomin', 'abhor', 'abid', 'abil', 'abj', 'abl', 'abliti', 'abnorm', 'aboard', 'abod']

```
In [73]: from tqdm import tqdm
Train_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Train_sent):
    sent_vec=np.zeros(50)
    weight=0
    for word in sent:
        try :
            w2v_vec=model.wv[word]
            tfidf_vec=X_Train_Tfidf[row,TFIDF_Feature.index(word)]
            sent_vec+=(w2v_vec*tfidf_vec)
            weight+=tfidf_vec

        except :
            pass
    if weight!=0:
        sent_vec/=weight
    Train_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1
```

```
100% |██████████████████████████████| 13400/13400 [02:49<00:00, 79.23it/  
s]
```

```
In [74]: Test_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Test_sent):
    sent_vec=np.zeros(50)
    weight=0

    for word in sent:
        try:
            w2v_vec=model.wv[word]
            tfidf_vec=X_Test_Tfidf(row,TFIDF_Feature.index(word))
            sent_vec+=(w2v_vec*tfidf_vec)
            weight+=tfidf

        except :
            pass

    if weight!=0:
        sent_vec/=weight
    Test_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1
```

```
100%|███████████████████████████████| 6600/6600 [00:56<00:00, 117.44it/  
s]
```

```
In [75]: Train_TFIDF_W2V_Vectors = np.nan_to_num(Train_TFIDF_W2V_Vectors)
Test_TFIDF_W2V_Vectors = np.nan_to_num(Test_TFIDF_W2V_Vectors)
```

```
In [76]: X_Train_TfidfW2v=Train_TFIDF_W2V_Vectors
X_Test_TfidfW2v=Test_TFIDF_W2V_Vectors
```



```
In [77]: print(X_Train_TfidfW2v.shape, Y_train.shape)
print(X_Test_TfidfW2v.shape, Y_test.shape)
```

```
(13400, 50) (13400,)
(6600, 50) (6600,)
```

Finding the best hyperparameter (C)

```
In [78]: gsv=Grid_SearchCV_RBF(X_Train_TfidfW2v,Y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 5 folds for each of 9 candidates, totalling 45 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
 [Parallel(n_jobs=-1)]: Done 45 out of 45 | elapsed: 5.8min finished

Best HyperParameter: {'C': 10}
 Best Accuracy: 83.63%

Plot

```
In [79]: plot_r(gsv)
```



Training the model

```
In [136]: Best_Param=gsv.best_params_
C = Best_Param['C']


Model_TfidfW2v=SVC(C=C,probability=True,class_weight='balanced')
Model_TfidfW2v.fit(X_Train_TfidfW2v,Y_train)
```

```
Out[136]: SVC(C=10, cache_size=200, class_weight='balanced', coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto_deprecated',
kernel='rbf', max_iter=-1, probability=True, random_state=None,
shrinking=True, tol=0.001, verbose=False)
```

Evaluating the performance of model


```
In [137]: trainconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Train set



```
In [ ]: testconfusionmatrix(Model_TfidfW2v,X_Test_TfidfW2v,Y_test)
```

```
In [139]: plot_auc_roc(Model_TfidfW2v,X_Train_TfidfW2v,X_Test_TfidfW2v,Y_train,Y_test)
```



```
In [140]: print("Classification Report: \n")
y_pred=Model_TfidfW2v.predict(X_Test_TfidfW2v)
print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.11	1.00	0.19	705
1	0.00	0.00	0.00	5895
micro avg	0.11	0.11	0.11	6600
macro avg	0.05	0.50	0.10	6600
weighted avg	0.01	0.11	0.02	6600

1.Report On Different Vectorizer Method and RBF Kernel

```
In [75]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Hyperparameter(C)", "Train AUC", "Test AUC", "F1-Score"]

x.add_row(["BOW", 100, 0.90, 0.87, 0.86])
x.add_row(["TF-IDF", 1000, 0.97, 0.91, 0.94])
x.add_row(["Avg W2V", 100, 0.88, 0.84, 0.85])
x.add_row(["TF-IDF W2V", 1000, 0.90, 0.50, 0.84])

print(x)
```

Vectorizer	Hyperparameter(C)	Train AUC	Test AUC	F1-Score
BOW	100	0.9	0.87	0.86
TF-IDF	1000	0.97	0.91	0.94
Avg W2V	100	0.88	0.84	0.85
TF-IDF W2V	1000	0.9	0.5	0.84

2.Report On Different Vectorizer Method and Linear Kernel (L1 Regularisation Parameter)

```
In [50]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Hyperparameter(Alpha)", "Train AUC", "Test AUC", "F1-Score"]

x.add_row(["BOW", 0.0001, 0.93, 0.92, 0.87])
x.add_row(["TF-IDF", 0.0001, 0.93, 0.93, 0.87])
x.add_row(["Avg W2V", 0.001, 0.89, 0.89, 0.83])
x.add_row(["TF-IDF W2V", 0.001, 0.86, 0.50, 0.80])

print(x)
```

Vectorizer	Hyperparameter(Alpha)	Train AUC	Test AUC	F1-Score
BOW	0.0001	0.93	0.92	0.87
TF-IDF	0.0001	0.93	0.93	0.87
Avg W2V	0.001	0.89	0.89	0.83
TF-IDF W2V	0.001	0.86	0.5	0.8

3.Report On Different Vectorizer Method and Linear Kernel (L2 Regularisation Parameter)

```
In [51]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Hyperparameter(Alpha)", "Train AUC", "Test AUC",
"F1-Score"]

x.add_row(["BOW", 0.0001, 0.94, 0.93, 0.88])
x.add_row(["TF-IDF", 0.0001, 0.96, 0.94, 0.88])
x.add_row(["Avg W2V", 0.001, 0.89, 0.89, 0.80])
x.add_row(["TF-IDF W2V", 0.001, 0.86, 0.50, 0.80])

print(x)
```

Vectorizer	Hyperparameter(Alpha)	Train AUC	Test AUC	F1-Score
BOW	0.0001	0.94	0.93	0.88
TF-IDF	0.0001	0.96	0.94	0.88
Avg W2V	0.001	0.89	0.89	0.8
TF-IDF W2V	0.001	0.86	0.5	0.8

4. I have used SGDClassifier for Linear SVM on 100K DataSet and SVC for RBF SVM on 20K DataSet.

5. Since data is unbalanced , i did time based splitting and used roc_auc metric as scoring parameter in GridsearchCV .

6. In case of RBF SVM , TFIDF is performing better than other.

7. In Case of Linear SVM , TFIDF-W2V is overfitting.