

# Assignment-9: Apply Random Forest On Amazon Fine Food Reviews DataSet ¶

## Introduction

(i).A random forest consists of multiple random decision trees. Two types of randomnesses are built into the trees. First, each tree is built on a random sample from the original data. Second, at each tree node, a subset of features are randomly selected to generate the best split.

## Objective

To Predict the Polarity of Amazon Fine Food Review Using Random Forest Algorithm.

## Importing All Required Library

```

In [1]: %matplotlib inline
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
import math

from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.ensemble import RandomForestClassifier
from sklearn.datasets import make_classification

from sklearn.metrics import classification_report
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import TimeSeriesSplit

from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import cross_val_score
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score

from gensim.models import Word2Vec
from gensim.models import KeyedVectors

import pickle

from tqdm import tqdm
import os
import warnings
warnings.filterwarnings("ignore")

```

C:\Users\User\Anaconda3\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows; aliasing chunkize to chunkize\_serial  
 warnings.warn("detected Windows; aliasing chunkize to chunkize\_serial")

## Importing Amazon Fine Food Review Dataset

```

In [2]: if os.path.isfile("final.sqlite"):
        conn=sqlite3.connect("final.sqlite")
        Data=pd.read_sql_query("select * from Reviews where Score!=3",conn)
        conn.close()
    else :
        print("Error Importing the file")

```

In [3]: *# Printing some data of DataFrame*

```
Data['Score'].value_counts()
```

Out[3]:

1	307061
0	57110

Name: Score, dtype: int64

## Information About DataSet

```
In [4]: print("\nNumber of Reviews: ",Data["Text"].count())
print("\nNumber of Users: ",len(Data["UserId"].unique())) # Unique returns 1-D
array of unique values of DataFrame.
print("\nNumber of Products: ",len(Data["ProductId"].unique()))
print("\nShape of Data: ", Data.shape)
print("\nColumn Name of DataSet : ",Data.columns)
print("\n\nNumber of Attributes/Columns in data: 12")
print("\nNumber of Positive Reviews : ", Data['Score'].value_counts()[1])
print("\nNumber of Negative Reviews : ", Data['Score'].value_counts()[0])
```

Number of Reviews: 364171

Number of Users: 243414

Number of Products: 65442

Shape of Data: (364171, 12)

Column Name of DataSet : Index(['index', 'Id', 'ProductId', 'UserId', 'ProfileName',  
'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',  
'Summary', 'Text', 'CleanedText'],  
dtype='object')

Number of Attributes/Columns in data: 12

Number of Positive Reviews : 307061

Number of Negative Reviews : 57110

```
In [5]: print("\nNumber of Reviews: ",Data["Text"].count())
```

Number of Reviews: 364171

## Attribute Information About DataSet

- 1.Id - A unique value starts from 1
- 2.ProductId - A unique identifier for the product
- 3.UserId - A unique identifier for the user
- 4.ProfileName - Name of user profile
- 5.HelpfulnessNumerator - Number of users who found the review helpful
- 6.HelpfulnessDenominator - Number of users who indicated whether they found the review helpful or not
- 7.Score - Rating 0 or 1
- 8.Time - Timestamp for the review
- 9.Summary - Brief summary of the review
- 10.Text - Text of the review
- 11.Cleaned Text - Text that only alphabets

```
In [6]: # Sorting on the basis of Time Parameter
Data.sort_values('Time', inplace=True)
```

```
In [7]: Data=Data.head(100000)
```

```
In [8]: Y = Data['Score']
X = Data['CleanedText']
```

## Splitting DataSet into Train and Test Data

```
In [9]: from sklearn.model_selection import train_test_split
# X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, shuffle=False): this is for time series split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33, shuffle=False) # this is random splitting

print("Shape of Train and Test Dataset for 100k points")
print(X_train.shape, Y_train.shape)
print(X_test.shape, Y_test.shape)
```

```
Shape of Train and Test Dataset for 100k points
(67000,) (67000,)
(33000,) (33000,)
```

## Defining Some Function

## Train Data Confusion Matrix Plot

```
In [10]: def trainconfusionmatrix(model,X_train,y_train):
    print("Confusion Matrix for Train set")
    cm=confusion_matrix(y_train, model.predict(X_train))
    class_label = ["negative", "positive"]
    df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
    sns.heatmap(df_cm, annot = True, fmt = "d")
    plt.title("Train Confusiion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()
```

## Test Data Confusion Matrix Plot

```
In [11]: def testconfusionmatrix(model,X_test,y_test):
    print("Confusion Matrix for Test set")
    cm=confusion_matrix(y_test, model.predict(X_test))
    class_label = ["negative", "positive"]
    df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
    sns.heatmap(df_cm, annot = True, fmt = "d")
    plt.title("Test Confusiion Matrix")
    plt.xlabel("Predicted Label")
    plt.ylabel("True Label")
    plt.show()
```

## ROC-AUC Curve

```
In [12]: def plot_auc_roc(model,X_train,X_test,y_train,y_test):
    train_fpr, train_tpr, thresholds = roc_curve(y_train,model.predict_proba(X_train)[:,:1])
    test_fpr, test_tpr, thresholds = roc_curve(y_test,model.predict_proba(X_test)[:,:1])

    plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_tpr)))
    plt.plot(test_fpr, test_tpr, label="test AUC =" +str(auc(test_fpr, test_tpr)))
    plt.legend()
    plt.xlabel("Hyperameter (C)")
    plt.ylabel("AUC")
    plt.title("ROC CURVE PLOTS")
    plt.show()
```

## GridSearchCV

```
In [23]: def Grid_SearchCV(X_train,Y_train):

    param_grid= {'max_depth':[1,5,10,50,100,500,1000],
                  'n_estimators' :[5,10,20,50,100]
                }
    tscv = TimeSeriesSplit(n_splits=10)
    clf = RandomForestClassifier(class_weight='balanced')
    gsv=GridSearchCV(clf,param_grid, scoring = 'roc_auc', cv=tscv,n_jobs=-1,pre_dispatch=2)
    gsv.fit(X_train,Y_train)

    return gsv
```

## Train Auc Heatmap

```
In [26]: def Train_Auc_Heatmap(gsv):
    results = gsv.cv_results_
    train = np.array(results['mean_train_score'])
    train = train.reshape(7,5)
    sns.heatmap(train,annot =True,yticklabels = [1,5,10,50,100,500,1000],xticklabels = [5,10,20,50,100],linewidths=.5,cmap="YlGnBu" )
```

## Test Auc Heatmap

```
In [27]: def Cv_Auc_Heatmap(gsv):
    results = gsv.cv_results_
    cv = np.array(results['mean_test_score'])
    cv = cv.reshape(7,5)
    sns.heatmap(cv,annot =True,yticklabels = [1,5,10,50,100,500,1000],xticklabels = [5,10,20,50,100],linewidths=.5 ,cmap= 'YlGnBu' )
```

## Bags of Words Vectorizer

```
In [17]: vectorizer = CountVectorizer()
vectorizer.fit(X_train) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_bow = vectorizer.transform(X_train)
X_train_bow=preprocessing.normalize(X_train_bow)

X_test_bow = vectorizer.transform(X_test)
X_test_bow=preprocessing.normalize(X_test_bow)

print("Shape of Train and Test Data After vectorizations")
print(X_train_bow.shape, Y_train.shape)
print(X_test_bow.shape, Y_test.shape)
```

Shape of Train and Test Data After vectorizations  
 (67000, 30737) (67000,)  
 (33000, 30737) (33000,)

## Finding the best value Of hyperparameter


```
In [24]: gsv=Grid_SearchCV(X_train_bow,Y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Best HyperParameter: {'max\_depth': 500, 'n\_estimators': 100}  
 Best Accuracy: 90.96%


## Train Auc HeatMap

```
In [28]: Train_Auc_Heatmap(gsv)
```



## Test Auc Heatmap

In [29]: Train\_Auc\_Heatmap(gsv)



## Training the model

```
In [30]: Best_Param=gsv.best_params_
depth=Best_Param['max_depth']
n_estimator =Best_Param['n_estimators']

Model_Bow=RandomForestClassifier(max_depth=depth,n_estimators=n_estimator,class_weight="balanced")
Model_Bow.fit(X_train_bow,Y_train)
```

```
Out[30]: RandomForestClassifier(bootstrap=True, class_weight='balanced',
                                criterion='gini', max_depth=500, max_features='auto',
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                n_estimators=100, n_jobs=None, oob_score=False,
                                random_state=None, verbose=0, warm_start=False)
```

```
In [44]: from sklearn.externals import joblib
joblib.dump(Model_Bow,"BOW_Model.pkl")
```

```
Out[44]: ['BOW_Model.pkl']
```


```
In [18]: from sklearn.externals import joblib
Model_Bow=joblib.load("BOW_Model.pkl")
```

## Evaluating the performance of model




```
In [31]: trainconfusionmatrix(Model_Bow,X_train_bow,Y_train)
```

Confusion Matrix for Train set




```
In [32]: testconfusionmatrix(Model_Bow,X_test_bow,Y_test)
```

Confusion Matrix for Test set



```
In [34]: plot_auc_roc(Model_Bow,X_train_bow,X_test_bow,Y_train,Y_test)
```



```
In [35]: print("Classification Report: \n")
y_pred=Model_Bow.predict(X_test_bow)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.12	0.22	4478
1	0.88	1.00	0.94	28522
micro avg	0.88	0.88	0.88	33000
macro avg	0.92	0.56	0.58	33000
weighted avg	0.89	0.88	0.84	33000


## Important Feature

```
In [39]: from wordcloud import WordCloud, STOPWORDS

features =vectorizer.get_feature_names()
coef = Model_Bow.feature_importances_
coef_df = pd.DataFrame({'word': features, 'coefficient': coef}, index = None)
df = coef_df.sort_values("coefficient", ascending = False)[:100]
cloud = " ".join(word for word in df.word)
stopwords = set(STOPWORDS)
wordcloud = WordCloud(width = 1000, height = 600, background_color = 'white', s
topwords = stopwords).generate(cloud)

# plot the WordCloud image
plt.figure(figsize = (10, 8))
plt.imshow(wordcloud, interpolation = 'bilinear')
plt.axis("off")
#plt.title("Top 100 most important features\n")
plt.tight_layout(pad = 0)

plt.show()
```



## TF-IDF Vectorizer

```
In [40]: vectorizer_tfidf=TfidfVectorizer()
         vectorizer_tfidf.fit(X_train)
```

```
Out[40]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
                        dtype=<class 'numpy.float64'>, encoding='utf-8', input='content',
                        lowercase=True, max_df=1.0, max_features=None, min_df=1,
                        ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_idf=True,
                        stop_words=None, strip_accents=None, sublinear_tf=False,
                        token_pattern='(?u)\\b\\w\\w+\\b', tokenizer=None, use_idf=True,
                        vocabulary=None)
```

```
In [41]: X_Train_Tfidf=vectorizer_tfidf.transform(X_train)
         X_Train_Tfidf=preprocessing.normalize(X_Train_Tfidf)

         X_Test_Tfidf=vectorizer_tfidf.transform(X_test)
         X_Test_Tfidf=preprocessing.normalize(X_Test_Tfidf)
```

```
In [42]: print("Shape of Train and Test Data After vectorizations")
         print(X_Train_Tfidf.shape, Y_train.shape)
         print(X_Test_Tfidf.shape, Y_test.shape)
```

```
Shape of Train and Test Data After vectorizations
(67000, 30737) (67000,)
(33000, 30737) (33000,)
```

## Finding the best value Of hyperparameter


```
In [43]: gsv=Grid_SearchCV(X_Train_Tfidf,Y_train)

         print("Best HyperParameter: ",gsv.best_params_)
         print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

```
Best HyperParameter: {'max_depth': 1000, 'n_estimators': 100}
Best Accuracy: 90.99%
```


## Train Auc HeatMap

In [44]: Train\_Auc\_Heatmap(gsv)



## Test AUC Heatmap

In [45]: Cv\_Auc\_Heatmap(gsv)



## Training the model

```
In [47]: Best_Param=gsv.best_params_
depth=Best_Param['max_depth']
n_estimator =Best_Param['n_estimators']


Model_Tfidf=RandomForestClassifier(max_depth=depth,n_estimators=n_estimator,class_weight="balanced")
Model_Tfidf.fit(X_Train_Tfidf,Y_train)
```

```
Out[47]: RandomForestClassifier(bootstrap=True, class_weight='balanced',
                                criterion='gini', max_depth=1000, max_features='auto',
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                n_estimators=100, n_jobs=None, oob_score=False,
                                random_state=None, verbose=0, warm_start=False)
```

## Evaluating the performance of model


```
In [48]: trainconfusionmatrix(Model_Tfidf,X_Train_Tfidf,Y_train)
```

Confusion Matrix for Train set




```
In [49]: testconfusionmatrix(Model_Tfidf,X_Test_Tfidf,Y_test)
```

Confusion Matrix for Test set



```
In [50]: plot_auc_roc(Model_Tfidf,X_Train_Tfidf,X_Test_Tfidf,Y_train,Y_test)
```



```
In [51]: print("Classification Report: \n")
y_pred=Model_Tfidf.predict(X_Test_Tfidf)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.12	0.22	4478
1	0.88	1.00	0.93	28522
micro avg	0.88	0.88	0.88	33000
macro avg	0.91	0.56	0.58	33000
weighted avg	0.89	0.88	0.84	33000

## Important Features




```
In [53]: from wordcloud import WordCloud, STOPWORDS

features =vectorizer.get_feature_names()
coef = Model_Tfidf.feature_importances_
coef_df = pd.DataFrame({'word': features, 'coefficient': coef}, index = None)
df = coef_df.sort_values("coefficient", ascending = False)[:100]
cloud = " ".join(word for word in df.word)
stopwords = set(STOPWORDS)
wordcloud = WordCloud(width = 1000, height = 600, background_color = 'white', s
topwords = stopwords).generate(cloud)

# plot the WordCloud image
plt.figure(figsize = (10, 8))
plt.imshow(wordcloud, interpolation = 'bilinear')
plt.axis("off")
#plt.title("Top 100 most important features\n")
plt.tight_layout(pad = 0)

plt.show()
```



## Word To Vector

```
In [54]: list_of_Train_sent=[]
          list_of_Test_sent=[]

          for sent in X_train:
              list_of_Train_sent.append(sent.split())

          for sent in X_test:
              list_of_Test_sent.append(sent.split())
```

```
In [55]: model=Word2Vec(list_of_Train_sent,min_count=5,size=50, workers=4)
```

## Average Word To Vector

```
In [56]: import numpy as np

Train_vectors = []
for sent in list_of_Train_sent:
    sent_vec = np.zeros(50)
    cnt_words = 0
    for word in sent:
        try:
            vec = model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            pass
    if cnt_words!=0:
        sent_vec /= cnt_words
    Train_vectors.append(sent_vec)
Train_vectors = np.nan_to_num(Train_vectors)
```

```
In [57]: import numpy as np

Test_vectors=[]
for sent in list_of_Test_sent:
    sent_vec=np.zeros(50)
    cnt_words=0
    for word in sent:
        try:
            vec=model.wv[word]
            sent_vec+=vec
            cnt_words+=1
        except:
            pass
    if cnt_words!=0:
        sent_vec/=cnt_words
    Test_vectors.append(sent_vec)
Test_vectors=np.nan_to_num(Test_vectors)
```

```
In [58]: print("Shape of Test Vectors : ",Test_vectors.shape)
```

Shape of Test Vectors : (33000, 50)

```
In [59]: X_Train_Awv=Train_vectors
X_Test_Awv=Test_vectors
```

```
In [60]: print(X_Train_Awv.shape, Y_train.shape)
print(X_Test_Awv.shape, Y_test.shape)

(67000, 50) (67000,)
(33000, 50) (33000,)
```

## Finding the best value Of hyperparameter


```
In [61]: gsv=Grid_SearchCV(X_Train_Awv,Y_train)

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))

Best HyperParameter: {'max_depth': 50, 'n_estimators': 100}
Best Accuracy: 87.15%
```


## Train Auc Heatmap

```
In [62]: Train_Auc_Heatmap(gsv)
```



## Test Auc Heatmap

In [63]: Cv\_Auc\_Heatmap(gsv)



## Training Model

```
In [64]: Best_Param=gsv.best_params_
depth=Best_Param['max_depth']
n_estimator =Best_Param['n_estimators']


Model_Awv=RandomForestClassifier(max_depth=depth,n_estimators=n_estimator,class_weight="balanced")
Model_Awv.fit(X_Train_Awv,Y_train)
```

```
Out[64]: RandomForestClassifier(bootstrap=True, class_weight='balanced',
                                criterion='gini', max_depth=50, max_features='auto',
                                max_leaf_nodes=None, min_impurity_decrease=0.0,
                                min_impurity_split=None, min_samples_leaf=1,
                                min_samples_split=2, min_weight_fraction_leaf=0.0,
                                n_estimators=100, n_jobs=None, oob_score=False,
                                random_state=None, verbose=0, warm_start=False)
```

## Evaluating the performance of model


```
In [65]: trainconfusionmatrix(Model_Awv,X_Train_Awv,Y_train)
```

Confusion Matrix for Train set

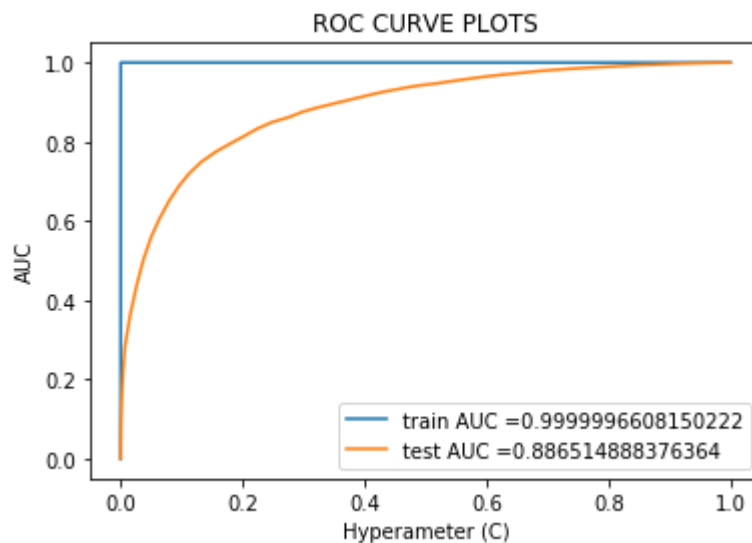


```
In [66]: trainconfusionmatrix(Model_Awv,X_Test_Awv,Y_test)
```

Confusion Matrix for Train set



```
In [69]: plot_auc_roc(Model_Awv,X_Train_Awv,X_Test_Awv,Y_train,Y_test)
```



```
In [68]: print("Classification Report: \n")
y_pred=Model_Awv.predict(X_Test_Awv)
print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.77	0.17	0.28	4478
1	0.88	0.99	0.93	28522
micro avg	0.88	0.88	0.88	33000
macro avg	0.83	0.58	0.61	33000
weighted avg	0.87	0.88	0.85	33000

## TF-IDF Word To Vector

```
In [70]: TFIDF_Feature=vectorizer_tfidf.get_feature_names()
print(len(TFIDF_Feature))
print(TFIDF_Feature[0:20])
```

30737

```
['aaa', 'aaaaaaaaagghh', 'aaaaah', 'aaaaahhhhhhhhhhhhhhhhh', 'aaaah', 'aaah',
'aachen', 'aad', 'aadp', 'aafco', 'aagh', 'aah', 'aahh', 'aand', 'aardvark',
'ab', 'aback', 'abandon', 'abaolut', 'abattoir']
```

```
In [71]: w2v_words = list(model.wv.vocab)
print("Number of words that occurred minimum 5 times ", len(w2v_words))
print("Sample words ", w2v_words[0:50])
```

Number of words that occurred minimum 5 times 10488

Sample words ['witti', 'littl', 'book', 'make', 'son', 'laugh', 'loud', 'ca  
r', 'drive', 'along', 'always', 'sing', 'refrain', 'hes', 'learn', 'whale', 'i  
ndia', 'droop', 'love', 'new', 'word', 'introduc', 'silli', 'classic', 'wil  
l', 'bet', 'still', 'abl', 'memori', 'colleg', 'rememb', 'see', 'show', 'ai  
r', 'televis', 'year', 'ago', 'child', 'sister', 'later', 'bought', 'day', 't  
hirti', 'someth', 'use', 'seri', 'song', 'student', 'teach', 'preschool']

```
In [72]: from tqdm import tqdm
Train_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Train_sent):
    sent_vec=np.zeros(50)
    weight=0
    for word in sent:
        try :
            if word in TFIDF_Feature and word in w2v_words:

                w2v_vec=model.wv[word]
                tfidf_vec=X_Train_Tfidf[row,TFIDF_Feature.index(word)]
                sent_vec+=(w2v_vec*tfidf_vec)
                weight+=tfidf_vec

        except :
            pass
    if weight!=0:
        sent_vec/=weight
    Train_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1
```

```
100%|███████████████████████| 67000/67000 [1:00:12<00:00, 18.55it/
```

s]

```
In [73]: Test_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Test_sent):
    sent_vec=np.zeros(50)
    weight=0

    for word in sent:
        try:
            if word in TFIDF_Feature and word in w2v_words:

                w2v_vec=model.wv[word]
                tfidf_vec=X_Test_Tfidf(row,TFIDF_Feature.index(word))
                sent_vec+=(w2v_vec*tfidf_vec)
                weight+=tfidf

        except :
            pass

    if weight!=0:
        sent_vec/=weight
    Test_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1
```

100%|██| 33000/33000 [23:28<00:00, 23.42it/s]

```
In [74]: Train_TFIDF_W2V_Vectors = np.nan_to_num(Train_TFIDF_W2V_Vectors)
Test_TFIDF_W2V_Vectors = np.nan_to_num(Test_TFIDF_W2V_Vectors)
```

```
In [75]: X_Train_TfidfW2v=Train_TFIDF_W2V_Vectors
X_Test_TfidfW2v=Test_TFIDF_W2V_Vectors
```

```
In [76]: print(X_Train_TfidfW2v.shape, Y_train.shape)
print(X_Test_TfidfW2v.shape, Y_test.shape)

(67000, 50) (67000,)
(33000, 50) (33000,)
```

## Finding the best value Of hyperparameter

```
In [77]: gsv=Grid_SearchCV(X_Train_TfidfW2v,Y_train)


print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))

Best HyperParameter: {'max_depth': 500, 'n_estimators': 100}
Best Accuracy: 84.17%
```

## Train Auc Heatmap




In [78]: Train\_Auc\_Heatmap(gsv)



## Test AUC Heatmap

In [79]: Cv\_Auc\_Heatmap(gsv)



## Training Model

```
In [80]: Best_Param=gsv.best_params_
n_estimator =Best_Param['n_estimators']


Model_TfidfW2v=RandomForestClassifier(max_depth=depth,n_estimators=n_estimator
,class_weight="balanced")
Model_TfidfW2v.fit(X_Train_TfidfW2v,Y_train)
```

```
Out[80]: RandomForestClassifier(bootstrap=True, class_weight='balanced',
criterion='gini', max_depth=50, max_features='auto',
max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=100, n_jobs=None, oob_score=False,
random_state=None, verbose=0, warm_start=False)
```

## Evaluating the performance of model


```
In [81]: trainconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Train set




```
In [82]: testconfusionmatrix(Model_TfidfW2v,X_Test_TfidfW2v,Y_test)
```

Confusion Matrix for Test set



```
In [83]: plot_auc_roc(Model_TfidfW2v,X_Train_TfidfW2v,X_Test_TfidfW2v,Y_train,Y_test)
```



```
In [84]: print("Classification Report: \n")
y_pred=Model_TfidfW2v.predict(X_Test_TfidfW2v)

print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	4478
1	0.86	1.00	0.93	28522
micro avg	0.86	0.86	0.86	33000
macro avg	0.43	0.50	0.46	33000
weighted avg	0.75	0.86	0.80	33000

## Conclusion

### 1. Report on different vectorizer method

```
In [2]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Hyperparameter(max_depth)", "Hyperparameter(n_e
stimulator)", "Train AUC", "Test AUC", "F1-Score"]

x.add_row(["BOW", 500, 100, 0.99, 0.92, 0.84])
x.add_row(["TF-IDF", 1000, 100, 0.99, 0.92, 0.84])
x.add_row(["Avg W2V", 50, 100, 0.99, 0.88, 0.85])
x.add_row(["TF-IDF W2V", 500, 100, 0.99, 0.50, 0.80])

print(x)
```

```
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+
| Vectorizer | Hyperparameter(max_depth) | Hyperparameter(n_estimator) | Train AUC | Test AUC | F1-Score |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+
| BOW | 500 | 100 | 0.99 | 0.92 | 0.84 |
| TF-IDF | 1000 | 100 | 0.99 | 0.92 | 0.84 |
| Avg W2V | 50 | 100 | 0.99 | 0.88 | 0.85 |
| TF-IDF W2V | 500 | 100 | 0.99 | 0.50 | 0.80 |
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+
+-----+-----+-----+
```

1. I have taken two hyper-parameter to train model that is max\_depth and n\_estimator.
2. In BOW and TF-IDF , model is performing well .
3. Model is Overfit in case of TF-IDF W2v .