

## Assignment-12 : Apply Keras on MNIST DataSet

```
In [0]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal

%matplotlib notebook
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
```

```
In [0]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
def plot_loss(x, vy, ty, xlabel='Epoch', ylabel='Categorical Crossentropy Loss'):
    _, = plt.plot(x, vy, 'b', label="Validation Loss")
    _, = plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel(xlabel)
    plt.ylabel(ylabel)

    plt.grid()
    plt.legend()
    plt.grid()
    plt.show()
```

```
In [3]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()

Downloading data from https://s3.amazonaws.com/img-datasets/mnist.npz
11493376/11490434 [=====] - 2s 0us/step
```

```
In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
 Number of training examples : 10000 and each image is of shape (28, 28)

```
In [0]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [6]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of
shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of
shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)  
 Number of training examples : 10000 and each image is of shape (784)

```
In [0]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255

X_train = X_train/255
X_test = X_test/255
```

```
In [8]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# Lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5  
 After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```
In [0]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

## Multi-Layer Perceptron With 2 hidden layer

## 1. MLP + Relu Activation + Adam Optimizer (H1: 720 , H2 : 200)

```
In [0]: from keras.models import Sequential
        from keras.layers import Dense, Activation, Dropout, BatchNormalization
```

```
In [13]: %%time
def Build_NN_2(input_dim, output_dim=10):
    model = Sequential()
    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))

    model.add(Dense(200, activation='relu', kernel_initializer='he_normal'))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 4 µs, sys: 0 ns, total: 4 µs  
Wall time: 23.8 µs

```
In [17]: model=Build_NN_2(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
```

```
W0713 10:20:14.764899 140220561528704 deprecation.py:323] From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use tf.where in 2.0, which has the same broadcast rule as np.where
```

Layer (type)	Output Shape	Param #
dense_10 (Dense)	(None, 720)	565200
dense_11 (Dense)	(None, 200)	144200
dense_12 (Dense)	(None, 10)	2010
Total params: 711,410		
Trainable params: 711,410		
Non-trainable params: 0		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 5s 127us/step - loss: 0.2554 - acc: 0.9246 - val\_loss: 0.1375 - val\_acc: 0.9583

Epoch 2/20

42000/42000 [=====] - 1s 34us/step - loss: 0.0917 - acc: 0.9724 - val\_loss: 0.1157 - val\_acc: 0.9638

Epoch 3/20

42000/42000 [=====] - 1s 34us/step - loss: 0.0564 - acc: 0.9823 - val\_loss: 0.1064 - val\_acc: 0.9690

Epoch 4/20

42000/42000 [=====] - 1s 35us/step - loss: 0.0345 - acc: 0.9894 - val\_loss: 0.1036 - val\_acc: 0.9711

Epoch 5/20

42000/42000 [=====] - 2s 36us/step - loss: 0.0250 - acc: 0.9921 - val\_loss: 0.0896 - val\_acc: 0.9744

Epoch 6/20

42000/42000 [=====] - 2s 37us/step - loss: 0.0181 - acc: 0.9939 - val\_loss: 0.0993 - val\_acc: 0.9742

Epoch 7/20

42000/42000 [=====] - 2s 37us/step - loss: 0.0199 - acc: 0.9936 - val\_loss: 0.0978 - val\_acc: 0.9754

Epoch 8/20

42000/42000 [=====] - 2s 36us/step - loss: 0.0182 - acc: 0.9944 - val\_loss: 0.1027 - val\_acc: 0.9735

Epoch 9/20

42000/42000 [=====] - 2s 37us/step - loss: 0.0075 - acc: 0.9978 - val\_loss: 0.1142 - val\_acc: 0.9750

Epoch 10/20

42000/42000 [=====] - 1s 35us/step - loss: 0.0122 - acc: 0.9955 - val\_loss: 0.1199 - val\_acc: 0.9717

Epoch 11/20

42000/42000 [=====] - 2s 36us/step - loss: 0.0101 - acc: 0.9967 - val\_loss: 0.1259 - val\_acc: 0.9741

Epoch 12/20

42000/42000 [=====] - 1s 34us/step - loss: 0.0095 - acc: 0.9972 - val\_loss: 0.1256 - val\_acc: 0.9704

Epoch 13/20

42000/42000 [=====] - 1s 34us/step - loss: 0.0118 - acc: 0.9961 - val\_loss: 0.1064 - val\_acc: 0.9773

Epoch 14/20

42000/42000 [=====] - 1s 34us/step - loss: 0.0103 -

```

acc: 0.9965 - val_loss: 0.1131 - val_acc: 0.9760
Epoch 15/20
42000/42000 [=====] - 1s 35us/step - loss: 0.0072 -
acc: 0.9978 - val_loss: 0.1271 - val_acc: 0.9744
Epoch 16/20
42000/42000 [=====] - 1s 34us/step - loss: 0.0073 -
acc: 0.9974 - val_loss: 0.1220 - val_acc: 0.9767
Epoch 17/20
42000/42000 [=====] - 1s 34us/step - loss: 0.0044 -
acc: 0.9986 - val_loss: 0.1356 - val_acc: 0.9746
Epoch 18/20
42000/42000 [=====] - 1s 35us/step - loss: 0.0108 -
acc: 0.9965 - val_loss: 0.1343 - val_acc: 0.9759
Epoch 19/20
42000/42000 [=====] - 1s 34us/step - loss: 0.0070 -
acc: 0.9977 - val_loss: 0.1302 - val_acc: 0.9754
Epoch 20/20
42000/42000 [=====] - 1s 34us/step - loss: 0.0049 -
acc: 0.9986 - val_loss: 0.1358 - val_acc: 0.9744

```

## Plot

```


In [32]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)-

```

Test Score: 0.11015889572340877

Test Accuracy: 0.9787



```

In [33]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## 2. MLP + Relu Activation + Adam Optimizer + Batch Normalization (H1: 720 , H2: 200 )



```
In [34]: %%time
def Build_NN_2(input_dim, output_dim=10):
    model = Sequential()
    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
    model.add(BatchNormalization())

    model.add(Dense(200, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 8  $\mu$ s, sys: 0 ns, total: 8  $\mu$ s

Wall time: 12.6  $\mu$ s

```
In [37]: model=Build_NN_2(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
```

Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 720)	565200
batch_normalization_1 (Batch Normalization)	(None, 720)	2880
dense_15 (Dense)	(None, 200)	144200
batch_normalization_2 (Batch Normalization)	(None, 200)	800
dense_16 (Dense)	(None, 10)	2010
Total params: 715,090		
Trainable params: 713,250		
Non-trainable params: 1,840		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 3s 74us/step - loss: 0.1973 - acc: 0.9413 - val\_loss: 0.1216 - val\_acc: 0.9645

Epoch 2/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0669 - acc: 0.9805 - val\_loss: 0.1032 - val\_acc: 0.9691

Epoch 3/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0423 - acc: 0.9869 - val\_loss: 0.1019 - val\_acc: 0.9706

Epoch 4/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0281 - acc: 0.9911 - val\_loss: 0.1041 - val\_acc: 0.9691

Epoch 5/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0236 - acc: 0.9928 - val\_loss: 0.1242 - val\_acc: 0.9671

Epoch 6/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0170 - acc: 0.9949 - val\_loss: 0.1174 - val\_acc: 0.9693

Epoch 7/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0176 - acc: 0.9939 - val\_loss: 0.1178 - val\_acc: 0.9687

Epoch 8/20

42000/42000 [=====] - 2s 53us/step - loss: 0.0189 - acc: 0.9940 - val\_loss: 0.1072 - val\_acc: 0.9727

Epoch 9/20

42000/42000 [=====] - 2s 55us/step - loss: 0.0182 - acc: 0.9942 - val\_loss: 0.1138 - val\_acc: 0.9711

Epoch 10/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0109 - acc: 0.9966 - val\_loss: 0.0973 - val\_acc: 0.9764

Epoch 11/20

42000/42000 [=====] - 2s 54us/step - loss: 0.0108 - acc: 0.9965 - val\_loss: 0.1255 - val\_acc: 0.9703

Epoch 12/20

42000/42000 [=====] - 2s 53us/step - loss: 0.0145 - acc: 0.9954 - val\_loss: 0.1061 - val\_acc: 0.9752

Epoch 13/20

```

42000/42000 [=====] - 2s 54us/step - loss: 0.0110 -
acc: 0.9963 - val_loss: 0.1055 - val_acc: 0.9760
Epoch 14/20
42000/42000 [=====] - 2s 54us/step - loss: 0.0101 -
acc: 0.9967 - val_loss: 0.1100 - val_acc: 0.9742
Epoch 15/20
42000/42000 [=====] - 2s 54us/step - loss: 0.0148 -
acc: 0.9954 - val_loss: 0.1037 - val_acc: 0.9747
Epoch 16/20
42000/42000 [=====] - 2s 54us/step - loss: 0.0153 -
acc: 0.9948 - val_loss: 0.0960 - val_acc: 0.9768
Epoch 17/20
42000/42000 [=====] - 2s 54us/step - loss: 0.0065 -
acc: 0.9979 - val_loss: 0.0935 - val_acc: 0.9789
Epoch 18/20
42000/42000 [=====] - 2s 54us/step - loss: 0.0032 -
acc: 0.9993 - val_loss: 0.0913 - val_acc: 0.9788
Epoch 19/20
42000/42000 [=====] - 2s 54us/step - loss: 0.0049 -
acc: 0.9987 - val_loss: 0.1083 - val_acc: 0.9769
Epoch 20/20
42000/42000 [=====] - 2s 53us/step - loss: 0.0099 -
acc: 0.9970 - val_loss: 0.1097 - val_acc: 0.9766

```

```


In [38]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)

```

Test Score: 0.08973172496785191

Test Accuracy: 0.9797



```

In [39]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



### 3. MLP + Relu Activation + Adam Optimizer + Dropout (H1: 720 , H2: 200)

```
In [40]: %%time
def Build_NN_2(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
    model.add(Dropout(0.5))

    model.add(Dense(200, activation='relu', kernel_initializer='he_normal'))
    model.add(Dropout(0.5))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 4  $\mu$ s, sys: 1e+03 ns, total: 5  $\mu$ s

Wall time: 7.63  $\mu$ s

```
In [41]: model=Build_NN_2(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
```

```
W0713 10:58:05.901090 140220561528704 deprecation.py:506] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.  
Instructions for updating:  
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.
```



Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 720)	565200
dropout_1 (Dropout)	(None, 720)	0
dense_18 (Dense)	(None, 200)	144200
dropout_2 (Dropout)	(None, 200)	0
dense_19 (Dense)	(None, 10)	2010
Total params: 711,410		
Trainable params: 711,410		
Non-trainable params: 0		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 2s 53us/step - loss: 0.4603 - acc: 0.8564 - val\_loss: 0.1784 - val\_acc: 0.9467

Epoch 2/20

42000/42000 [=====] - 2s 38us/step - loss: 0.2027 - acc: 0.9388 - val\_loss: 0.1210 - val\_acc: 0.9636

Epoch 3/20

42000/42000 [=====] - 2s 38us/step - loss: 0.1512 - acc: 0.9545 - val\_loss: 0.1053 - val\_acc: 0.9689

Epoch 4/20

42000/42000 [=====] - 2s 38us/step - loss: 0.1256 - acc: 0.9629 - val\_loss: 0.0953 - val\_acc: 0.9715

Epoch 5/20

42000/42000 [=====] - 2s 38us/step - loss: 0.1081 - acc: 0.9675 - val\_loss: 0.0913 - val\_acc: 0.9722

Epoch 6/20

42000/42000 [=====] - 2s 38us/step - loss: 0.0946 - acc: 0.9708 - val\_loss: 0.0898 - val\_acc: 0.9736

Epoch 7/20

42000/42000 [=====] - 2s 38us/step - loss: 0.0853 - acc: 0.9739 - val\_loss: 0.0831 - val\_acc: 0.9765

Epoch 8/20

42000/42000 [=====] - 2s 38us/step - loss: 0.0795 - acc: 0.9755 - val\_loss: 0.0820 - val\_acc: 0.9758

Epoch 9/20

42000/42000 [=====] - 2s 38us/step - loss: 0.0743 - acc: 0.9770 - val\_loss: 0.0861 - val\_acc: 0.9761

Epoch 10/20

42000/42000 [=====] - 2s 37us/step - loss: 0.0671 - acc: 0.9791 - val\_loss: 0.0797 - val\_acc: 0.9770

Epoch 11/20

42000/42000 [=====] - 2s 38us/step - loss: 0.0605 - acc: 0.9804 - val\_loss: 0.0853 - val\_acc: 0.9775

Epoch 12/20

42000/42000 [=====] - 2s 38us/step - loss: 0.0608 - acc: 0.9811 - val\_loss: 0.0874 - val\_acc: 0.9771

Epoch 13/20

```

42000/42000 [=====] - 2s 38us/step - loss: 0.0577 -
acc: 0.9816 - val_loss: 0.0828 - val_acc: 0.9778
Epoch 14/20
42000/42000 [=====] - 2s 38us/step - loss: 0.0530 -
acc: 0.9832 - val_loss: 0.0869 - val_acc: 0.9776
Epoch 15/20
42000/42000 [=====] - 2s 38us/step - loss: 0.0515 -
acc: 0.9836 - val_loss: 0.0838 - val_acc: 0.9784
Epoch 16/20
42000/42000 [=====] - 2s 38us/step - loss: 0.0497 -
acc: 0.9837 - val_loss: 0.0791 - val_acc: 0.9793
Epoch 17/20
42000/42000 [=====] - 2s 37us/step - loss: 0.0485 -
acc: 0.9844 - val_loss: 0.0853 - val_acc: 0.9783
Epoch 18/20
42000/42000 [=====] - 2s 39us/step - loss: 0.0451 -
acc: 0.9857 - val_loss: 0.0814 - val_acc: 0.9798
Epoch 19/20
42000/42000 [=====] - 2s 38us/step - loss: 0.0409 -
acc: 0.9865 - val_loss: 0.0827 - val_acc: 0.9794
Epoch 20/20
42000/42000 [=====] - 2s 38us/step - loss: 0.0428 -
acc: 0.9864 - val_loss: 0.0891 - val_acc: 0.9784

```


```

In [43]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)

```

Test Score: 0.07207091313775582  
Test Accuracy: 0.9807



```

In [44]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



#### 4. MLP + Relu Activation + Adam Optimizer + Batch Normalization + Dropout (H1: 720 , H2: 200)

```
In [45]: %%time
def Build_NN_2(input_dim, output_dim=10):
    model = Sequential()
    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))

    model.add(Dense(200, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 10 µs, sys: 2 µs, total: 12 µs  
Wall time: 16.5 µs

```
In [46]: model=Build_NN_2(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_20 (Dense)	(None, 720)	565200
batch_normalization_3 (Batch Normalization)	(None, 720)	2880
dropout_3 (Dropout)	(None, 720)	0
dense_21 (Dense)	(None, 200)	144200
batch_normalization_4 (Batch Normalization)	(None, 200)	800
dropout_4 (Dropout)	(None, 200)	0
dense_22 (Dense)	(None, 10)	2010
Total params: 715,090		
Trainable params: 713,250		
Non-trainable params: 1,840		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 3s 79us/step - loss: 0.4287 - acc: 0.8700 - val\_loss: 0.1650 - val\_acc: 0.9502

Epoch 2/20

42000/42000 [=====] - 2s 56us/step - loss: 0.2026 - acc: 0.9368 - val\_loss: 0.1310 - val\_acc: 0.9621

Epoch 3/20

42000/42000 [=====] - 2s 57us/step - loss: 0.1581 - acc: 0.9517 - val\_loss: 0.1083 - val\_acc: 0.9676

Epoch 4/20

42000/42000 [=====] - 2s 56us/step - loss: 0.1290 - acc: 0.9601 - val\_loss: 0.0974 - val\_acc: 0.9701

Epoch 5/20

42000/42000 [=====] - 2s 56us/step - loss: 0.1146 - acc: 0.9644 - val\_loss: 0.0926 - val\_acc: 0.9729

Epoch 6/20

42000/42000 [=====] - 2s 57us/step - loss: 0.1031 - acc: 0.9670 - val\_loss: 0.0872 - val\_acc: 0.9742

Epoch 7/20

42000/42000 [=====] - 2s 56us/step - loss: 0.0968 - acc: 0.9692 - val\_loss: 0.0922 - val\_acc: 0.9726

Epoch 8/20

42000/42000 [=====] - 2s 57us/step - loss: 0.0834 - acc: 0.9737 - val\_loss: 0.0845 - val\_acc: 0.9750

Epoch 9/20

42000/42000 [=====] - 2s 57us/step - loss: 0.0772 - acc: 0.9759 - val\_loss: 0.0797 - val\_acc: 0.9771

Epoch 10/20

42000/42000 [=====] - 2s 58us/step - loss: 0.0737 - acc: 0.9767 - val\_loss: 0.0857 - val\_acc: 0.9752

Epoch 11/20

42000/42000 [=====] - 2s 58us/step - loss: 0.0731 - acc: 0.9771 - val\_loss: 0.0794 - val\_acc: 0.9771


```
Epoch 12/20
42000/42000 [=====] - 2s 57us/step - loss: 0.0653 -
acc: 0.9791 - val_loss: 0.0829 - val_acc: 0.9775
Epoch 13/20
42000/42000 [=====] - 2s 56us/step - loss: 0.0631 -
acc: 0.9795 - val_loss: 0.0811 - val_acc: 0.9769
Epoch 14/20
42000/42000 [=====] - 2s 56us/step - loss: 0.0577 -
acc: 0.9811 - val_loss: 0.0737 - val_acc: 0.9796
Epoch 15/20
42000/42000 [=====] - 2s 56us/step - loss: 0.0589 -
acc: 0.9807 - val_loss: 0.0864 - val_acc: 0.9760
Epoch 16/20
42000/42000 [=====] - 2s 57us/step - loss: 0.0568 -
acc: 0.9815 - val_loss: 0.0794 - val_acc: 0.9781
Epoch 17/20
42000/42000 [=====] - 2s 56us/step - loss: 0.0566 -
acc: 0.9820 - val_loss: 0.0762 - val_acc: 0.9794
Epoch 18/20
42000/42000 [=====] - 2s 56us/step - loss: 0.0507 -
acc: 0.9829 - val_loss: 0.0813 - val_acc: 0.9772
Epoch 19/20
42000/42000 [=====] - 2s 56us/step - loss: 0.0469 -
acc: 0.9842 - val_loss: 0.0782 - val_acc: 0.9790
Epoch 20/20
42000/42000 [=====] - 2s 57us/step - loss: 0.0451 -
acc: 0.9853 - val_loss: 0.0776 - val_acc: 0.9796
```

```
In [47]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)
```

Test Score: 0.06645472465842323

Test Accuracy: 0.9822





```

In [48]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## Multi-Layer Perceptron With 3 hidden layer

### 1. MLP + Relu Activation + Adam Optimizer (H1: 700 , H2 : 360 , H3: 180)

```
In [49]: %%time
def Build_NN_3(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(700, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))

    model.add(Dense(360, activation='relu', kernel_initializer='he_normal'))

    model.add(Dense(180, activation='relu', kernel_initializer='he_normal'))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 5 µs, sys: 0 ns, total: 5 µs

Wall time: 9.54 µs

```
In [57]: model=Build_NN_3(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 700)	549500
batch_normalization_9 (Batch Normalization)	(None, 700)	2800
dense_30 (Dense)	(None, 360)	252360
batch_normalization_10 (Batch Normalization)	(None, 360)	1440
dense_31 (Dense)	(None, 180)	64980
batch_normalization_11 (Batch Normalization)	(None, 180)	720
dense_32 (Dense)	(None, 10)	1810
Total params: 873,610		
Trainable params: 871,130		
Non-trainable params: 2,480		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 4s 101us/step - loss: 0.1929 - acc: 0.9418 - val\_loss: 0.1217 - val\_acc: 0.9631

Epoch 2/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0718 - acc: 0.9776 - val\_loss: 0.1013 - val\_acc: 0.9697

Epoch 3/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0444 - acc: 0.9862 - val\_loss: 0.1031 - val\_acc: 0.9703

Epoch 4/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0429 - acc: 0.9857 - val\_loss: 0.1023 - val\_acc: 0.9693

Epoch 5/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0375 - acc: 0.9877 - val\_loss: 0.0931 - val\_acc: 0.9747

Epoch 6/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0217 - acc: 0.9929 - val\_loss: 0.0917 - val\_acc: 0.9756

Epoch 7/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0171 - acc: 0.9947 - val\_loss: 0.1140 - val\_acc: 0.9703

Epoch 8/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0270 - acc: 0.9909 - val\_loss: 0.0941 - val\_acc: 0.9746

Epoch 9/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0137 - acc: 0.9953 - val\_loss: 0.1088 - val\_acc: 0.9757

Epoch 10/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0147 - acc: 0.9948 - val\_loss: 0.1233 - val\_acc: 0.9704

Epoch 11/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0155 - acc: 0.9948 - val\_loss: 0.1002 - val\_acc: 0.9758


```
Epoch 12/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0133 -
acc: 0.9956 - val_loss: 0.1110 - val_acc: 0.9739
Epoch 13/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0190 -
acc: 0.9934 - val_loss: 0.1125 - val_acc: 0.9716
Epoch 14/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0119 -
acc: 0.9961 - val_loss: 0.1202 - val_acc: 0.9730
Epoch 15/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0122 -
acc: 0.9959 - val_loss: 0.1025 - val_acc: 0.9761
Epoch 16/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0097 -
acc: 0.9965 - val_loss: 0.1077 - val_acc: 0.9761
Epoch 17/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0107 -
acc: 0.9967 - val_loss: 0.1074 - val_acc: 0.9763
Epoch 18/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0170 -
acc: 0.9942 - val_loss: 0.1056 - val_acc: 0.9750
Epoch 19/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0116 -
acc: 0.9968 - val_loss: 0.0988 - val_acc: 0.9777
Epoch 20/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0094 -
acc: 0.9970 - val_loss: 0.0998 - val_acc: 0.9785
```

```
In [58]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)
```

Test Score: 0.09005934175993825

Test Accuracy: 0.981



```

In [59]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## 2. MLP + Relu Activation + Adam Optimizer + Batch Normalization (H1: 700 , H2: 360 , H3: 180)

```
In [60]: %%time
def Build_NN_3(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(700, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
    model.add(BatchNormalization())

    model.add(Dense(360, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())

    model.add(Dense(180, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 10  $\mu$ s, sys: 0 ns, total: 10  $\mu$ s  
Wall time: 14.3  $\mu$ s



```
In [61]: model=Build_NN_3(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
          validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_33 (Dense)	(None, 700)	549500
batch_normalization_12 (Batch Normalization)	(None, 700)	2800
dense_34 (Dense)	(None, 360)	252360
batch_normalization_13 (Batch Normalization)	(None, 360)	1440
dense_35 (Dense)	(None, 180)	64980
batch_normalization_14 (Batch Normalization)	(None, 180)	720
dense_36 (Dense)	(None, 10)	1810
Total params: 873,610		
Trainable params: 871,130		
Non-trainable params: 2,480		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 4s 105us/step - loss: 0.1984 - acc: 0.9399 - val\_loss: 0.1328 - val\_acc: 0.9596

Epoch 2/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0689 - acc: 0.9785 - val\_loss: 0.1093 - val\_acc: 0.9658

Epoch 3/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0495 - acc: 0.9847 - val\_loss: 0.1037 - val\_acc: 0.9711

Epoch 4/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0369 - acc: 0.9879 - val\_loss: 0.0927 - val\_acc: 0.9742

Epoch 5/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0356 - acc: 0.9882 - val\_loss: 0.0964 - val\_acc: 0.9733

Epoch 6/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0253 - acc: 0.9914 - val\_loss: 0.1031 - val\_acc: 0.9722

Epoch 7/20

42000/42000 [=====] - 3s 68us/step - loss: 0.0213 - acc: 0.9931 - val\_loss: 0.0923 - val\_acc: 0.9758

Epoch 8/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0210 - acc: 0.9932 - val\_loss: 0.1169 - val\_acc: 0.9687

Epoch 9/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0154 - acc: 0.9952 - val\_loss: 0.1039 - val\_acc: 0.9731

Epoch 10/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0159 - acc: 0.9948 - val\_loss: 0.1059 - val\_acc: 0.9739

Epoch 11/20

42000/42000 [=====] - 3s 67us/step - loss: 0.0172 - acc: 0.9942 - val\_loss: 0.1093 - val\_acc: 0.9746


```
Epoch 12/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0110 -
acc: 0.9964 - val_loss: 0.1062 - val_acc: 0.9739
Epoch 13/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0179 -
acc: 0.9942 - val_loss: 0.0878 - val_acc: 0.9795
Epoch 14/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0141 -
acc: 0.9957 - val_loss: 0.1068 - val_acc: 0.9751
Epoch 15/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0092 -
acc: 0.9971 - val_loss: 0.1050 - val_acc: 0.9762
Epoch 16/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0160 -
acc: 0.9947 - val_loss: 0.1053 - val_acc: 0.9758
Epoch 17/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0159 -
acc: 0.9948 - val_loss: 0.1017 - val_acc: 0.9761
Epoch 18/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0073 -
acc: 0.9978 - val_loss: 0.0908 - val_acc: 0.9790
Epoch 19/20
42000/42000 [=====] - 3s 67us/step - loss: 0.0072 -
acc: 0.9979 - val_loss: 0.0987 - val_acc: 0.9777
Epoch 20/20
42000/42000 [=====] - 3s 68us/step - loss: 0.0082 -
acc: 0.9970 - val_loss: 0.1068 - val_acc: 0.9775
```

```
In [62]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)
```

Test Score: 0.08819649735184722

Test Accuracy: 0.9787



```

In [63]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



### 3. MLP + Relu Activation + Adam Optimizer + Dropout (H1: 700 , H2: 360 , H3: 180)

```
In [64]: %time
def Build_NN_3(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(700, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))

    model.add(Dense(360, activation='relu', kernel_initializer='he_normal'))

    model.add(Dense(180, activation='relu', kernel_initializer='he_normal'))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 3  $\mu$ s, sys: 1e+03 ns, total: 4  $\mu$ s  
Wall time: 7.15  $\mu$ s

```
In [65]: model=Build_NN_3(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_37 (Dense)	(None, 700)	549500
dense_38 (Dense)	(None, 360)	252360
dense_39 (Dense)	(None, 180)	64980
dense_40 (Dense)	(None, 10)	1810
Total params: 868,650		
Trainable params: 868,650		
Non-trainable params: 0		

None

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 3s 72us/step - loss: 0.2423 - acc: 0.9256 - val\_loss: 0.1264 - val\_acc: 0.9612

Epoch 2/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0896 - acc: 0.9723 - val\_loss: 0.1058 - val\_acc: 0.9699

Epoch 3/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0524 - acc: 0.9838 - val\_loss: 0.1206 - val\_acc: 0.9635

Epoch 4/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0399 - acc: 0.9864 - val\_loss: 0.1011 - val\_acc: 0.9712

Epoch 5/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0250 - acc: 0.9921 - val\_loss: 0.1080 - val\_acc: 0.9724

Epoch 6/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0246 - acc: 0.9912 - val\_loss: 0.1324 - val\_acc: 0.9676

Epoch 7/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0226 - acc: 0.9927 - val\_loss: 0.1324 - val\_acc: 0.9698

Epoch 8/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0211 - acc: 0.9930 - val\_loss: 0.1179 - val\_acc: 0.9712

Epoch 9/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0154 - acc: 0.9949 - val\_loss: 0.1241 - val\_acc: 0.9739

Epoch 10/20

42000/42000 [=====] - 2s 40us/step - loss: 0.0218 - acc: 0.9925 - val\_loss: 0.1305 - val\_acc: 0.9701

Epoch 11/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0125 - acc: 0.9958 - val\_loss: 0.1172 - val\_acc: 0.9747

Epoch 12/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0097 - acc: 0.9965 - val\_loss: 0.1398 - val\_acc: 0.9723

Epoch 13/20

42000/42000 [=====] - 2s 39us/step - loss: 0.0137 - acc: 0.9957 - val\_loss: 0.1286 - val\_acc: 0.9738



```

Epoch 14/20
42000/42000 [=====] - 2s 39us/step - loss: 0.0111 -
acc: 0.9963 - val_loss: 0.1207 - val_acc: 0.9779
Epoch 15/20
42000/42000 [=====] - 2s 40us/step - loss: 0.0132 -
acc: 0.9959 - val_loss: 0.1233 - val_acc: 0.9749
Epoch 16/20
42000/42000 [=====] - 2s 39us/step - loss: 0.0137 -
acc: 0.9964 - val_loss: 0.1285 - val_acc: 0.9734
Epoch 17/20
42000/42000 [=====] - 2s 40us/step - loss: 0.0081 -
acc: 0.9973 - val_loss: 0.1169 - val_acc: 0.9784
Epoch 18/20
42000/42000 [=====] - 2s 39us/step - loss: 0.0109 -
acc: 0.9965 - val_loss: 0.1368 - val_acc: 0.9756
Epoch 19/20
42000/42000 [=====] - 2s 40us/step - loss: 0.0056 -
acc: 0.9985 - val_loss: 0.1886 - val_acc: 0.9686
Epoch 20/20
42000/42000 [=====] - 2s 39us/step - loss: 0.0095 -
acc: 0.9972 - val_loss: 0.1356 - val_acc: 0.9761

```

```


In [66]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)

```

Test Score: 0.11028124448086878

Test Accuracy: 0.9769



```

In [67]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



#### 4. MLP + Relu Activation + Adam Optimizer + Batch Normalization + Dropout (H1: 700 , H2: 360 , H3: 180)

```
In [0]: def Build_NN_3(input_dim, output_dim=10):
        model = Sequential()
        model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
        model.add(BatchNormalization())
        model.add(Dropout(0.5))

        model.add(Dense(200, activation='relu', kernel_initializer='he_normal'))
        model.add(BatchNormalization())
        model.add(Dropout(0.5))

        model.add(Dense(output_dim, activation='softmax'))

        model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
        print(model.summary())

        return model
```

```
In [69]: model=Build_NN_3(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 720)	565200
batch_normalization_15 (Batch Normalization)	(None, 720)	2880
dropout_9 (Dropout)	(None, 720)	0
dense_42 (Dense)	(None, 200)	144200
batch_normalization_16 (Batch Normalization)	(None, 200)	800
dropout_10 (Dropout)	(None, 200)	0
dense_43 (Dense)	(None, 10)	2010
Total params: 715,090		
Trainable params: 713,250		
Non-trainable params: 1,840		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 4s 96us/step - loss: 0.4246 - acc: 0.8723 - val\_loss: 0.1577 - val\_acc: 0.9528

Epoch 2/20

42000/42000 [=====] - 3s 60us/step - loss: 0.2050 - acc: 0.9397 - val\_loss: 0.1266 - val\_acc: 0.9621

Epoch 3/20

42000/42000 [=====] - 3s 60us/step - loss: 0.1575 - acc: 0.9512 - val\_loss: 0.1114 - val\_acc: 0.9671

Epoch 4/20

42000/42000 [=====] - 3s 60us/step - loss: 0.1341 - acc: 0.9597 - val\_loss: 0.0980 - val\_acc: 0.9701

Epoch 5/20

42000/42000 [=====] - 3s 61us/step - loss: 0.1157 - acc: 0.9634 - val\_loss: 0.0908 - val\_acc: 0.9726

Epoch 6/20

42000/42000 [=====] - 3s 60us/step - loss: 0.1023 - acc: 0.9679 - val\_loss: 0.0927 - val\_acc: 0.9742

Epoch 7/20

42000/42000 [=====] - 2s 58us/step - loss: 0.0950 - acc: 0.9703 - val\_loss: 0.0869 - val\_acc: 0.9761

Epoch 8/20

42000/42000 [=====] - 2s 58us/step - loss: 0.0882 - acc: 0.9718 - val\_loss: 0.0897 - val\_acc: 0.9743

Epoch 9/20

42000/42000 [=====] - 2s 58us/step - loss: 0.0839 - acc: 0.9744 - val\_loss: 0.0816 - val\_acc: 0.9764

Epoch 10/20

42000/42000 [=====] - 2s 58us/step - loss: 0.0788 - acc: 0.9747 - val\_loss: 0.0806 - val\_acc: 0.9767

Epoch 11/20

42000/42000 [=====] - 2s 59us/step - loss: 0.0692 - acc: 0.9781 - val\_loss: 0.0815 - val\_acc: 0.9769


```
Epoch 12/20
42000/42000 [=====] - 3s 60us/step - loss: 0.0635 -
acc: 0.9792 - val_loss: 0.0824 - val_acc: 0.9772
Epoch 13/20
42000/42000 [=====] - 3s 60us/step - loss: 0.0661 -
acc: 0.9787 - val_loss: 0.0818 - val_acc: 0.9777
Epoch 14/20
42000/42000 [=====] - 2s 58us/step - loss: 0.0582 -
acc: 0.9810 - val_loss: 0.0822 - val_acc: 0.9775
Epoch 15/20
42000/42000 [=====] - 2s 58us/step - loss: 0.0540 -
acc: 0.9820 - val_loss: 0.0851 - val_acc: 0.9777
Epoch 16/20
42000/42000 [=====] - 2s 58us/step - loss: 0.0568 -
acc: 0.9820 - val_loss: 0.0820 - val_acc: 0.9786
Epoch 17/20
42000/42000 [=====] - 2s 58us/step - loss: 0.0534 -
acc: 0.9832 - val_loss: 0.0803 - val_acc: 0.9790
Epoch 18/20
42000/42000 [=====] - 2s 58us/step - loss: 0.0506 -
acc: 0.9837 - val_loss: 0.0808 - val_acc: 0.9790
Epoch 19/20
42000/42000 [=====] - 2s 58us/step - loss: 0.0498 -
acc: 0.9838 - val_loss: 0.0784 - val_acc: 0.9802
Epoch 20/20
42000/42000 [=====] - 2s 58us/step - loss: 0.0485 -
acc: 0.9830 - val_loss: 0.0780 - val_acc: 0.9803
```

```
In [70]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)
```

Test Score: 0.06508726586559788

Test Accuracy: 0.9819



```

In [71]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## Multi-Layer Perceptron With 5 hidden layer

### 1. MLP + Relu Activation + Adam Optimizer (H1: 720 , H2 : 360 , H3: 180 , H4: 90 , H5: 45)



```
In [72]: %%time
def Build_NN_5(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))

    model.add(Dense(360, activation='relu', kernel_initializer='he_normal'))
    model.add(Dense(180, activation='relu', kernel_initializer='he_normal'))
    model.add(Dense(90, activation='relu', kernel_initializer='he_normal'))
    model.add(Dense(45, activation='relu', kernel_initializer='he_normal'))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 5  $\mu$ s, sys: 0 ns, total: 5  $\mu$ s

Wall time: 10.5  $\mu$ s

```
In [73]: model=Build_NN_5(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_44 (Dense)	(None, 720)	565200
dense_45 (Dense)	(None, 360)	259560
dense_46 (Dense)	(None, 180)	64980
dense_47 (Dense)	(None, 90)	16290
dense_48 (Dense)	(None, 45)	4095
dense_49 (Dense)	(None, 10)	460
Total params: 910,585		
Trainable params: 910,585		
Non-trainable params: 0		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 4s 87us/step - loss: 0.2700 - acc: 0.9175 - val\_loss: 0.1370 - val\_acc: 0.9585

Epoch 2/20

42000/42000 [=====] - 2s 47us/step - loss: 0.1000 - acc: 0.9695 - val\_loss: 0.1245 - val\_acc: 0.9629

Epoch 3/20

42000/42000 [=====] - 2s 46us/step - loss: 0.0632 - acc: 0.9804 - val\_loss: 0.1086 - val\_acc: 0.9690

Epoch 4/20

42000/42000 [=====] - 2s 46us/step - loss: 0.0461 - acc: 0.9850 - val\_loss: 0.1095 - val\_acc: 0.9692

Epoch 5/20

42000/42000 [=====] - 2s 45us/step - loss: 0.0389 - acc: 0.9879 - val\_loss: 0.1203 - val\_acc: 0.9688

Epoch 6/20

42000/42000 [=====] - 2s 45us/step - loss: 0.0330 - acc: 0.9895 - val\_loss: 0.1003 - val\_acc: 0.9739

Epoch 7/20

42000/42000 [=====] - 2s 46us/step - loss: 0.0276 - acc: 0.9915 - val\_loss: 0.1175 - val\_acc: 0.9709

Epoch 8/20

42000/42000 [=====] - 2s 45us/step - loss: 0.0224 - acc: 0.9923 - val\_loss: 0.1141 - val\_acc: 0.9722

Epoch 9/20

42000/42000 [=====] - 2s 45us/step - loss: 0.0247 - acc: 0.9922 - val\_loss: 0.0996 - val\_acc: 0.9757

Epoch 10/20

42000/42000 [=====] - 2s 46us/step - loss: 0.0186 - acc: 0.9938 - val\_loss: 0.1086 - val\_acc: 0.9741

Epoch 11/20

42000/42000 [=====] - 2s 47us/step - loss: 0.0186 - acc: 0.9942 - val\_loss: 0.1003 - val\_acc: 0.9785

Epoch 12/20

42000/42000 [=====] - 2s 45us/step - loss: 0.0152 -

```

acc: 0.9951 - val_loss: 0.1275 - val_acc: 0.9745
Epoch 13/20
42000/42000 [=====] - 2s 45us/step - loss: 0.0176 -
acc: 0.9945 - val_loss: 0.1158 - val_acc: 0.9747
Epoch 14/20
42000/42000 [=====] - 2s 45us/step - loss: 0.0186 -
acc: 0.9940 - val_loss: 0.1110 - val_acc: 0.9771
Epoch 15/20
42000/42000 [=====] - 2s 45us/step - loss: 0.0134 -
acc: 0.9959 - val_loss: 0.1244 - val_acc: 0.9718
Epoch 16/20
42000/42000 [=====] - 2s 46us/step - loss: 0.0109 -
acc: 0.9965 - val_loss: 0.1117 - val_acc: 0.9760
Epoch 17/20
42000/42000 [=====] - 2s 45us/step - loss: 0.0117 -
acc: 0.9965 - val_loss: 0.1298 - val_acc: 0.9733
Epoch 18/20
42000/42000 [=====] - 2s 46us/step - loss: 0.0129 -
acc: 0.9960 - val_loss: 0.1011 - val_acc: 0.9780
Epoch 19/20
42000/42000 [=====] - 2s 46us/step - loss: 0.0126 -
acc: 0.9963 - val_loss: 0.1372 - val_acc: 0.9706
Epoch 20/20
42000/42000 [=====] - 2s 45us/step - loss: 0.0110 -
acc: 0.9967 - val_loss: 0.1305 - val_acc: 0.9772

```

```

In [74]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')


x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)

```

```

Test Score: 0.08912124511943212
Test Accuracy: 0.9823

```



```

In [75]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



**2. MLP + Relu Activation + Adam Optimizer + Batch Normalization (H1: 720 , H2 : 360 , H3: 180 , H4: 90 , H5: 45)**

```
In [76]: %%time
def Build_NN_5(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
    model.add(BatchNormalization())

    model.add(Dense(360, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())

    model.add(Dense(180, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())

    model.add(Dense(90, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())

    model.add(Dense(45, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 0 ns, sys: 9 µs, total: 9 µs  
Wall time: 13.1 µs

```
In [77]: model=Build_NN_5(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
          validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_50 (Dense)	(None, 720)	565200
batch_normalization_17 (Batch Normalization)	(None, 720)	2880
dense_51 (Dense)	(None, 360)	259560
batch_normalization_18 (Batch Normalization)	(None, 360)	1440
dense_52 (Dense)	(None, 180)	64980
batch_normalization_19 (Batch Normalization)	(None, 180)	720
dense_53 (Dense)	(None, 90)	16290
batch_normalization_20 (Batch Normalization)	(None, 90)	360
dense_54 (Dense)	(None, 45)	4095
batch_normalization_21 (Batch Normalization)	(None, 45)	180
dense_55 (Dense)	(None, 10)	460
Total params: 916,165		
Trainable params: 913,375		
Non-trainable params: 2,790		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 6s 154us/step - loss: 0.2544 - acc: 0.9275 - val\_loss: 0.1409 - val\_acc: 0.9579

Epoch 2/20

42000/42000 [=====] - 4s 93us/step - loss: 0.0949 - acc: 0.9723 - val\_loss: 0.1084 - val\_acc: 0.9677

Epoch 3/20

42000/42000 [=====] - 4s 94us/step - loss: 0.0569 - acc: 0.9834 - val\_loss: 0.1041 - val\_acc: 0.9702

Epoch 4/20

42000/42000 [=====] - 4s 95us/step - loss: 0.0509 - acc: 0.9838 - val\_loss: 0.1150 - val\_acc: 0.9663

Epoch 5/20

42000/42000 [=====] - 4s 94us/step - loss: 0.0418 - acc: 0.9865 - val\_loss: 0.1010 - val\_acc: 0.9714

Epoch 6/20

42000/42000 [=====] - 4s 94us/step - loss: 0.0401 - acc: 0.9879 - val\_loss: 0.0937 - val\_acc: 0.9753

Epoch 7/20

42000/42000 [=====] - 4s 96us/step - loss: 0.0284 - acc: 0.9905 - val\_loss: 0.1068 - val\_acc: 0.9724

Epoch 8/20

42000/42000 [=====] - 4s 95us/step - loss: 0.0287 - acc: 0.9906 - val\_loss: 0.1064 - val\_acc: 0.9718

Epoch 9/20




```
42000/42000 [=====] - 4s 93us/step - loss: 0.0256 -  
acc: 0.9918 - val_loss: 0.0921 - val_acc: 0.9753  
Epoch 10/20  
42000/42000 [=====] - 4s 94us/step - loss: 0.0272 -  
acc: 0.9915 - val_loss: 0.0980 - val_acc: 0.9744  
Epoch 11/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0294 -  
acc: 0.9900 - val_loss: 0.1029 - val_acc: 0.9743  
Epoch 12/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0219 -  
acc: 0.9927 - val_loss: 0.1064 - val_acc: 0.9718  
Epoch 13/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0167 -  
acc: 0.9944 - val_loss: 0.1004 - val_acc: 0.9759  
Epoch 14/20  
42000/42000 [=====] - 4s 94us/step - loss: 0.0154 -  
acc: 0.9951 - val_loss: 0.0965 - val_acc: 0.9766  
Epoch 15/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0140 -  
acc: 0.9955 - val_loss: 0.0932 - val_acc: 0.9782  
Epoch 16/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0199 -  
acc: 0.9934 - val_loss: 0.1133 - val_acc: 0.9726  
Epoch 17/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0172 -  
acc: 0.9945 - val_loss: 0.0950 - val_acc: 0.9779  
Epoch 18/20  
42000/42000 [=====] - 4s 94us/step - loss: 0.0116 -  
acc: 0.9964 - val_loss: 0.0938 - val_acc: 0.9765  
Epoch 19/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0198 -  
acc: 0.9937 - val_loss: 0.0997 - val_acc: 0.9764  
Epoch 20/20  
42000/42000 [=====] - 4s 93us/step - loss: 0.0101 -  
acc: 0.9970 - val_loss: 0.0982 - val_acc: 0.9769
```

```
In [78]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)
```

Test Score: 0.09106637627696619

Test Accuracy: 0.9772



```

In [79]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



### 3. MLP + Relu Activation + Adam Optimizer + Dropout (H1: 720 , H2: 360 , H3: 180 , H4: 90 , H5: 45)

```
In [80]: %%time
def Build_NN_5(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
    model.add(Dropout(0.5))

    model.add(Dense(360, activation='relu', kernel_initializer='he_normal'))
    model.add(Dropout(0.5))

    model.add(Dense(180, activation='relu', kernel_initializer='he_normal'))
    model.add(Dropout(0.5))

    model.add(Dense(90, activation='relu', kernel_initializer='he_normal'))
    model.add(Dropout(0.5))

    model.add(Dense(45, activation='relu', kernel_initializer='he_normal'))
    model.add(Dropout(0.5))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 6 µs, sys: 0 ns, total: 6 µs  
Wall time: 9.78 µs

```
In [81]: model=Build_NN_5(input_dim)
         print()
         history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
         validation_split=0.3, verbose=1)
         print()
```

Layer (type)	Output Shape	Param #
dense_56 (Dense)	(None, 720)	565200
dropout_11 (Dropout)	(None, 720)	0
dense_57 (Dense)	(None, 360)	259560
dropout_12 (Dropout)	(None, 360)	0
dense_58 (Dense)	(None, 180)	64980
dropout_13 (Dropout)	(None, 180)	0
dense_59 (Dense)	(None, 90)	16290
dropout_14 (Dropout)	(None, 90)	0
dense_60 (Dense)	(None, 45)	4095
dropout_15 (Dropout)	(None, 45)	0
dense_61 (Dense)	(None, 10)	460
Total params: 910,585		
Trainable params: 910,585		
Non-trainable params: 0		

None

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 4s 100us/step - loss: 1.7393 - acc: 0.3657 - val\_loss: 0.6480 - val\_acc: 0.8299

Epoch 2/20

42000/42000 [=====] - 2s 51us/step - loss: 0.7234 - acc: 0.7688 - val\_loss: 0.2972 - val\_acc: 0.9284

Epoch 3/20

42000/42000 [=====] - 2s 51us/step - loss: 0.4746 - acc: 0.8765 - val\_loss: 0.2142 - val\_acc: 0.9476

Epoch 4/20

42000/42000 [=====] - 2s 52us/step - loss: 0.3734 - acc: 0.9103 - val\_loss: 0.1886 - val\_acc: 0.9560

Epoch 5/20

42000/42000 [=====] - 2s 53us/step - loss: 0.3239 - acc: 0.9246 - val\_loss: 0.1751 - val\_acc: 0.9594

Epoch 6/20

42000/42000 [=====] - 2s 52us/step - loss: 0.2874 - acc: 0.9342 - val\_loss: 0.1633 - val\_acc: 0.9612

Epoch 7/20

42000/42000 [=====] - 2s 52us/step - loss: 0.2557 - acc: 0.9405 - val\_loss: 0.1579 - val\_acc: 0.9654

Epoch 8/20

42000/42000 [=====] - 2s 52us/step - loss: 0.2316 - acc: 0.9480 - val\_loss: 0.1607 - val\_acc: 0.9656

Epoch 9/20


```
42000/42000 [=====] - 2s 51us/step - loss: 0.2238 -  
acc: 0.9493 - val_loss: 0.1692 - val_acc: 0.9669  
Epoch 10/20  
42000/42000 [=====] - 2s 52us/step - loss: 0.2101 -  
acc: 0.9530 - val_loss: 0.1490 - val_acc: 0.9666  
Epoch 11/20  
42000/42000 [=====] - 2s 52us/step - loss: 0.1951 -  
acc: 0.9570 - val_loss: 0.1288 - val_acc: 0.9718  
Epoch 12/20  
42000/42000 [=====] - 2s 52us/step - loss: 0.1886 -  
acc: 0.9581 - val_loss: 0.1527 - val_acc: 0.9694  
Epoch 13/20  
42000/42000 [=====] - 2s 51us/step - loss: 0.1723 -  
acc: 0.9620 - val_loss: 0.1479 - val_acc: 0.9713  
Epoch 14/20  
42000/42000 [=====] - 2s 51us/step - loss: 0.1686 -  
acc: 0.9625 - val_loss: 0.1383 - val_acc: 0.9724  
Epoch 15/20  
42000/42000 [=====] - 2s 52us/step - loss: 0.1632 -  
acc: 0.9637 - val_loss: 0.1408 - val_acc: 0.9738  
Epoch 16/20  
42000/42000 [=====] - 2s 51us/step - loss: 0.1544 -  
acc: 0.9666 - val_loss: 0.1496 - val_acc: 0.9709  
Epoch 17/20  
42000/42000 [=====] - 2s 52us/step - loss: 0.1463 -  
acc: 0.9678 - val_loss: 0.1291 - val_acc: 0.9730  
Epoch 18/20  
42000/42000 [=====] - 2s 51us/step - loss: 0.1386 -  
acc: 0.9688 - val_loss: 0.1261 - val_acc: 0.9745  
Epoch 19/20  
42000/42000 [=====] - 2s 52us/step - loss: 0.1402 -  
acc: 0.9695 - val_loss: 0.1256 - val_acc: 0.9744  
Epoch 20/20  
42000/42000 [=====] - 2s 51us/step - loss: 0.1322 -  
acc: 0.9706 - val_loss: 0.1360 - val_acc: 0.9751
```

```
In [82]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)
```

Test Score: 0.11893713512013138

Test Accuracy: 0.9776





```

In [83]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



**4. MLP + Relu Activation + Adam Optimizer + Batch Normalization + Dropout (H1: 720 , H2: 360 , H3: 180 , H4: 180 , H5: 45)**

```
In [84]: %%time
def Build_NN_5(input_dim, output_dim=10):
    model = Sequential()

    model.add(Dense(720, activation='relu', kernel_initializer='he_normal', input_shape=(input_dim,)))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))

    model.add(Dense(360, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))

    model.add(Dense(180, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))

    model.add(Dense(90, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))

    model.add(Dense(45, activation='relu', kernel_initializer='he_normal'))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))

    model.add(Dense(output_dim, activation='softmax'))

    model.compile(optimizer='Adam', loss='categorical_crossentropy', metrics=['accuracy'])
    print(model.summary())

    return model
```

CPU times: user 7  $\mu$ s, sys: 1e+03 ns, total: 8  $\mu$ s

Wall time: 11.4  $\mu$ s

```
In [85]: model=Build_NN_5(input_dim)
          print()
          history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
                              validation_split=0.3, verbose=1)
          print()
```

Layer (type)	Output Shape	Param #
dense_62 (Dense)	(None, 720)	565200
batch_normalization_22 (Batch Normalization)	(None, 720)	2880
dropout_16 (Dropout)	(None, 720)	0
dense_63 (Dense)	(None, 360)	259560
batch_normalization_23 (Batch Normalization)	(None, 360)	1440
dropout_17 (Dropout)	(None, 360)	0
dense_64 (Dense)	(None, 180)	64980
batch_normalization_24 (Batch Normalization)	(None, 180)	720
dropout_18 (Dropout)	(None, 180)	0
dense_65 (Dense)	(None, 90)	16290
batch_normalization_25 (Batch Normalization)	(None, 90)	360
dropout_19 (Dropout)	(None, 90)	0
dense_66 (Dense)	(None, 45)	4095
batch_normalization_26 (Batch Normalization)	(None, 45)	180
dropout_20 (Dropout)	(None, 45)	0
dense_67 (Dense)	(None, 10)	460
Total params: 916,165		
Trainable params: 913,375		
Non-trainable params: 2,790		
None		

Train on 42000 samples, validate on 18000 samples

Epoch 1/20

42000/42000 [=====] - 7s 174us/step - loss: 1.4528 - acc: 0.5301 - val\_loss: 0.3335 - val\_acc: 0.9094

Epoch 2/20

42000/42000 [=====] - 4s 99us/step - loss: 0.5557 - acc: 0.8435 - val\_loss: 0.2039 - val\_acc: 0.9429

Epoch 3/20

42000/42000 [=====] - 4s 100us/step - loss: 0.3801 - acc: 0.8986 - val\_loss: 0.1790 - val\_acc: 0.9519

Epoch 4/20

42000/42000 [=====] - 4s 99us/step - loss: 0.3034 - acc: 0.9224 - val\_loss: 0.1574 - val\_acc: 0.9589

Epoch 5/20

42000/42000 [=====] - 4s 98us/step - loss: 0.2567 - acc: 0.9350 - val\_loss: 0.1489 - val\_acc: 0.9623


```
Epoch 6/20
42000/42000 [=====] - 4s 100us/step - loss: 0.2315 -
acc: 0.9423 - val_loss: 0.1436 - val_acc: 0.9649
Epoch 7/20
42000/42000 [=====] - 4s 101us/step - loss: 0.2072 -
acc: 0.9472 - val_loss: 0.1352 - val_acc: 0.9662
Epoch 8/20
42000/42000 [=====] - 4s 101us/step - loss: 0.1939 -
acc: 0.9512 - val_loss: 0.1239 - val_acc: 0.9695
Epoch 9/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1764 -
acc: 0.9554 - val_loss: 0.1240 - val_acc: 0.9707
Epoch 10/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1626 -
acc: 0.9611 - val_loss: 0.1164 - val_acc: 0.9731
Epoch 11/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1725 -
acc: 0.9576 - val_loss: 0.1039 - val_acc: 0.9752
Epoch 12/20
42000/42000 [=====] - 4s 100us/step - loss: 0.1498 -
acc: 0.9632 - val_loss: 0.1109 - val_acc: 0.9731
Epoch 13/20
42000/42000 [=====] - 4s 100us/step - loss: 0.1426 -
acc: 0.9650 - val_loss: 0.1057 - val_acc: 0.9747
Epoch 14/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1412 -
acc: 0.9649 - val_loss: 0.1155 - val_acc: 0.9740
Epoch 15/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1345 -
acc: 0.9671 - val_loss: 0.1104 - val_acc: 0.9739
Epoch 16/20
42000/42000 [=====] - 4s 100us/step - loss: 0.1268 -
acc: 0.9688 - val_loss: 0.1060 - val_acc: 0.9759
Epoch 17/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1247 -
acc: 0.9706 - val_loss: 0.1005 - val_acc: 0.9761
Epoch 18/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1169 -
acc: 0.9716 - val_loss: 0.0963 - val_acc: 0.9778
Epoch 19/20
42000/42000 [=====] - 4s 100us/step - loss: 0.1152 -
acc: 0.9719 - val_loss: 0.1065 - val_acc: 0.9750
Epoch 20/20
42000/42000 [=====] - 4s 99us/step - loss: 0.1069 -
acc: 0.9746 - val_loss: 0.1025 - val_acc: 0.9763
```

```
In [86]: score = model.evaluate(X_test, Y_test, verbose=0)
print(f'Test Score: {score[0]}')
print(f'Test Accuracy: {score[1]}\n')

x = list(range(1, nb_epoch+1))
vy = history.history['val_loss']
ty = history.history['loss']
plot_loss(x, vy, ty)
```

Test Score: 0.08941399474844802

Test Accuracy: 0.9788



```

In [87]: w_after = model.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



## Conclusion

1. I have tried 3 different architectures having layers 2, 3 and 5 respectively.
1. Subsequent increase in depth of the Network (3 and 5 Hidden Layers) shows not much improvement in the Test Accuracy.

1. Also, the Crossover point between Train and Validation Loss is increasing as the Networks depth increases. This is evident from the Loss vs Epoch curves.