

# Assignment-5 - Apply-Logistic \_Regression -On-Amazon-Review-Dataset

April 19, 2019

## 1 Assignment-5: Apply Logistic Regression On Amazon Fine Food Reviews DataSet

### 1.1 Introduction

(i).Logistic regression is used to describe data and to explain the relationship between one dependent binary variable and one or more nominal, ordinal, interval or ratio-level independent variables.

(ii).It's name is Regression but actually it is classification algorithm.

### 1.2 Objective

To Predict the Polarity of Amazon Fine Food Review Using Logistic Regression Algorithm.

### 1.3 Importing All Required Library

```
In [1]: %matplotlib inline
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
import math

from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.linear_model import LogisticRegression

from sklearn.metrics import classification_report
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import TimeSeriesSplit
```

```

from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import cross_val_score
from sklearn import preprocessing
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score

from gensim.models import Word2Vec
from gensim.models import KeyedVectors

import pickle

from tqdm import tqdm
import os
import warnings
warnings.filterwarnings("ignore")

```

C:\Users\User\Anaconda3\lib\site-packages\gensim\utils.py:1197: UserWarning: detected Windows; warnings.warn("detected Windows; aliasing chunkize to chunkize\_serial")

## 1.4 Importing Amazon Fine Food Review Dataset

```

In [2]: if os.path.isfile("final.sqlite"):
        conn=sqlite3.connect("final.sqlite")
        Data=pd.read_sql_query("select * from Logistic_Regression_Data where Score!=3",conn)
        conn.close()
    else :
        print("Error Importing the file")

```

In [3]: *# Printing some data of DataFrame*

```
Data['Score'].value_counts()
```

```

Out[3]: 1    126413
        0     23587
        Name: Score, dtype: int64

```

## 1.5 Information About DataSet

```

In [4]: print("\nNumber of Reviews: ",Data["Text"].count())
        print("\nNumber of Users: ",len(Data["UserId"].unique())) # Unique returns 1-D array o
        print("\nNumber of Products: ",len(Data["ProductId"].unique()))
        print("\nShape of Data: ", Data.shape)
        print("\nColumn Name of DataSet : ",Data.columns)
        print("\n\nNumber of Attributes/Columns in data: 12")
        print("\nNumber of Positive Reviews : ", Data['Score'].value_counts()[1])
        print("\nNumber of Negative Reviews : ", Data['Score'].value_counts()[0])

```

Number of Reviews: 150000

Number of Users: 115887

Number of Products: 42869

Shape of Data: (150000, 13)

```
Column Name of DataSet : Index(['level_0', 'index', 'Id', 'ProductId', 'UserId', 'ProfileName',
                                'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',
                                'Summary', 'Text', 'CleanedText'],
                                dtype='object')
```

Number of Attributes/Columns in data: 12

Number of Positive Reviews : 126413

Number of Negative Reviews : 23587

```
In [5]: print("\nNumber of Reviews: ",Data["Text"].count())
```

Number of Reviews: 150000

## 1.6 Attribute Information About DataSet

- 1.Id - A unique value starts from 1
- 2.ProductId - A unique identifier for the product
- 3.UserId - A unique identifier for the user
- 4.ProfileName - Name of user profile
- 5.HelpfulnessNumerator - Number of users who found the review helpful
- 6.HelpfulnessDenominator - Number of users who indicated whether they found the review helpful or not
- 7.Score - Rating 0 or 1
- 8.Time - Timestamp for the review
- 9.Summary - Brief summary of the review
- 10.Text - Text of the review
- 11.Cleaned Text - Text that only alphabets

```
In [6]: # Sorting on the basis of Time Parameter
        Data.sort_values('Time',inplace=True)
```

```
In [7]: Y = Data['Score'].values
        X = Data['CleanedText'].values
```

## 1.7 Splitting DataSet into Train and Test Data

```
In [8]: from sklearn.model_selection import train_test_split
        # X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, shuffle=False)
        X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33, shuffle=False)

        print("Shape of Train and Test Dataset for 50k points")
        print(X_train.shape, Y_train.shape)
        print(X_test.shape, Y_test.shape)
```

```
Shape of Train and Test Dataset for 50k points
(100500,) (100500,)
(49500,) (49500,)
```

## 1.8 Defining Some Function

### 1.8.1 Train Data Confusion Matrix Plot

```
In [9]: def trainconfusionmatrix(model,X_train,y_train):
        print("Confusion Matrix for Train set")
        cm=confusion_matrix(y_train, model.predict(X_train))
        class_label = ["negative", "positive"]
        df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
        sns.heatmap(df_cm, annot = True, fmt = "d")
        plt.title("Train Confusiion Matrix")
        plt.xlabel("Predicted Label")
        plt.ylabel("True Label")
        plt.show()
```

### 1.8.2 Test Data Confusion Matrix Plot

```
In [10]: def testconfusionmatrix(model,X_test,y_test):
        print("Confusion Matrix for Test set")
        cm=confusion_matrix(y_test, model.predict(X_test))
        class_label = ["negative", "positive"]
        df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
        sns.heatmap(df_cm, annot = True, fmt = "d")
        plt.title("Test Confusiion Matrix")
        plt.xlabel("Predicted Label")
        plt.ylabel("True Label")
        plt.show()
```

### 1.8.3 ROC-AUC Curve Plot

```
In [11]: def plot_auc_roc(model,X_train,X_test,y_train,y_test):
        train_fpr, train_tpr, thresholds = roc_curve(y_train, model.predict_proba(X_train)[:,1])
        test_fpr, test_tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test)[:,1])
```

```

plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))
plt.plot(test_fpr, test_tpr, label="test AUC =" + str(auc(test_fpr, test_tpr)))
plt.legend()
plt.xlabel("Hyperparameter (C)")
plt.ylabel("AUC")
plt.title("ROC CURVE PLOTS")
plt.show()

```

#### 1.8.4 Error Plot

In [12]: `def plot(gsv):`

```

    Res=gsv.cv_results_
    cv_auc=[(1-x) for x in Res['mean_test_score']]
    train_auc = [(1-x) for x in Res['mean_train_score']]

    x1=[10000,5000,1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0005,0.0001]
    log_c= [math.log10(x) for x in x1]

    plt.xlabel("Hyperparameter (Log(C))",fontsize=15)
    plt.ylabel("Missclassification Error",fontsize=15)
    plt.title('Misclassification Error v/s C',fontsize=15)
    plt.plot(log_c,cv_auc,label="Cross-Validation")
    plt.plot(log_c,train_auc,label="Train")
    plt.legend()
    plt.show()

```

#### 1.8.5 GridSearchCV

In [13]: `def Grid_SearchCV(X_train,Y_train,regularisation_parameter):`

```

    param_grid = {'C':[10000,5000,1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0005,0.0001],
                  'penalty':[regularisation_parameter]}
    tscv = TimeSeriesSplit(n_splits=10)
    LR = LogisticRegression(class_weight='balanced')

    gsv=GridSearchCV(LR,param_grid,cv=tscv,verbose=1,scoring='roc_auc')
    gsv.fit(X_train,Y_train)

    return gsv

```

#### 1.8.6 30 Informative Feature

In [14]: `def show_30_informative_feature(vectorizer,model,n=30):`

```

Weights_Index = model.coef_[0].argsort()
length = len(Weights_Index)
# For Negative Class
neg_class=Weights_Index[:30]

neg_feat=[vectorizer.get_feature_names()[x] for x in neg_class]
neg_prob=[model.coef_[0][x] for x in neg_class]

neg_zip=list(zip(neg_feat,neg_prob))
neg_zip.sort()

# For Positive Class
pos_class = Weights_Index[-30:]

pos_feat=[vectorizer.get_feature_names()[x] for x in pos_class]
pos_prob=[model.coef_[0][x] for x in pos_class]
pos_zip=list(zip(pos_feat,pos_prob))
pos_zip.sort()

top=zip(pos_zip,neg_zip)

print("{0:20}{1:55}{2:20}".format("S.N", "Positive", "Negative"))
print("_"*90)
i=1
for (fn_1,coef_1), (fn_2,coef_2) in top:
    print("%d.\t\t%.3f\t%-30s\t\t%.3f\t%s" % (i,coef_1, fn_1, coef_2, fn_2))
    i+=1

```

## 1.9 Bags of Words Vectorizer

```

In [15]: vectorizer = CountVectorizer()
vectorizer.fit(X_train) # fit has to happen only on train data

# we use the fitted CountVectorizer to convert the text to vector
X_train_bow = vectorizer.transform(X_train)
X_train_bow=preprocessing.normalize(X_train_bow)

X_test_bow = vectorizer.transform(X_test)
X_test_bow=preprocessing.normalize(X_test_bow)

print("Shape of Train and Test Data After vectorizations")
print(X_train_bow.shape, Y_train.shape)
print(X_test_bow.shape, Y_test.shape)

```

Shape of Train and Test Data After vectorizations  
(100500, 37994) (100500,)

(49500, 37994) (49500,)

## 1.10 Part 1 : Taking L1 as a Regularisation Parameter

### 1.10.1 Finding the best value Of hyperparameter (C or 1/Lambda)

```
In [18]: gsv=Grid_SearchCV(X_train_bow,Y_train,'l1')

        print("Best HyperParameter: ",gsv.best_params_)
        print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 10 folds for each of 17 candidates, totalling 170 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 23.5min finished
```

```
Best HyperParameter: {'C': 5, 'penalty': 'l1'}
Best Accuracy: 93.25%
```

### 1.10.2 Error-Plot

```
In [45]: plot(gsv)
```




### 1.10.3 Training the model

```
In [25]: Best_Param=gsv.best_params_  
         C=Best_Param['C']  
         Penalty = Best_Param['penalty']  
  
         Model_Bow=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')  
         Model_Bow.fit(X_train_bow,Y_train)  
  
Out[25]: LogisticRegression(C=5, class_weight=None, dual=False, fit_intercept=True,  
                             intercept_scaling=1, max_iter=100, multi_class='warn',  
                             n_jobs=None, penalty='l1', random_state=None, solver='warn',  
                             tol=0.0001, verbose=0, warm_start=False)
```

### 1.10.4 Evaluating the performance of model

```
In [26]: trainconfusionmatrix(Model_Bow,X_train_bow,Y_train)
```


Confusion Matrix for Train set




```
In [27]: testconfusionmatrix(Model_Bow,X_test_bow,Y_test)
```

Confusion Matrix for Test set





In [28]: `plot_auc_roc(Model_Bow,X_train_bow,X_test_bow,Y_train,Y_test)`



```
In [29]: print("Classification Report: \n")
         y_pred=Model_Bow.predict(X_test_bow)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.81	0.68	0.74	8692
1	0.93	0.97	0.95	40808
micro avg	0.92	0.92	0.92	49500
macro avg	0.87	0.82	0.85	49500
weighted avg	0.91	0.92	0.91	49500

### 1.10.5 Displaying 30 most informative features

```
In [30]: show_30_informative_feature(vectorizer,Model_Bow)
```

S.N	Positive	Negative	
1.	10.751	addict	-15.221
2.	14.850	bravo	-13.469
3.	11.311	cujo	-12.922
4.	10.830	delici	-21.729
5.	16.410	downsid	-15.091
6.	12.746	fascin	-12.848
7.	14.110	filbert	-13.724
8.	12.056	finest	-16.663
9.	12.521	goshoptnt	-13.234
10.	14.539	gripe	-14.645
11.	12.249	heal	-19.385
12.	12.816	hook	-13.507
13.	11.312	morsel	-12.876
14.	13.046	narrow	-12.526
15.	14.304	nevertheless	-17.224
16.	11.407	penal	-12.653
17.	12.362	rater	-12.851
18.	10.762	reassur	-12.464
19.	11.579	refresh	-12.738
20.	14.879	skeptic	-12.960
21.	10.681	solv	-20.452
22.	11.757	sooth	-14.868

23.	14.262	spectacular	-12.400
24.	10.961	steal	-14.206
25.	11.518	stumbl	-13.515
26.	13.801	tastey	-13.100
27.	12.057	versatil	-19.808
28.	12.000	whim	-12.651
29.	11.250	yay	-17.415
30.	12.017	yum	-18.358

## 1.11 Part 2 : Taking L2 as a Regularisation Parameter

### 1.11.1 Finding the best value Of hyperparameter (C or 1/Lambda)

In [46]: `gsv=Grid_SearchCV(X_train_bow,Y_train,'l2')`

```
print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 10 folds for each of 17 candidates, totalling 170 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 18.2min finished
```

```
Best HyperParameter: {'C': 5, 'penalty': 'l2'}
Best Accuracy: 93.73%
```

### 1.11.2 Error Plot

In [47]: `plot(gsv)`



### 1.11.3 Training the model

```
In [49]: Best_Param=gsv.best_params_
         C=Best_Param['C']
         Penalty = Best_Param['penalty']


         Model_Bow=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')
         Model_Bow.fit(X_train_bow,Y_train)

Out[49]: LogisticRegression(C=5, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

### 1.11.4 Evaluating the performance of model


```
In [50]: trainconfusionmatrix(Model_Bow,X_train_bow,Y_train)
```

Confusion Matrix for Train set




```
In [51]: testconfusionmatrix(Model_Bow,X_test_bow,Y_test)
```

Confusion Matrix for Test set



In [52]: `plot_auc_roc(Model_Bow,X_train_bow,X_test_bow,Y_train,Y_test)`



```
In [53]: print("Classification Report: \n")
         y_pred=Model_Bow.predict(X_test_bow)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.61	0.87	0.72	8692
1	0.97	0.88	0.92	40808
micro avg	0.88	0.88	0.88	49500
macro avg	0.79	0.88	0.82	49500
weighted avg	0.91	0.88	0.89	49500

```
In [54]: show_30_informative_feature(vectorizer,Model_Bow)
```

S.N	Positive		Negative	
1.	9.411	addict	-11.813	aw
2.	9.014	amaz	-8.823	blar
3.	7.990	awesom	-8.392	can
4.	8.843	beat	-8.716	con
5.	6.319	beauti	-10.212	dis
6.	7.859	best	-9.389	dis
7.	7.196	complaint	-6.951	ear
8.	9.800	delici	-7.319	fai
9.	7.950	downsid	-7.624	fla
10.	5.720	drawback	-7.934	gr
11.	8.252	excel	-9.971	hor
12.	6.118	fabul	-7.615	in
13.	7.002	fantast	-7.501	me
14.	6.565	glad	-6.946	po
15.	6.892	great	-7.751	re
16.	5.885	habit	-8.956	re
17.	6.093	happier	-8.352	ru
18.	6.132	heaven	-7.053	sha
19.	9.308	hook	-7.148	sta
20.	8.343	perfect	-9.299	ta
21.	7.814	refresh	-11.563	te
22.	7.599	satisfi	-9.097	th
23.	8.586	skeptic	-7.799	un
24.	6.738	smooth	-7.521	un

25.	6.210	sooth	-8.705	un
26.	6.343	terrif	-9.674	unp
27.	6.375	uniqu	-7.638	wea
28.	7.160	worri	-8.367	wor
29.	8.218	yum	-15.006	w
30.	8.129	yummi	-7.478	yu

## 1.12 TF-IDF Vectorizer

```
In [30]: vectorizer_tfidf=TfidfVectorizer()
         vectorizer_tfidf.fit(X_train)
```

```
Out[30]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
                        dtype=<class 'numpy.float64'>, encoding='utf-8', input='content',
                        lowercase=True, max_df=1.0, max_features=None, min_df=1,
                        ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_idf=True,
                        stop_words=None, strip_accents=None, sublinear_tf=False,
                        token_pattern='(?u)\\b\\w\\w+\\b', tokenizer=None, use_idf=True,
                        vocabulary=None)
```

```
In [31]: X_Train_Tfidf=vectorizer_tfidf.transform(X_train)
         X_Train_Tfidf=preprocessing.normalize(X_Train_Tfidf)
```

```
X_Test_Tfidf=vectorizer_tfidf.transform(X_test)
X_Test_Tfidf=preprocessing.normalize(X_Test_Tfidf)
```

```
In [32]: print("Shape of Train and Test Data After vectorizations")
         print(X_Train_Tfidf.shape, Y_train.shape)
         print(X_Test_Tfidf.shape, Y_test.shape)
```

```
Shape of Train and Test Data After vectorizations
(100500, 37994) (100500,)
(49500, 37994) (49500,)
```

## 1.13 Taking L1 as a Regularisation Parameter

### 1.13.1 Finding the best value Of hyperparameter (C or 1/Lambda)

```
In [21]: gsv=Grid_SearchCV(X_Train_Tfidf,Y_train,'l1')

         print("Best HyperParameter: ",gsv.best_params_)
         print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Fitting 10 folds for each of 17 candidates, totalling 170 fits


```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 163.2min finished
```



Best HyperParameter: {'C': 1, 'penalty': 'l1'}  
Best Accuracy: 93.30%

### 1.13.2 Error Plot

```
In [22]: plot(gsv)
```




### 1.13.3 Training the model

```
In [23]: Best_Param=gsv.best_params_  
         C=Best_Param['C']  
         Penalty = Best_Param['penalty']  
  
         Model_Tfidf=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')  
         Model_Tfidf.fit(X_train_bow,Y_train)  
  
Out[23]: LogisticRegression(C=1, class_weight='balanced', dual=False,  
                             fit_intercept=True, intercept_scaling=1, max_iter=100,  
                             multi_class='warn', n_jobs=None, penalty='l1', random_state=None,  
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

#### 1.13.4 Evaluating the performance of model


```
In [24]: trainconfusionmatrix(Model_Tfidf,X_Train_Tfidf,Y_train)
```

Confusion Matrix for Train set




```
In [25]: testconfusionmatrix(Model_Tfidf,X_Test_Tfidf,Y_test)
```

Confusion Matrix for Test set



In [26]: `plot_auc_roc(Model_Tfidf,X_Train_Tfidf,X_Test_Tfidf,Y_train,Y_test)`



```
In [27]: print("Classification Report: \n")
        y_pred=Model_Tfidf.predict(X_Test_Tfidf)

        print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.51	0.92	0.66	8692
1	0.98	0.81	0.89	40808
micro avg	0.83	0.83	0.83	49500
macro avg	0.75	0.87	0.77	49500
weighted avg	0.90	0.83	0.85	49500

### 1.13.5 Displaying 30 most informative features

```
In [28]: show_30_informative_feature(vectorizer_tfidf,Model_Tfidf)
```

S.N	Positive		Negative	
1.	10.139	addict	-13.064	av
2.	9.406	amaz	-9.133	blan
3.	8.423	awesom	-10.132	can
4.	10.021	beat	-10.651	co
5.	6.980	beauti	-7.857	dece
6.	7.909	best	-9.879	dece
7.	7.294	complaint	-10.159	dis
8.	9.996	delici	-11.139	dis
9.	11.857	downsid	-7.996	dr
10.	8.466	excel	-8.134	fa
11.	6.967	fabul	-10.004	f
12.	7.052	fantast	-8.553	gro
13.	6.808	glad	-10.879	h
14.	6.894	great	-10.254	in
15.	7.186	habit	-10.983	me
16.	7.372	happier	-8.612	opp
17.	6.660	heaven	-8.446	re
18.	11.047	hook	-9.221	re
19.	8.251	perfect	-9.101	ru
20.	8.721	refresh	-10.548	ta
21.	7.890	satisfi	-12.189	te
22.	12.359	skeptic	-9.873	tl

23.	6.920	smooth	-10.156	tl
24.	8.476	sooth	-9.988	una
25.	7.458	uniqu	-12.696	un
26.	7.559	versatil	-9.074	un.
27.	7.442	whim	-11.399	un
28.	7.932	worri	-9.091	wor
29.	9.618	yum	-17.436	w
30.	8.364	yummi	-8.111	yu

## 1.14 Taking L2 as a Regularisation Parameter

### 1.14.1 Finding the best value Of hyperparameter (C or 1/Lambda)

In [29]: `gsv=Grid_SearchCV(X_Train_Tfidf,Y_train,'l2')`

```
print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 10 folds for each of 17 candidates, totalling 170 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 15.0min finished
```

```
Best HyperParameter: {'C': 1, 'penalty': 'l2'}
Best Accuracy: 93.68%
```

### 1.14.2 Error Plot

In [30]: `plot(gsv)`



### 1.14.3 Training the mode

```
In [31]: Best_Param=gsv.best_params_
         C=Best_Param['C']
         Penalty = Best_Param['penalty']


         Model_Tfidf=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')
         Model_Tfidf.fit(X_train_bow,Y_train)

Out[31]: LogisticRegression(C=1, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

### 1.14.4 Evaluating the performance of model


```
In [32]: trainconfusionmatrix(Model_Tfidf,X_Train_Tfidf,Y_train)
```

Confusion Matrix for Train set




```
In [33]: testconfusionmatrix(Model_Tfidf,X_Test_Tfidf,Y_test)
```

Confusion Matrix for Test set



In [34]: `plot_auc_roc(Model_Tfidf,X_Train_Tfidf,X_Test_Tfidf,Y_train,Y_test)`





```
In [35]: print("Classification Report: \n")
         y_pred=Model_Tfidf.predict(X_Test_Tfidf)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.51	0.93	0.66	8692
1	0.98	0.81	0.89	40808
micro avg	0.83	0.83	0.83	49500
macro avg	0.75	0.87	0.77	49500
weighted avg	0.90	0.83	0.85	49500

```
In [36]: show_30_informative_feature(vectorizer_tfidf,Model_Tfidf)
```

S.N	Positive		Negative	
1.	6.398	addict	-8.666	aw
2.	7.030	amaz	-4.353	bar
3.	5.944	awesom	-4.465	bew
4.	5.792	beat	-6.507	blan
5.	4.120	beauti	-4.733	cand
6.	7.024	best	-9.024	disa
7.	4.659	complaint	-6.022	disp
8.	8.489	delici	-4.433	ear
9.	4.560	easi	-5.182	gros
10.	7.182	excel	-7.371	hor
11.	5.105	fantast	-4.307	mor
12.	4.673	favorit	-5.133	po
13.	5.131	glad	-5.292	re
14.	6.400	great	-7.271	re
15.	4.215	happi	-5.036	ru
16.	5.426	hook	-4.640	sac
17.	4.936	love	-4.423	sha
18.	4.570	nice	-4.696	so
19.	7.153	perfect	-5.470	sta
20.	4.929	refresh	-4.798	st
21.	5.876	satisfi	-5.644	ta
22.	5.373	smooth	-8.374	te
23.	4.051	tasti	-5.834	th
24.	4.327	thank	-7.005	un

25.	3.958	uniqu	-5.636	un
26.	3.879	without	-4.787	was
27.	4.456	wonder	-5.962	wea
28.	4.659	worri	-5.332	wo
29.	4.800	yum	-10.212	w
30.	6.071	yummi	-4.410	yu

## 1.15 Word To Vector

```
In [15]: list_of_Train_sent=[]
        list_of_Test_sent=[]

        for sent in X_train:
            list_of_Train_sent.append(sent.split())

        for sent in X_test:
            list_of_Test_sent.append(sent.split())

In [16]: model=Word2Vec(list_of_Train_sent,min_count=5,size=50, workers=4)
```

## 1.16 Average Word To Vector

```
In [17]: import numpy as np

        Train_vectors = []
        for sent in list_of_Train_sent:
            sent_vec = np.zeros(50)
            cnt_words = 0
            for word in sent:
                try:
                    vec = model.wv[word]
                    sent_vec += vec
                    cnt_words += 1
                except:
                    pass
            if cnt_words!=0:
                sent_vec /= cnt_words
            Train_vectors.append(sent_vec)
        Train_vectors = np.nan_to_num(Train_vectors)

In [18]: import numpy as np

        Test_vectors=[]
        for sent in list_of_Test_sent:
            sent_vec=np.zeros(50)
            cnt_words=0
            for word in sent:
```

```

        try:
            vec=model.wv[word]
            sent_vec+=vec
            cnt_words+=1
        except:
            pass
    if cnt_words!=0:
        sent_vec/=cnt_words
    Test_vectors.append(sent_vec)
Test_vectors=np.nan_to_num(Test_vectors)

```

```
In [19]: print("Shape of Test Vectors : ",Test_vectors.shape)
```

```
Shape of Test Vectors :  (49500, 50)
```

```
In [20]: X_Train_Awv=Train_vectors
        X_Test_Awv=Test_vectors
```

```
In [21]: print(X_Train_Awv.shape, Y_train.shape)
        print(X_Test_Awv.shape, Y_test.shape)
```

```
(100500, 50) (100500,)
(49500, 50) (49500,)
```

## 1.17 Taking L1 as a Regularisation Parameter

### 1.17.1 Finding the best value Of hyperparameter (C or 1/Lambda)

```
In [23]: gsv=Grid_SearchCV(X_Train_Awv,Y_train,"l1")

        print("Best HyperParameter: ",gsv.best_params_)
        print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 10 folds for each of 17 candidates, totalling 170 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 52.6min finished
```

```
Best HyperParameter:  {'C': 1, 'penalty': 'l1'}
Best Accuracy: 90.46%
```

### 1.17.2 Error Plot

```
In [24]: plot(gsv)
```



### 1.17.3 Training Model

```
In [25]: Best_Param=gsv.best_params_
         C=Best_Param['C']
         Penalty = Best_Param['penalty']


         Model_Awv=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')
         Model_Awv.fit(X_Train_Awv,Y_train)

Out[25]: LogisticRegression(C=1, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l1', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

### 1.17.4 Evaluating the performance of model


```
In [26]: trainconfusionmatrix(Model_Awv,X_Train_Awv,Y_train)
```

Confusion Matrix for Train set




```
In [27]: trainconfusionmatrix(Model_Awv,X_Test_Awv,Y_test)
```

Confusion Matrix for Train set



In [28]: `plot_auc_roc(Model_Awv,X_Train_Awv,X_Test_Awv,Y_train,Y_test)`



```
In [29]: print("Classification Report: \n")
         y_pred=Model_Awv.predict(X_Test_Awv)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.48	0.85	0.61	8692
1	0.96	0.80	0.87	40808
micro avg	0.81	0.81	0.81	49500
macro avg	0.72	0.83	0.74	49500
weighted avg	0.88	0.81	0.83	49500

## 1.18 Taking L2 as a Regularisation Parameter

### 1.18.1 Finding the best value Of hyperparameter (C or 1/Lambda)

```
In [22]: gsv=Grid_SearchCV(X_Train_Awv,Y_train,"l2")

         print("Best HyperParameter: ",gsv.best_params_)
         print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

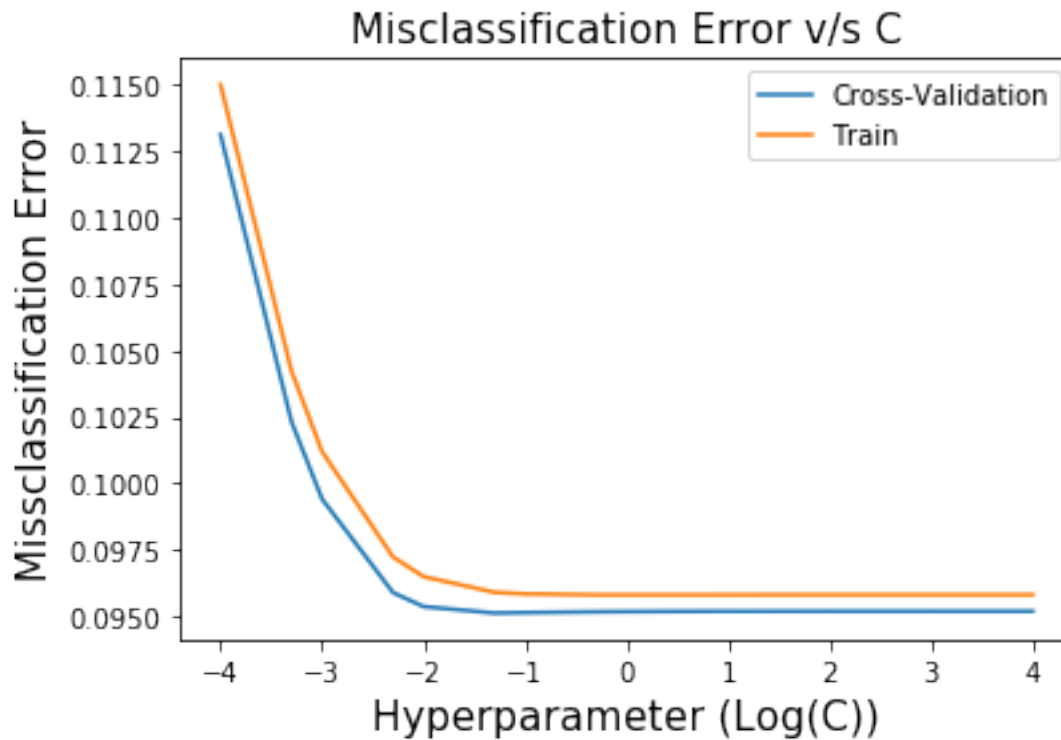
Fitting 10 folds for each of 17 candidates, totalling 170 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 3.2min finished
```

```
Best HyperParameter: {'C': 0.05, 'penalty': 'l2'}
Best Accuracy: 90.49%
```

### 1.18.2 Error Plot

```
In [23]: plot(gsv)
```



### 1.18.3 Training the model

```
In [24]: Best_Param=gsv.best_params_
         C=Best_Param['C']
         Penalty = Best_Param['penalty']

         Model_Awv=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')
         Model_Awv.fit(X_Train_Awv,Y_train)


Out[24]: LogisticRegression(C=0.05, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l2', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

### 1.18.4 Evaluating the performance of model

```
In [25]: trainconfusionmatrix(Model_Awv,X_Train_Awv,Y_train)
```


Confusion Matrix for Train set






```
In [26]: trainconfusionmatrix(Model_Awv,X_Test_Awv,Y_test)
```

Confusion Matrix for Train set



```
In [27]: plot_auc_roc(Model_Awv,X_Train_Awv,X_Test_Awv,Y_train,Y_test)
```



```
In [28]: print("Classification Report: \n")
        y_pred=Model_Awv.predict(X_Test_Awv)

        print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.48	0.85	0.61	8692
1	0.96	0.80	0.88	40808
micro avg	0.81	0.81	0.81	49500
macro avg	0.72	0.83	0.74	49500
weighted avg	0.88	0.81	0.83	49500

## 1.19 TF-IDF Word To Vector

```
In [33]: TFIDF_Feature=vectorizer_tfidf.get_feature_names()
        print(len(TFIDF_Feature))
        print(TFIDF_Feature[0:20])
```

37994

['aaa', 'aaa', 'aaaaaaaaaaaaaaaaaaaaaargh', 'aaaaaaaaaag

```
In [34]: from tqdm import tqdm
        Train_TFIDF_W2V_Vectors=[]
        row=0
        for sent in tqdm(list_of_Train_sent):
            sent_vec=np.zeros(50)
            weight=0
            for word in sent:
                try :
                    w2v_vec=model.wv[word]
                    tfidf_vec=X_Train_Tfidf[row,TFIDF_Feature.index(word)]
                    sent_vec+=(w2v_vec*tfidf_vec)
                    weight+=tfidf_vec

                except :
                    pass
            if weight!=0:
                sent_vec/=weight
            Train_TFIDF_W2V_Vectors.append(sent_vec)
            row+=1
```

100%|| 100500/100500 [1:04:56<00:00, 25.79it/s]

```
In [35]: Test_TFIDF_W2V_Vectors=[]
row=0
for sent in tqdm(list_of_Test_sent):
    sent_vec=np.zeros(50)
    weight=0

    for word in sent:
        try:
            w2v_vec=model.wv[word]
            tfidf_vec=X_Test_Tfidf(row,TFIDF_Feature.index(word))
            sent_vec+=(w2v_vec*tfidf_vec)
            weight+=tfidf

        except :
            pass

    if weight!=0:
        sent_vec/=weight
    Test_TFIDF_W2V_Vectors.append(sent_vec)
    row+=1
```

100%|| 49500/49500 [24:01<00:00, 34.33it/s]

```
In [36]: Train_TFIDF_W2V_Vectors = np.nan_to_num(Train_TFIDF_W2V_Vectors)
Test_TFIDF_W2V_Vectors = np.nan_to_num(Test_TFIDF_W2V_Vectors)
```

```
In [37]: X_Train_TfidfW2v=Train_TFIDF_W2V_Vectors
X_Test_TfidfW2v=Test_TFIDF_W2V_Vectors
```

```
In [38]: print(X_Train_TfidfW2v.shape, Y_train.shape)
print(X_Test_TfidfW2v.shape, Y_test.shape)
```

```
(100500, 50) (100500,)
(49500, 50) (49500,)
```

## 1.20 Taking L1 as a Regularisation Parameter

### 1.20.1 Finding the best value Of hyperparameter (C or 1/Lambda)

```
In [39]: gsv=Grid_SearchCV(X_Train_TfidfW2v,Y_train,"l1")

print("Best HyperParameter: ",gsv.best_params_)
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

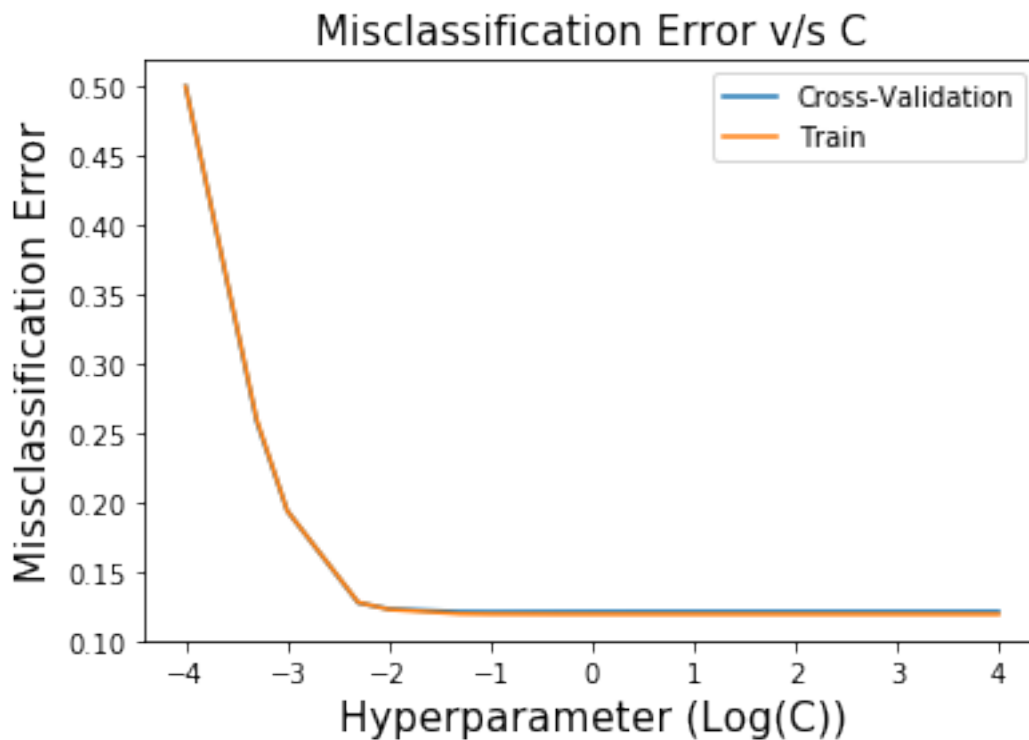
Fitting 10 folds for each of 17 candidates, totalling 170 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 26.1min finished
```

```
Best HyperParameter: {'C': 0.1, 'penalty': 'l1'}  
Best Accuracy: 87.85%
```

### 1.20.2 Error Plot

```
In [40]: plot(gsv)
```




### 1.20.3 Training Model

```
In [41]: Best_Param=gsv.best_params_  
         C=Best_Param['C']  
         Penalty = Best_Param['penalty']  
  
         Model_TfidfW2v=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')  
         Model_TfidfW2v.fit(X_Train_TfidfW2v,Y_train)  
  
Out[41]: LogisticRegression(C=0.1, class_weight='balanced', dual=False,  
                             fit_intercept=True, intercept_scaling=1, max_iter=100,  
                             multi_class='warn', n_jobs=None, penalty='l1', random_state=None,  
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

#### 1.20.4 Evaluating the performance of model


```
In [42]: trainconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Train set




```
In [43]: testconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Test set



In [44]: `plot_auc_roc(Model_TfidfW2v,X_Train_TfidfW2v,X_Test_TfidfW2v,Y_train,Y_test)`



```
In [45]: print("Classification Report: \n")
         y_pred=Model_TfidfW2v.predict(X_Test_TfidfW2v)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.18	1.00	0.30	8692
1	0.00	0.00	0.00	40808
micro avg	0.18	0.18	0.18	49500
macro avg	0.09	0.50	0.15	49500
weighted avg	0.03	0.18	0.05	49500

## 1.21 Taking L2 as a Regularisation Parameter

### 1.21.1 Finding the best value Of hyperparameter (C or 1/Lambda)

```
In [46]: gsv=Grid_SearchCV(X_Train_TfidfW2v,Y_train,"l1")

         print("Best HyperParameter: ",gsv.best_params_)
         print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Fitting 10 folds for each of 17 candidates, totalling 170 fits


```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 27.0min finished
```

```
Best HyperParameter: {'C': 0.1, 'penalty': 'l1'}
Best Accuracy: 87.85%
```

### 1.21.2 Error Plot

```
In [47]: plot(gsv)
```





### 1.21.3 Training the model

```
In [48]: Best_Param=gsv.best_params_
         C=Best_Param['C']
         Penalty = Best_Param['penalty']


         Model_TfidfW2v=LogisticRegression(C=C,penalty=Penalty,class_weight='balanced')
         Model_TfidfW2v.fit(X_Train_TfidfW2v,Y_train)
```

```
Out[48]: LogisticRegression(C=0.1, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l1', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

### 1.21.4 Evaluating the performance of model


```
In [49]: trainconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Train set




```
In [50]: testconfusionmatrix(Model_TfidfW2v,X_Train_TfidfW2v,Y_train)
```

Confusion Matrix for Test set



In [51]: `plot_auc_roc(Model_TfidfW2v,X_Train_TfidfW2v,X_Test_TfidfW2v,Y_train,Y_test)`



```
In [52]: print("Classification Report: \n")
        y_pred=Model_TfidfW2v.predict(X_Test_TfidfW2v)

        print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.18	1.00	0.30	8692
1	0.00	0.00	0.00	40808
micro avg	0.18	0.18	0.18	49500
macro avg	0.09	0.50	0.15	49500
weighted avg	0.03	0.18	0.05	49500

## 1.22 Perturbation Test On BOW Vectorizer and Regularisation Parameter L2

```
In [62]: from scipy.sparse import find
        #Weights before adding random noise
        weights1 = find(Model_Bow.coef_[0])[2]
        print(weights1[:50])
```

```
[ 1.13345591  0.02002588  0.00200432  0.03735237  0.00537296  0.00278505
 -0.05079941  0.00931602  0.00537296  0.13805403  0.03798705  0.0069488
  0.03927589  0.01479231  0.02671619  0.0565606  -0.40902605  0.03473881
  0.00330004  0.07520482 -0.05522391  0.06312428  0.02666827  0.032443
  0.04260278 -0.03394333  0.04648075  0.00561827  0.28007135  0.18782385
 -0.34749559  0.05698905  0.00788936 -1.170488  0.52939662  0.20041705
 -0.07370247 -0.23254724  0.23001907  0.09660233  0.3921331  0.35992047
 -0.27294236  0.17301799  0.80156947  0.07930274 -0.76391136  0.07604107
  0.0629005  0.02145765]
```

```
In [63]: X_train_t = X_train_bow
        #Random noise
        epsilon = np.random.uniform(low=-0.0001, high=0.0001, size=(find(X_train_t)[0].size,))
        #Getting the postions(row and column) and value of non-zero datapoints
        a,b,c = find(X_train_t)

        #Introducing random noise to non-zero datapoints
        X_train_t[a,b] = epsilon + X_train_t[a,b]
```

```
In [64]: Model_Pert= LogisticRegression(C=C, penalty= 'l2',class_weight='balanced')
        Model_Pert.fit(X_train_t,Y_train)
```

```

y_pred = Model_Pert.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("Non Zero weights:", np.count_nonzero(Model_Pert.coef_))

```

Accuracy on test set: 87.933%  
Non Zero weights: 37994

```

In [65]: #Weights after adding random noise
weights2 = find(Model_Pert.coef_[0])[2]
print(weights2[:50])

```

```

[ 1.13328701  0.02007606  0.00200724  0.03734596  0.00538038  0.00278589
 -0.05092452  0.00931236  0.00538486  0.13790361  0.03810089  0.00694868
  0.03934439  0.0148218   0.02677143  0.0565591  -0.40895739  0.03471803
  0.00329922  0.07520774 -0.05585209  0.06320037  0.02665189  0.03241966
  0.04256979 -0.0338506   0.04647984  0.00561043  0.28022148  0.18792607
 -0.34781192  0.05681231  0.00787586 -1.17053264  0.52951101  0.20055738
 -0.07368803 -0.23281588  0.22997003  0.09659027  0.39222921  0.35980129
 -0.27276686  0.1729645   0.80172465  0.07946077 -0.76348883  0.07607611
  0.06293845  0.02146781]

```

```

In [66]: weights_diff = (abs(weights1 - weights2)/weights1) * 100

```

```

In [67]: print(weights_diff[np.where(weights_diff > 30)].size)

```

3

### 1.22.1 Showing How Sparsity increases as we decrease C or increase Lambada(1/C) When L1 Regularisation is used

```

In [55]: Model_Sparse= LogisticRegression(C= 1000, penalty= 'l1',class_weight='balanced')
Model_Sparse.fit(X_train_bow,Y_train)
y_pred = Model_Sparse.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("F1-Score on test set: %0.3f"%(f1_score(Y_test, y_pred)))
print("Non Zero weights:", np.count_nonzero(Model_Sparse.coef_))

```

Accuracy on test set: 85.586%  
F1-Score on test set: 0.910  
Non Zero weights: 18732

```

In [56]: Model_Sparse= LogisticRegression(C= 100, penalty= 'l1',class_weight='balanced')
Model_Sparse.fit(X_train_bow,Y_train)
y_pred = Model_Sparse.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("F1-Score on test set: %0.3f"%(f1_score(Y_test, y_pred)))
print("Non Zero weights:", np.count_nonzero(Model_Sparse.coef_))

```

Accuracy on test set: 86.901%  
F1-Score on test set: 0.918  
Non Zero weights: 14916

```
In [57]: Model_Sparse= LogisticRegression(C= 10, penalty= 'l1',class_weight='balanced')
Model_Sparse.fit(X_train_bow,Y_train)
y_pred = Model_Sparse.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("F1-Score on test set: %0.3f"%(f1_score(Y_test, y_pred)))
print("Non Zero weights:",np.count_nonzero(Model_Sparse.coef_))
```

Accuracy on test set: 88.222%  
F1-Score on test set: 0.926  
Non Zero weights: 7780

```
In [58]: Model_Sparse= LogisticRegression(C= 1, penalty= 'l1',class_weight='balanced')
Model_Sparse.fit(X_train_bow,Y_train)
y_pred = Model_Sparse.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("F1-Score on test set: %0.3f"%(f1_score(Y_test, y_pred)))
print("Non Zero weights:",np.count_nonzero(Model_Sparse.coef_))
```

Accuracy on test set: 87.299%  
F1-Score on test set: 0.919  
Non Zero weights: 1931

```
In [59]: Model_Sparse= LogisticRegression(C= 0.1, penalty= 'l1',class_weight='balanced')
Model_Sparse.fit(X_train_bow,Y_train)
y_pred = Model_Sparse.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("F1-Score on test set: %0.3f"%(f1_score(Y_test, y_pred)))
print("Non Zero weights:",np.count_nonzero(Model_Sparse.coef_))
```

Accuracy on test set: 83.978%  
F1-Score on test set: 0.896  
Non Zero weights: 409

```
In [60]: Model_Sparse= LogisticRegression(C= 0.01, penalty= 'l1')
Model_Sparse.fit(X_train_bow,Y_train)
y_pred = Model_Sparse.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("F1-Score on test set: %0.3f"%(f1_score(Y_test, y_pred)))
print("Non Zero weights:",np.count_nonzero(Model_Sparse.coef_))
```

Accuracy on test set: 82.828%  
F1-Score on test set: 0.906

Non Zero weights: 21

```
In [61]: Model_Sparse= LogisticRegression(C= 0.001, penalty= 'l1',class_weight='balanced')
Model_Sparse.fit(X_train_bow,Y_train)
y_pred = Model_Sparse.predict(X_test_bow)
print("Accuracy on test set: %0.3f%%"%(accuracy_score(Y_test, y_pred)*100))
print("F1-Score on test set: %0.3f"%(f1_score(Y_test, y_pred)))
print("Non Zero weights:",np.count_nonzero(Model_Sparse.coef_))
```

Accuracy on test set: 38.505%

F1-Score on test set: 0.423

Non Zero weights: 1

1. Decrease in value of C decrease the non of non zeros weight decrease , which means it is increasing the sparsity(No of Zeros)

### 1.22.2 Addition of another column length

```
In [18]: Train_len=[]
Test_len=[]
for i in X_train:
    Train_len.append(len(i))

for i in X_test:
    Test_len.append(len(i))

In [19]: Train_len=np.array(Train_len)
Test_len=np.array(Test_len)

In [20]: Train_len=Train_len[:,np.newaxis]
Test_len=Test_len[:,np.newaxis]
```

### Bag Of Word Vectorizer

```
In [21]: X_Train_BOW=X_train_bow.todense()

In [22]: X_Train_New=np.append(X_Train_BOW,Train_len,axis=1)

In [23]: from scipy.sparse import csr_matrix
X_Train_New= csr_matrix(X_Train_New)

In [24]: print("Shape of Train Data Before Adding length column ")
print(X_train_bow.shape)

print("\nShape of Train Data After Adding length column ")
print(X_Train_New.shape)
```

Shape of Train Data Before Adding length column  
(100500, 37994)

Shape of Train Data After Adding length column  
(100500, 37995)

### Finding the best value Of hyperparameter (C or 1/lambda)

```
In [29]: gsv=Grid_SearchCV(X_Train_New,Y_train,"l1")
```

```
print("Best HyperParameter: ",gsv.best_params_)  
print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```


Fitting 10 folds for each of 17 candidates, totalling 170 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done 170 out of 170 | elapsed: 110.4min finished
```

Best HyperParameter: {'C': 1, 'penalty': 'l1'}  
Best Accuracy: 93.26%

### Error Plot

```
In [30]: plot(gsv)
```





## Training the Model

```
In [25]: X_Test_Bow=X_test_bow.todense()

In [26]: X_Test_New=np.append(X_Test_Bow,Test_len,axis=1)

In [27]: from scipy.sparse import csr_matrix
         X_Test_New= csr_matrix(X_Test_New)

In [31]: Best_Param=gsv.best_params_
         C=Best_Param['C']
         Penalty = Best_Param['penalty']


In [32]: Model_LBow = LogisticRegression(C=C,penalty = Penalty,class_weight='balanced')
         Model_LBow.fit(X_Train_New,Y_train)

Out[32]: LogisticRegression(C=1, class_weight='balanced', dual=False,
                             fit_intercept=True, intercept_scaling=1, max_iter=100,
                             multi_class='warn', n_jobs=None, penalty='l1', random_state=None,
                             solver='warn', tol=0.0001, verbose=0, warm_start=False)
```

## Evaluating the performance of the model


```
In [33]: trainconfusionmatrix(Model_LBow,X_Train_New,Y_train)
```

Confusion Matrix for Train set




```
In [34]: testconfusionmatrix(Model_LBow,X_Test_New,Y_test)
```

Confusion Matrix for Test set



```
In [35]: plot_auc_roc(Model_LBow,X_Train_New,X_Test_New,Y_train,Y_test)
```



```
In [36]: print("Classification Report: \n")
         y_pred=Model_LBow.predict(X_Test_New)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.59	0.88	0.71	8692
1	0.97	0.87	0.92	40808
micro avg	0.87	0.87	0.87	49500
macro avg	0.78	0.88	0.81	49500
weighted avg	0.90	0.87	0.88	49500

### 1.22.3 Displaying 30 most informative features

```
In [37]: show_30_informative_feature(vectorizer,Model_LBow)
```

S.N	Positive	Negative
-----	----------	----------

1.	10.165	addict	-13.051	av
2.	9.406	amaz	-9.123	blan
3.	8.438	awesom	-10.127	car
4.	10.038	beat	-10.607	co
5.	7.003	beauti	-7.860	dece
6.	7.925	best	-9.884	dece
7.	7.292	complaint	-10.151	dis
8.	10.010	delici	-11.124	d.
9.	11.888	downsid	-7.976	dr
10.	8.479	excel	-8.112	fa
11.	6.976	fabul	-10.002	fi
12.	7.065	fantast	-8.534	gr
13.	6.811	glad	-10.867	he
14.	6.903	great	-10.231	in
15.	7.235	habit	-10.951	me
16.	7.399	happier	-8.610	opp
17.	6.696	heaven	-8.455	re
18.	11.030	hook	-9.208	re
19.	8.266	perfect	-9.089	ru
20.	8.723	refresh	-10.560	ta
21.	7.902	satisfi	-12.177	ta
22.	12.386	skeptic	-9.787	tl
23.	6.934	smooth	-10.159	tl
24.	8.470	sooth	-9.990	una
25.	7.506	uniqu	-12.711	un
26.	7.592	versatil	-9.064	un
27.	7.450	whim	-11.392	un
28.	7.948	worri	-9.066	wor
29.	9.642	yum	-17.412	w
30.	8.376	yummi	-8.098	yuc

## 1.Report On Different Vectorizer Method and Regularisation Parameter L1

```
In [1]: from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["Vectorizer", "Hyperparameter(C or 1/Lambda)", "Train AUC", "Test AUC", "
```

```
x.add_row(["BOW", 5, 0.96, 0.94, 0.91])
```

```
x.add_row(["TF-IDF", 1, 0.95, 0.94, 0.85])
```

```
x.add_row(["Avg W2V", 1, 0.90, 0.90, 0.83])
```

```
x.add_row(["TF-IDF W2V", 0.1, 0.88, 0.50, 0.05])
```

```
print(x)
```

```
+-----+-----+-----+-----+-----+
| Vectorizer | Hyperparameter(C or 1/Lambda) | Train AUC | Test AUC | F1-Score |
```

Vectorizer	Hyperparameter(C or 1/Lambda)	Train AUC	Test AUC	F1-Score
BOW	5	0.96	0.94	0.91
TF-IDF	1	0.95	0.94	0.85
Avg W2V	1	0.9	0.9	0.83
TF-IDF W2V	0.1	0.88	0.5	0.05

## 2.Report On Different Vectorizer Method and Regularisation Parameter L2

In [2]: `from prettytable import PrettyTable`

```
x = PrettyTable()
```

```
x.field_names = ["Vectorizer", "Hyperparameter(C or 1/Lambda)", "Train AUC", "Test AUC", "F1-Score"]
```

```
x.add_row(["BOW", 5, 0.96, 0.94, 0.89])
```

```
x.add_row(["TF-IDF", 1, 0.96, 0.94, 0.85])
```

```
x.add_row(["Avg W2V", 0.05, 0.90, 0.90, 0.83])
```

```
x.add_row(["TF-IDF W2V", 0.1, 0.88, 0.50, 0.05])
```

```
print(x)
```

Vectorizer	Hyperparameter(C or 1/Lambda)	Train AUC	Test AUC	F1-Score
BOW	5	0.96	0.94	0.89
TF-IDF	1	0.96	0.94	0.85
Avg W2V	0.05	0.9	0.9	0.83
TF-IDF W2V	0.1	0.88	0.5	0.05

## 3.Report On Different Vectorizer Method After Addition Of Length as another Column

In [3]: `from prettytable import PrettyTable`

```
x = PrettyTable()
```

```
x.field_names = ["Vectorizer", "Hyperparameter(C or 1/Lambda)", "Train AUC", "Test AUC", "F1-Score"]
```

```
x.add_row(["BOW", 1, 0.95, 0.94, 0.88])
```

```
print(x)
```

Vectorizer	Hyperparameter(C or 1/Lambda)	Train AUC	Test AUC	F1-Score
BOW	1	0.95	0.94	0.88

BOW	1	0.95	0.94	0.88
-----	---	------	------	------

4. I have taken considerable amount of data but it did not take long time in execution .
5. Since data is unbalanced , i did time based splitting and used roc\_auc metric as scoring parameter in GridsearchCV .
6. After adding Length as another column , there is no any improvement.
7. TF-IDF W2V is not performing well .