

Assignment-4 - Apply-Naive Bayes-On-Amazon-Review-Dataset

March 30, 2019

1 Assignment-4: Apply Naive Bayes On Amazon Fine Food Reviews DataSet

1.1 Introduction

(i).Naive Bayes is a simple, yet effective and commonly-used, machine learning classifier. It is a probabilistic classifier that makes classifications using the Maximum A Posteriori decision rule in a Bayesian setting

(ii).Naive Bayes classifiers have been especially popular for text classification, and are a traditional solution for problems such as spam detection.

1.2 Objective

To Predict the Polarity of Amazon Fine Food Review Using Naive Bayes Algorithm.

1.3 Importing All Required Library

```
In [1]: %matplotlib inline
import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
import math

from sklearn.model_selection import GridSearchCV
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB

from sklearn.metrics import classification_report
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
```

```

from sklearn.model_selection import TimeSeriesSplit

from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import cross_val_score
from sklearn import preprocessing

import pickle

from tqdm import tqdm
import os
import warnings
warnings.filterwarnings("ignore")

```

1.4 Importing Amazon Fine Food Review Dataset

```

In [2]: if os.path.isfile("final.sqlite"):
        conn=sqlite3.connect("final.sqlite")
        Data=pd.read_sql_query("select * from NaiveBayes where Score!=3",conn)
        conn.close()
    else :
        print("Error Importing the file")

```

1.5 Taking 150K Random Points

```

In [ ]: Data=Data.sample(n=150000)

In [3]: # Printing some data of DataFrame

        Data['Score'].value_counts()

```

```

Out[3]: 1    126439
        0     23561
        Name: Score, dtype: int64

```

1.6 Information About DataSet

```

In [4]: print("\nNumber of Reviews: ",Data["Text"].count())
        print("\nNumber of Users: ",len(Data["UserId"].unique())) # Unique returns 1-D array o
        print("\nNumber of Products: ",len(Data["ProductId"].unique()))
        print("\nShape of Data: ", Data.shape)
        print("\nColumn Name of DataSet : ",Data.columns)
        print("\n\nNumber of Attributes/Columns in data: 12")
        print("\nNumber of Positive Reviews : ", Data['Score'].value_counts()[1])
        print("\nNumber of Negative Reviews : ", Data['Score'].value_counts()[0])

```

Number of Reviews: 150000

Number of Users: 115632

Number of Products: 43130

Shape of Data: (150000, 12)

```
Column Name of DataSet : Index(['index', 'Id', 'ProductId', 'UserId', 'ProfileName',  
                                'HelpfulnessNumerator', 'HelpfulnessDenominator', 'Score', 'Time',  
                                'Summary', 'Text', 'CleanedText'],  
                                dtype='object')
```

Number of Attributes/Columns in data: 12

Number of Positive Reviews : 126439

Number of Negative Reviews : 23561

```
In [5]: print("\nNumber of Reviews: ",Data["Text"].count())
```

Number of Reviews: 150000

1.7 Attribute Information About DataSet

1.Id - A unique value starts from 1

2.ProductId - A unique identifier for the product

3.UserId - A unique identifier for the user

4.ProfileName - Name of user profile

5.HelpfulnessNumerator - Number of users who found the review helpful

6.HelpfulnessDenominator - Number of users who indicated whether they found the review helpful or not

7.Score - Rating 0 or 1

8.Time - Timestamp for the review

9.Summary - Brief summary of the review

10.Text - Text of the review

11.Cleaned Text - Text that only alphabets

```
In [6]: # Sorting on the basis of Time Parameter  
        Data.sort_values('Time',inplace=True)
```

```
In [7]: Y = Data['Score'].values  
        X = Data['CleanedText'].values
```

1.8 Splitting DataSet into Train and Test Data

```
In [8]: from sklearn.model_selection import train_test_split
        # X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, shuffle=False)
        X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33) # this is random

        print("Shape of Train and Test Dataset for 50k points")
        print(X_train.shape, Y_train.shape)
        print(X_test.shape, Y_test.shape)
```

```
Shape of Train and Test Dataset for 50k points
(100500,) (100500,)
(49500,) (49500,)
```

1.9 Defining Some Function

1.9.1 Train Data Confusion Matrix Plot

```
In [9]: def trainconfusionmatrix(model,X_train,y_train):
        print("Confusion Matrix for Train set")
        cm=confusion_matrix(y_train, model.predict(X_train))
        class_label = ["negative", "positive"]
        df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
        sns.heatmap(df_cm, annot = True, fmt = "d")
        plt.title("Train Confusiion Matrix")
        plt.xlabel("Predicted Label")
        plt.ylabel("True Label")
        plt.show()
```

1.9.2 Test Data Confusion Matrix Plot

```
In [10]: def testconfusionmatrix(model,X_test,y_test):
        print("Confusion Matrix for Test set")
        cm=confusion_matrix(y_test, model.predict(X_test))
        class_label = ["negative", "positive"]
        df_cm = pd.DataFrame(cm, index = class_label, columns = class_label)
        sns.heatmap(df_cm, annot = True, fmt = "d")
        plt.title("Test Confusiion Matrix")
        plt.xlabel("Predicted Label")
        plt.ylabel("True Label")
        plt.show()
```

1.9.3 ROC-AUC Curve Plot

```
In [12]: def plot(Alpha,gsv):

        Res=gsv.cv_results_
        train_auc=Res['mean_train_score']
```

```

train_auc_std=Res['std_train_score']
cv_auc=Res['mean_test_score']
cv_auc_std=Res['std_test_score']

log_alpha=[math.log10(x) for x in Alpha ]
plt.plot(log_alpha, train_auc, label='Train AUC')
plt.gca().fill_between(log_alpha,train_auc - train_auc_std,train_auc + train_auc_std)

plt.plot(log_alpha, cv_auc, label='CV AUC')
plt.gca().fill_between(log_alpha,cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.5)
plt.legend()
plt.xlabel("Log(Alpha): hyperparameter")
plt.ylabel("AUC")
plt.title("Plot Between AUC & Log(Alpha)")
plt.show()

```

1.9.4 GridSearchCV

```

In [13]: def Grid_SearchCV(X_train,Y_train,alpha):

        tscv = TimeSeriesSplit(n_splits=10)
        M_NB = MultinomialNB()

        gsv=GridSearchCV(M_NB,alpha,cv=tscv,verbose=1,scoring='roc_auc')
        gsv.fit(X_train,Y_train)

        return gsv

```

1.9.5 30 Informative Feature

```

In [36]: def show_30_informative_feature(vectorizer,model,n=30):

        # For Negative Class
        neg_class_prob_sorted = model.feature_log_prob_[0, :].argsort()
        neg_feat=[vectorizer.get_feature_names()[x] for x in neg_class_prob_sorted[-n:]]
        neg_prob=[model.feature_log_prob_[0, :][x] for x in neg_class_prob_sorted[-n:]]

        neg_zip=list(zip(neg_feat,neg_prob))
        neg_zip.sort()

        # For Positive Class
        pos_class_prob_sorted = model.feature_log_prob_[1, :].argsort()
        pos_feat=[vectorizer.get_feature_names()[x] for x in pos_class_prob_sorted[-n:]]
        pos_prob=[model.feature_log_prob_[0, :][x] for x in pos_class_prob_sorted[-n:]]
        pos_zip=list(zip(pos_feat,pos_prob))
        pos_zip.sort()

```

```

top=zip(pos_zip,neg_zip)

print("{0:20}{1:55}{2:20}".format("S.N","Positive","Negative"))
print("_"*90)
i=1
for (fn_1,coef_1), (fn_2,coef_2) in top:
    print("%d.\t\t%.3f\t%-30s\t\t%.3f\t%s" % (i,coef_1, fn_1, coef_2, fn_2))
    i+=1

```

1.10 Bags of Words

```

In [15]: vectorizer = CountVectorizer()
         vectorizer.fit(X_train) # fit has to happen only on train data

         # we use the fitted CountVectorizer to convert the text to vector
         X_train_bow = vectorizer.transform(X_train)
         X_train_bow=preprocessing.normalize(X_train_bow)

         X_test_bow = vectorizer.transform(X_test)
         X_test_bow=preprocessing.normalize(X_test_bow)

         print("Shape of Train and Test Data After vectorizations")
         print(X_train_bow.shape, Y_train.shape)
         print(X_test_bow.shape, Y_test.shape)

```

```

Shape of Train and Test Data After vectorizations
(100500, 38189) (100500,)
(49500, 38189) (49500,)

```

1.10.1 Finding the best value Of hyperparameter (Alpha)

```

In [70]: Alpha={'alpha':[1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0005,0.0001]}
         gsv=Grid_SearchCV(X_train_bow,Y_train,Alpha)

         print("Best HyperParameter: ",gsv.best_params_)
         print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))

```

Fitting 10 folds for each of 15 candidates, totalling 150 fits

```

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 150 out of 150 | elapsed: 15.1s finished


```

```

Best HyperParameter: {'alpha': 0.05}
Best Accuracy: 92.22%

```

```
In [71]: plot(Alpha['alpha'],gsv)
```



1.10.2 Training the model


```
In [72]: model_FBOw=MultinomialNB(alpha=gsv.best_params_['alpha'])  
         model_FBOw.fit(X_train_bow,Y_train)
```

```
Out[72]: MultinomialNB(alpha=0.05, class_prior=None, fit_prior=True)
```

1.10.3 Evaluating the performance of model


```
In [73]: trainconfusionmatrix(model_FBOw,X_train_bow,Y_train)
```

Confusion Matrix for Train set




```
In [74]: testconfusionmatrix(model_FBOw,X_test_bow,Y_test)
```

Confusion Matrix for Test set



```
In [75]: plot_auc_roc(model_FBOw,X_train_bow,X_test_bow,Y_train,Y_test)
```



```
In [76]: print("Classification Report: \n")
         y_pred=model_FBOW.predict(X_test_bow)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.92 | 0.19 | 0.31 | 7990 |
| 1 | 0.86 | 1.00 | 0.93 | 41510 |
| micro avg | 0.87 | 0.87 | 0.87 | 49500 |
| macro avg | 0.89 | 0.59 | 0.62 | 49500 |
| weighted avg | 0.87 | 0.87 | 0.83 | 49500 |

1.10.4 Displaying 30 most informative feature

```
In [77]: show_30_informative_feature(vectorizer,model_FBOW)
```

| S.N | Positive | Negative | | |
|-----|----------|----------|--------|-----|
| 1. | -5.943 | also | -5.412 | ama |
| 2. | -5.412 | amazon | -5.505 | ba |
| 3. | -6.689 | best | -5.319 | bo |
| 4. | -5.004 | buy | -5.004 | bu |
| 5. | -5.043 | coffe | -5.043 | co |
| 6. | -5.238 | dont | -5.536 | di |
| 7. | -5.507 | eat | -5.238 | do |
| 8. | -6.103 | find | -5.507 | ea |
| 9. | -4.750 | flavor | -5.372 | ev |
| 10. | -5.395 | food | -4.750 | f |
| 11. | -5.145 | get | -5.395 | f |
| 12. | -5.004 | good | -5.145 | g |
| 13. | -5.886 | great | -5.004 | g |
| 14. | -4.260 | like | -4.260 | l |
| 15. | -6.045 | littl | -5.588 | l |
| 16. | -5.524 | love | -5.524 | l |
| 17. | -5.544 | make | -5.544 | m |
| 18. | -5.474 | much | -5.474 | m |
| 19. | -4.746 | one | -4.746 | o |
| 20. | -5.065 | order | -5.065 | o |
| 21. | -5.688 | price | -5.579 | p |
| 22. | -4.366 | product | -4.366 | p |

| | | | | |
|-----|--------|--------|--------|----|
| 23. | -5.509 | realli | -5.609 | pr |
| 24. | -5.986 | store | -5.509 | re |
| 25. | -4.148 | tast | -4.148 | ta |
| 26. | -5.423 | tea | -5.423 | te |
| 27. | -5.485 | time | -5.485 | ti |
| 28. | -4.885 | tri | -4.885 | tr |
| 29. | -5.072 | use | -5.072 | us |
| 30. | -4.857 | would | -4.857 | w |

1.11 TF-IDF

```
In [83]: vectorizer_tfidf=TfidfVectorizer()
         vectorizer_tfidf.fit(X_train)
```

```
Out[83]: TfidfVectorizer(analyzer='word', binary=False, decode_error='strict',
                        dtype=<class 'numpy.float64'>, encoding='utf-8', input='content',
                        lowercase=True, max_df=1.0, max_features=None, min_df=1,
                        ngram_range=(1, 1), norm='l2', preprocessor=None, smooth_idf=True,
                        stop_words=None, strip_accents=None, sublinear_tf=False,
                        token_pattern='(?u)\\b\\w\\w+\\b', tokenizer=None, use_idf=True,
                        vocabulary=None)
```

```
In [84]: X_Train_Tfidf=vectorizer_tfidf.transform(X_train)
         X_Train_Tfidf=preprocessing.normalize(X_Train_Tfidf)

         X_Test_Tfidf=vectorizer_tfidf.transform(X_test)
         X_Test_Tfidf=preprocessing.normalize(X_Test_Tfidf)
```

```
In [85]: print("Shape of Train and Test Data After vectorizations")
         print(X_Train_Tfidf.shape, Y_train.shape)
         print(X_Test_Tfidf.shape, Y_test.shape)
```

```
Shape of Train and Test Data After vectorizations
(100500, 38440) (100500,)
(49500, 38440) (49500,)
```

1.11.1 Finding the best value of hyperparameter(Alpha)

```
In [86]: Alpha={'alpha':[1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0005,0.0001]}
         gsv=Grid_SearchCV(X_Train_Tfidf,Y_train,Alpha)


         print("Best HyperParameter: ",gsv.best_params_)
         print("Best Accuracy: %.2f%%"%(gsv.best_score_*100))
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done 150 out of 150 | elapsed: 15.3s finished
```

```
Best HyperParameter: {'alpha': 0.1}  
Best Accuracy: 89.89%
```

```
In [87]: plot(Alpha['alpha'],gsv)
```



1.11.2 Training the model

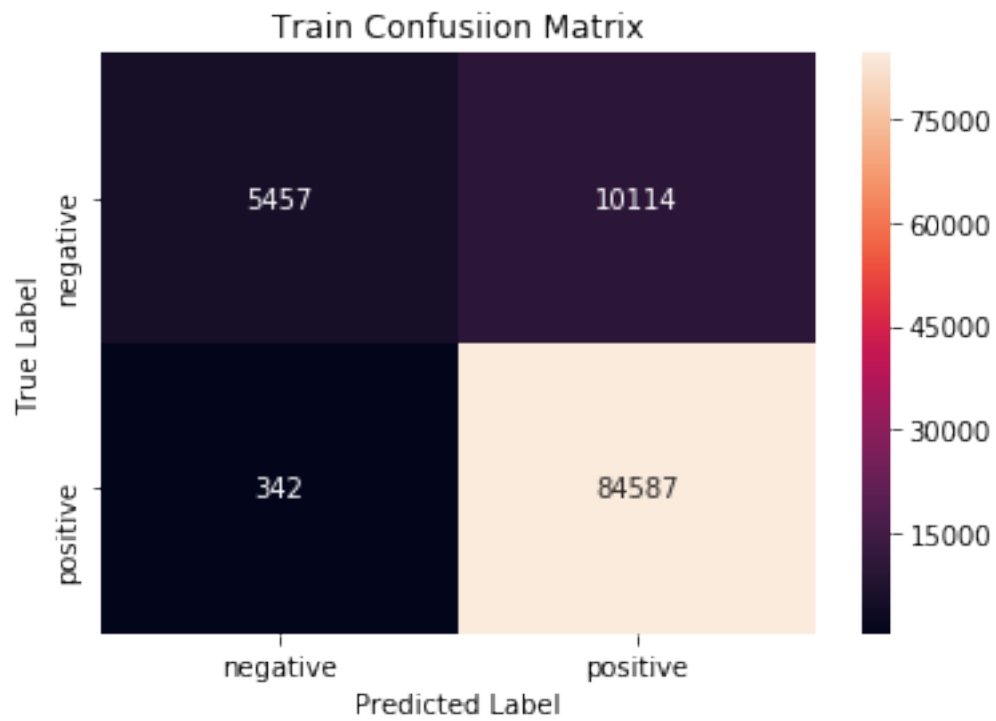
```
In [88]: model_Tfidf=MultinomialNB(alpha=gsv.best_params_['alpha'])  
         model_Tfidf.fit(X_Train_Tfidf,Y_train)
```

```
Out[88]: MultinomialNB(alpha=0.1, class_prior=None, fit_prior=True)
```

1.11.3 Evaluating the performance of model


```
In [89]: trainconfusionmatrix(model_Tfidf,X_Train_Tfidf,Y_train)
```

Confusion Matrix for Train set




```
In [90]: testconfusionmatrix(model_Tfidf,X_Test_Tfidf,Y_test)
```

Confusion Matrix for Test set



In [91]: `plot_auc_roc(model_Tfidf,X_Train_Tfidf,X_Test_Tfidf,Y_train,Y_test)`



```
In [92]: print("Classification Report: \n")
         y_pred=model_Tfidf.predict(X_Test_Tfidf)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.89 | 0.24 | 0.37 | 7990 |
| 1 | 0.87 | 0.99 | 0.93 | 41510 |
| micro avg | 0.87 | 0.87 | 0.87 | 49500 |
| macro avg | 0.88 | 0.61 | 0.65 | 49500 |
| weighted avg | 0.87 | 0.87 | 0.84 | 49500 |

1.11.4 Displaying 30 most informative feature

```
In [93]: show_30_informative_feature(vectorizer_tfidf,model_Tfidf)
```

| S.N | Positive | | Negative | |
|-----|----------|---------|----------|-----|
| 1. | -5.902 | amazon | -5.902 | ama |
| 2. | -7.158 | best | -5.850 | ba |
| 3. | -5.529 | buy | -5.856 | ba |
| 4. | -5.401 | coffe | -5.951 | bo |
| 5. | -5.940 | dog | -5.642 | bo |
| 6. | -6.222 | drink | -5.529 | bu |
| 7. | -5.965 | eat | -5.401 | co |
| 8. | -6.562 | find | -5.657 | di |
| 9. | -5.409 | flavor | -5.940 | do |
| 10. | -5.827 | food | -5.711 | d |
| 11. | -5.749 | get | -5.965 | ea |
| 12. | -5.729 | good | -5.813 | ex |
| 13. | -6.584 | great | -5.409 | f |
| 14. | -5.054 | like | -5.827 | f |
| 15. | -6.489 | littl | -5.749 | g |
| 16. | -6.247 | love | -5.729 | g |
| 17. | -6.146 | make | -5.054 | l |
| 18. | -5.945 | much | -5.953 | l |
| 19. | -5.428 | one | -5.968 | m |
| 20. | -5.528 | order | -5.945 | m |
| 21. | -6.138 | price | -5.428 | o |
| 22. | -5.048 | product | -5.528 | o |

| | | | | |
|-----|--------|--------|--------|----|
| 23. | -5.980 | realli | -5.914 | pa |
| 24. | -6.387 | store | -5.048 | pr |
| 25. | -4.916 | tast | -5.938 | pr |
| 26. | -5.728 | tea | -4.916 | ta |
| 27. | -5.999 | time | -5.728 | ta |
| 28. | -5.528 | tri | -5.528 | tr |
| 29. | -5.754 | use | -5.754 | us |
| 30. | -5.390 | would | -5.390 | w |

1.12 Addition of another column length

```
In [16]: Train_len=[]
        Test_len=[]
        for i in X_train:
            Train_len.append(len(i))

        for i in X_test:
            Test_len.append(len(i))

In [17]: Train_len=np.array(Train_len)
        Test_len=np.array(Test_len)

In [18]: Train_len=Train_len[:,np.newaxis]
        Test_len=Test_len[:,np.newaxis]
```

1.13 Bag Of Words

```
In [19]: X_Train_BOW=X_train_bow.todense()

In [20]: X_Train_New=np.append(X_Train_BOW,Train_len,axis=1)

In [21]: from scipy.sparse import csr_matrix
        X_Train_New= csr_matrix(X_Train_New)

In [22]: print("Shape of Train Data Before Adding length column ")
        print(X_train_bow.shape)

        print("\nShape of Train Data After Adding length column ")
        print(X_Train_New.shape)
```

Shape of Train Data Before Adding length column
(100500, 38189)

Shape of Train Data After Adding length column
(100500, 38190)

1.13.1 Finding the best value of hyperparameter (Alpha)

```
In [23]: Alpha={'alpha':[1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0005,0.0001]}
        gsv1=Grid_SearchCV(X_Train_New,Y_train,Alpha)


        print("Best HyperParameter: ",gsv1.best_params_)
        print("Best Accuracy: %.2f%%"%(gsv1.best_score_*100))
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.
[Parallel(n_jobs=1)]: Done 150 out of 150 | elapsed: 31.2s finished
```

```
Best HyperParameter: {'alpha': 0.05}
Best Accuracy: 91.83%
```

```
In [24]: plot(Alpha['alpha'],gsv1)
```



```
In [25]: X_Test_Bow=X_test_bow.todense()
```

```
In [26]: X_Test_New=np.append(X_Test_Bow,Test_len,axis=1)
```

```
In [27]: X_Test_New=csr_matrix(X_Test_New)

In [28]: print("Shape of Test Data Before Adding length column ")
          print(X_test_bow.shape)

          print("\nShape of Test Data After Adding length column ")
          print(X_Test_New.shape)
```

Shape of Test Data Before Adding length column
(49500, 38189)

Shape of Test Data After Adding length column
(49500, 38190)

1.13.2 Training the model


```
In [29]: model_New_Bow=MultinomialNB(alpha=gsv1.best_params_['alpha'])
          model_New_Bow.fit(X_Train_New,Y_train)
```

```
Out[29]: MultinomialNB(alpha=0.05, class_prior=None, fit_prior=True)
```

1.13.3 Evaluating the performance of model


```
In [30]: trainconfusionmatrix(model_New_Bow,X_Train_New,Y_train)
```

Confusion Matrix for Train set




```
In [31]: testconfusionmatrix(model_New_Bow,X_Train_New,Y_train)
```

Confusion Matrix for Test set



```
In [32]: plot_auc_roc(model_New_Bow,X_Train_New,X_Test_New,Y_train,Y_test)
```



```
In [33]: print("Classification Report: \n")
         y_pred=model_New_Bow.predict(X_Test_New)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.20 | 0.33 | 7793 |
| 1 | 0.87 | 0.99 | 0.93 | 41707 |
| micro avg | 0.87 | 0.87 | 0.87 | 49500 |
| macro avg | 0.88 | 0.60 | 0.63 | 49500 |
| weighted avg | 0.87 | 0.87 | 0.83 | 49500 |

```
In [37]: show_30_informative_feature(vectorizer,model_New_Bow)
```

| S.N | Positive | | Negative | |
|-----|----------|--------|----------|--------|
| 1. | -9.346 | amazon | -9.346 | amazon |
| 2. | -10.598 | best | -9.371 | best |

| | | | | |
|-----|--------|---------|--------|----|
| 3. | -8.908 | buy | -9.238 | bo |
| 4. | -8.942 | coffe | -8.908 | bu |
| 5. | -9.124 | dont | -8.942 | co |
| 6. | -9.384 | eat | -9.453 | di |
| 7. | -9.974 | find | -9.124 | do |
| 8. | -8.649 | flavor | -9.384 | ea |
| 9. | -9.312 | food | -9.299 | ev |
| 10. | -9.031 | get | -8.649 | f |
| 11. | -8.871 | good | -9.312 | f |
| 12. | -9.765 | great | -9.031 | g |
| 13. | -8.148 | like | -8.871 | g |
| 14. | -9.913 | littl | -8.148 | l |
| 15. | -9.423 | love | -9.468 | l |
| 16. | -9.437 | make | -9.423 | l |
| 17. | -9.363 | much | -9.437 | ma |
| 18. | -8.635 | one | -9.363 | mi |
| 19. | -8.943 | order | -8.635 | or |
| 20. | -9.581 | price | -8.943 | or |
| 21. | -8.261 | product | -9.479 | pa |
| 22. | -9.392 | realli | -8.261 | pa |
| 23. | -9.897 | store | -9.490 | pa |
| 24. | -8.042 | tast | -9.392 | re |
| 25. | -9.262 | tea | -8.042 | ta |
| 26. | -9.397 | time | -9.262 | ta |
| 27. | -8.790 | tri | -9.397 | ti |
| 28. | -8.961 | use | -8.790 | ti |
| 29. | -9.922 | well | -8.961 | us |
| 30. | -8.746 | would | -8.746 | w |

1.14 TF-IDF

```
In [30]: X_Train_New_Tf=X_Train_Tfidf.todense()
```

```
In [31]: X_Train_New_Tf1=np.append(X_Train_New_Tf,Train_len,axis=1)
```

```
In [32]: from scipy.sparse import csr_matrix
X_Train_New_Tf1=csr_matrix(X_Train_New_Tf1)
```

1.14.1 Finding the best value of hyperparameter (Alpha)

```
In [33]: Alpha={'alpha':[1000,500,100,50,10,5,1,0.5,0.1,0.05,0.01,0.005,0.001,0.0005,0.0001]}
gsv2=Grid_SearchCV(X_Train_New_Tf1,Y_train,Alpha)

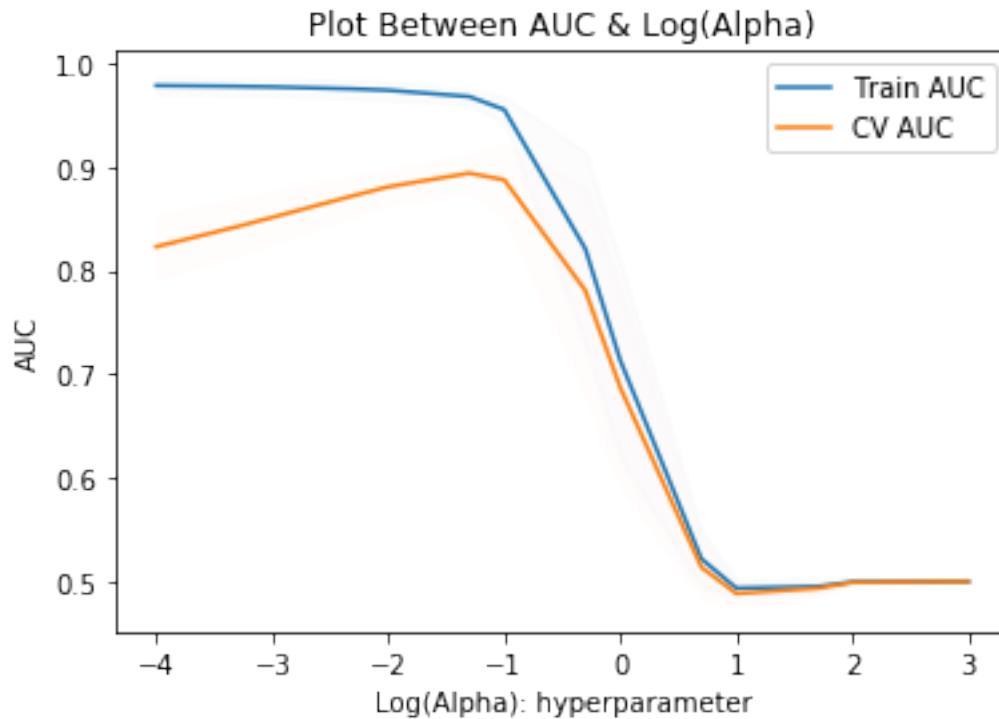
print("Best HyperParameter: ",gsv2.best_params_)
print("Best Accuracy: %.2f%%"%(gsv2.best_score_*100))
```

Fitting 10 folds for each of 15 candidates, totalling 150 fits

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done 150 out of 150 | elapsed: 24.7s finished
```

```
Best HyperParameter: {'alpha': 0.05}  
Best Accuracy: 89.40%
```

```
In [34]: plot(Alpha['alpha'],gsv2)
```



```
In [35]: X_Test_New_Tf1=X_Test_Tfidf.todense()
```

```
In [36]: X_Test_New_Tf1=np.append(X_Test_New_Tf1,Test_len,axis=1)
```

```
In [37]: from scipy.sparse import csr_matrix  
X_Test_New_Tf1=csr_matrix(X_Test_New_Tf1)
```

1.14.2 Training the model


```
In [38]: model_New_Tfidf1=MultinomialNB(alpha=gsv2.best_params_['alpha'])  
model_New_Tfidf1.fit(X_Train_New_Tf1,Y_train)
```

```
Out[38]: MultinomialNB(alpha=0.05, class_prior=None, fit_prior=True)
```

1.14.3 Evaluating the performance of model


```
In [39]: trainconfusionmatrix(model_New_Tfidf1,X_Train_New_Tf1,Y_train)
```

Confusion Matrix for Train set




```
In [40]: testconfusionmatrix(model_New_Tfidf1,X_Test_New_Tf1,Y_test)
```

Confusion Matrix for Test set



```
In [41]: plot_auc_roc(model_New_Tfidf1,X_Train_New_Tf1,X_Test_New_Tf1,Y_train,Y_test)
```




```
In [42]: print("Classification Report: \n")
         y_pred=model_New_Tfidf1.predict(X_Test_New_Tf1)

         print(classification_report(Y_test, y_pred))
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.25 | 0.39 | 7990 |
| 1 | 0.87 | 0.99 | 0.93 | 41510 |
| micro avg | 0.87 | 0.87 | 0.87 | 49500 |
| macro avg | 0.86 | 0.62 | 0.66 | 49500 |
| weighted avg | 0.87 | 0.87 | 0.84 | 49500 |

1.14.4 Displaying 30 most informative feature

```
In [67]: show_30_informative_feature(vectorizer_tfidf,model_New_Tfidf1)
```

| S.N | Positive | Negative | | |
|-----|----------|----------|--------|-----|
| 1. | -9.807 | amazon | -9.807 | ama |
| 2. | -11.064 | best | -9.755 | ba |
| 3. | -9.434 | buy | -9.760 | bag |
| 4. | -9.306 | coffe | -9.856 | bo |
| 5. | -9.845 | dog | -9.547 | bo |
| 6. | -10.127 | drink | -9.434 | bu |
| 7. | -9.870 | eat | -9.306 | co |
| 8. | -10.467 | find | -9.562 | d. |
| 9. | -9.313 | flavor | -9.845 | dog |
| 10. | -9.732 | food | -9.616 | do |
| 11. | -9.654 | get | -9.870 | ea |
| 12. | -9.634 | good | -9.717 | ev |
| 13. | -10.489 | great | -9.313 | : |
| 14. | -8.958 | like | -9.732 | f |
| 15. | -10.394 | littl | -9.654 | g |
| 16. | -10.152 | love | -9.634 | g |
| 17. | -10.050 | make | -8.958 | : |
| 18. | -9.850 | much | -9.858 | le |
| 19. | -9.333 | one | -9.872 | me |
| 20. | -9.433 | order | -9.850 | mu |
| 21. | -10.043 | price | -9.333 | e |
| 22. | -8.953 | product | -9.433 | or |

| | | | | |
|-----|---------|--------|--------|----|
| 23. | -9.885 | realli | -9.819 | pa |
| 24. | -10.292 | store | -8.953 | J |
| 25. | -8.821 | tast | -9.843 | pr |
| 26. | -9.633 | tea | -8.821 | ta |
| 27. | -9.904 | time | -9.633 | te |
| 28. | -9.433 | tri | -9.433 | ti |
| 29. | -9.659 | use | -9.659 | us |
| 30. | -9.295 | would | -9.295 | w |

2 Conclusion :

1.Report On Different Vectorizer Method

```
In [1]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Hyperparameter(Alpha)", "Train AUC", "Test AUC", "F1-Score"]

x.add_row(["BOW", 0.05, 0.95, 0.93, 0.83])
x.add_row(["TF-IDF", 0.1, 0.95, 0.91, 0.84])

print(x)
```

| Vectorizer | Hyperparameter(Alpha) | Train AUC | Test AUC | F1-Score |
|------------|-----------------------|-----------|----------|----------|
| BOW | 0.05 | 0.95 | 0.93 | 0.83 |
| TF-IDF | 0.1 | 0.95 | 0.91 | 0.84 |

2.Report On Different Vectorizer Method After Addition Of Length as another Column

```
In [2]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Vectorizer", "Hyperparameter(Alpha)", "Train AUC", "Test AUC", "F1-Score"]

x.add_row(["BOW", 0.05, 0.94, 0.92, 0.83])
x.add_row(["TF-IDF", 0.05, 0.95, 0.91, 0.84])

print(x)
```

| Vectorizer | Hyperparameter(Alpha) | Train AUC | Test AUC | F1-Score |
|------------|-----------------------|-----------|----------|----------|
| BOW | 0.05 | 0.94 | 0.92 | 0.83 |
| TF-IDF | 0.05 | 0.95 | 0.91 | 0.84 |

| | | | | | | |
|--|--------|--|------|--|------|--|
| | | | | | | |
| | BOW | | 0.05 | | 0.94 | |
| | TF-IDF | | 0.05 | | 0.95 | |
| | | | | | 0.92 | |
| | | | | | 0.83 | |
| | | | | | 0.91 | |
| | | | | | 0.84 | |

3. I have taken considerable amount of data but it did not take long time in execution .

4. Since data is unbalanced , i did time based splitting and used roc_auc metric as scoring parameter in GridsearchCV .

5. After adding Length as another column , there is no any improvement.